

# C++ Real-Time Audio Programming with Bela

Dr Andrew McPherson

---

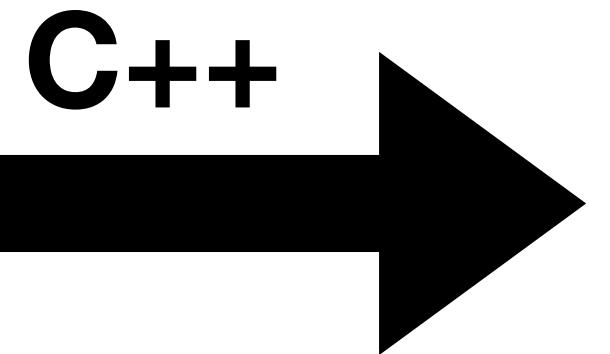
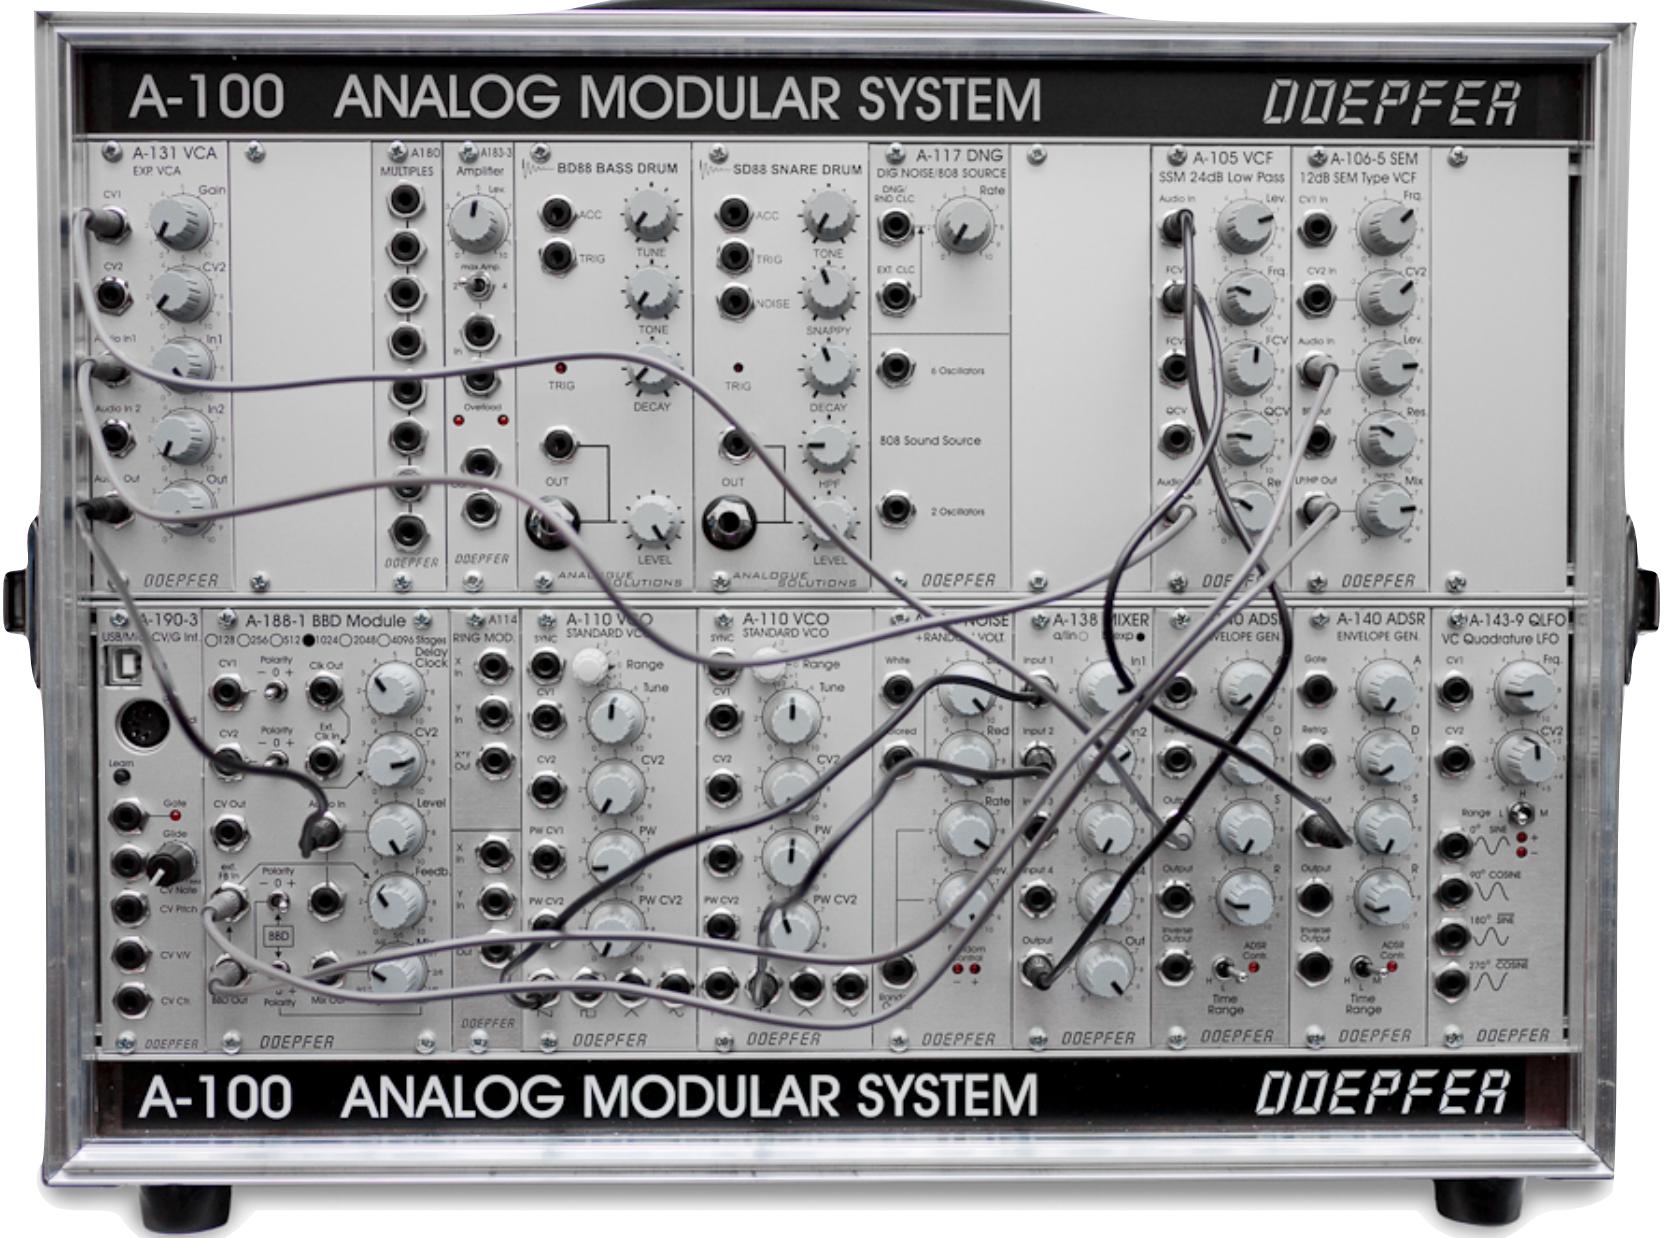
Centre for Digital Music  
School of Electronic Engineering and Computer Science  
Queen Mary University of London

---

Founder and Director, Bela

# C++ Real-Time Audio Programming with Bela

Modular synthesis



Embedded hardware

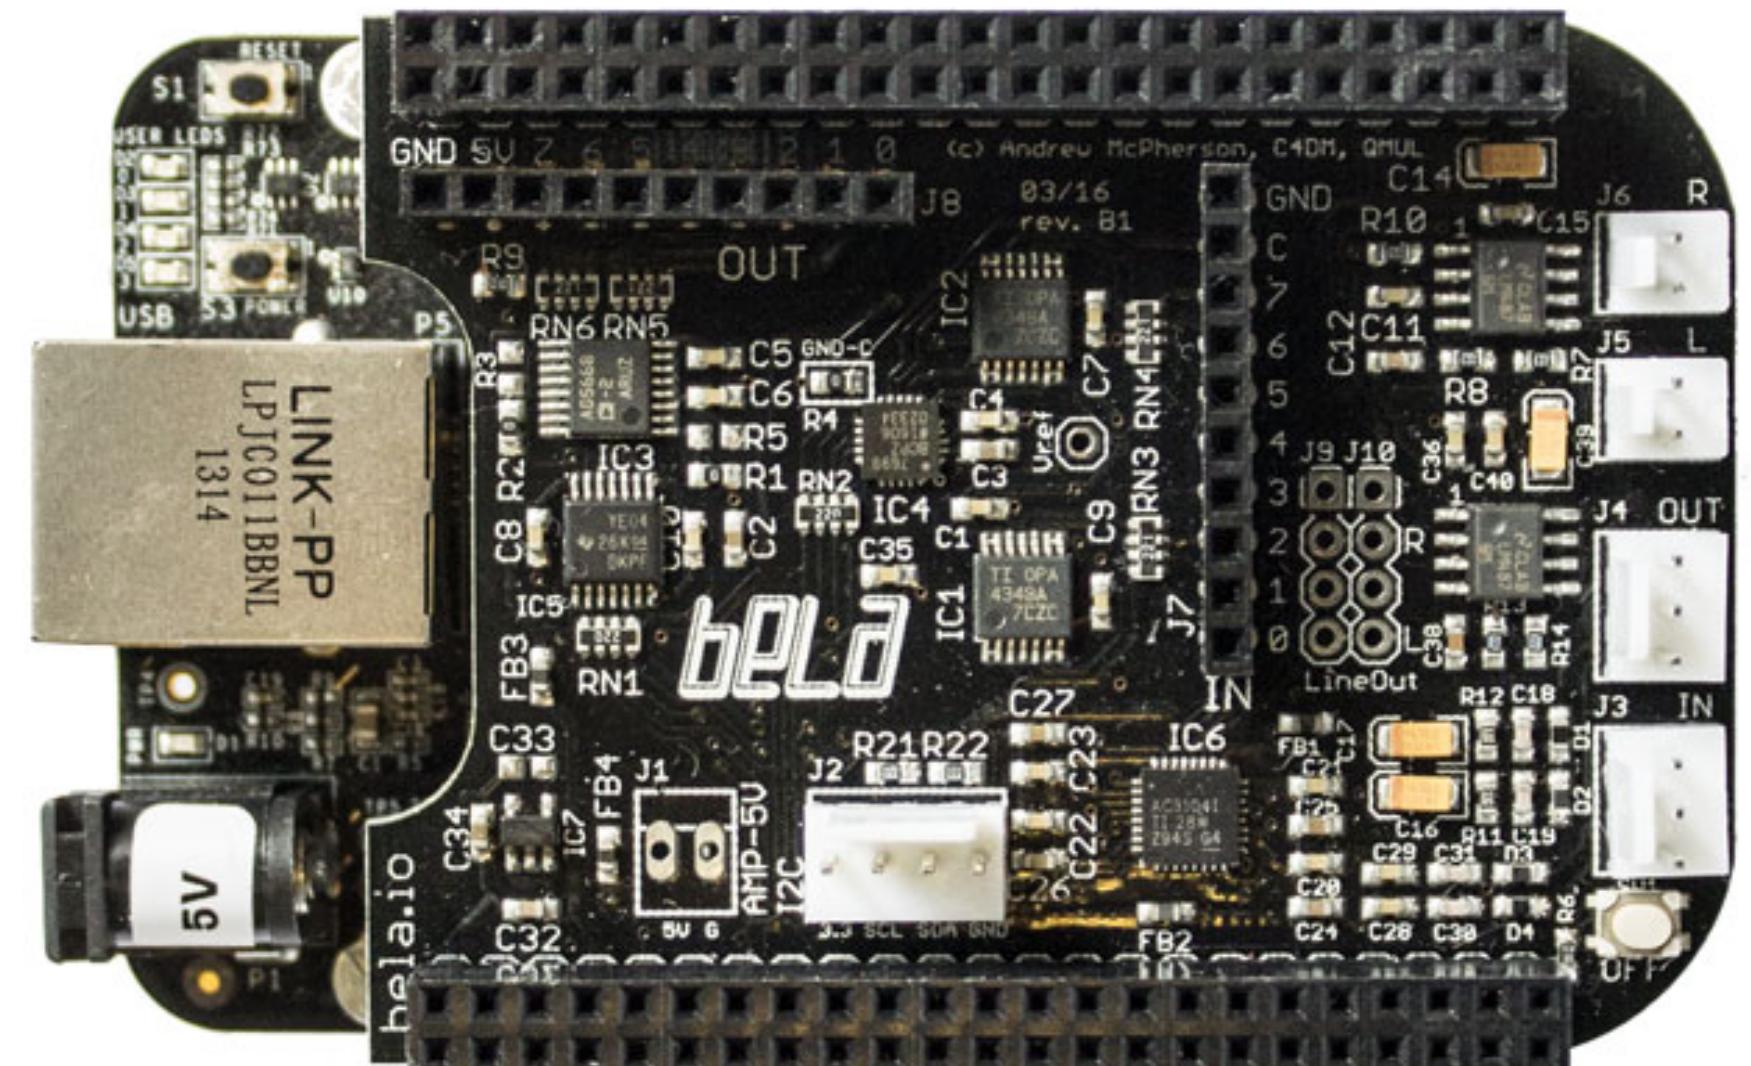


Image credit: Nina Richards, wikipedia (CC-BY 3.0)

# Course topics

## Programming topics

- Working in real time
- Buffers and arrays
- Parameter control
- Classes and objects
- Analog and digital I/O
- Filtering
- Circular buffers
- Timing in real time
- State machines
- MIDI
- Block-based processing
- Threads
- Fixed point arithmetic
- ARM assembly language



Today

## Music/audio topics

- Oscillators
- Samples
- Wavetables
- Filters
- Control voltages
- Gates and triggers
- Delays and delay-based effects
- Metronomes and clocks
- Envelopes
- ADSR
- MIDI
- Additive synthesis
- Phase vocoders
- Impulse reverb

# Lecture 3: Wavetables

**What you'll learn today:**

Working with buffers (arrays)

Interpolation

The Bela oscilloscope

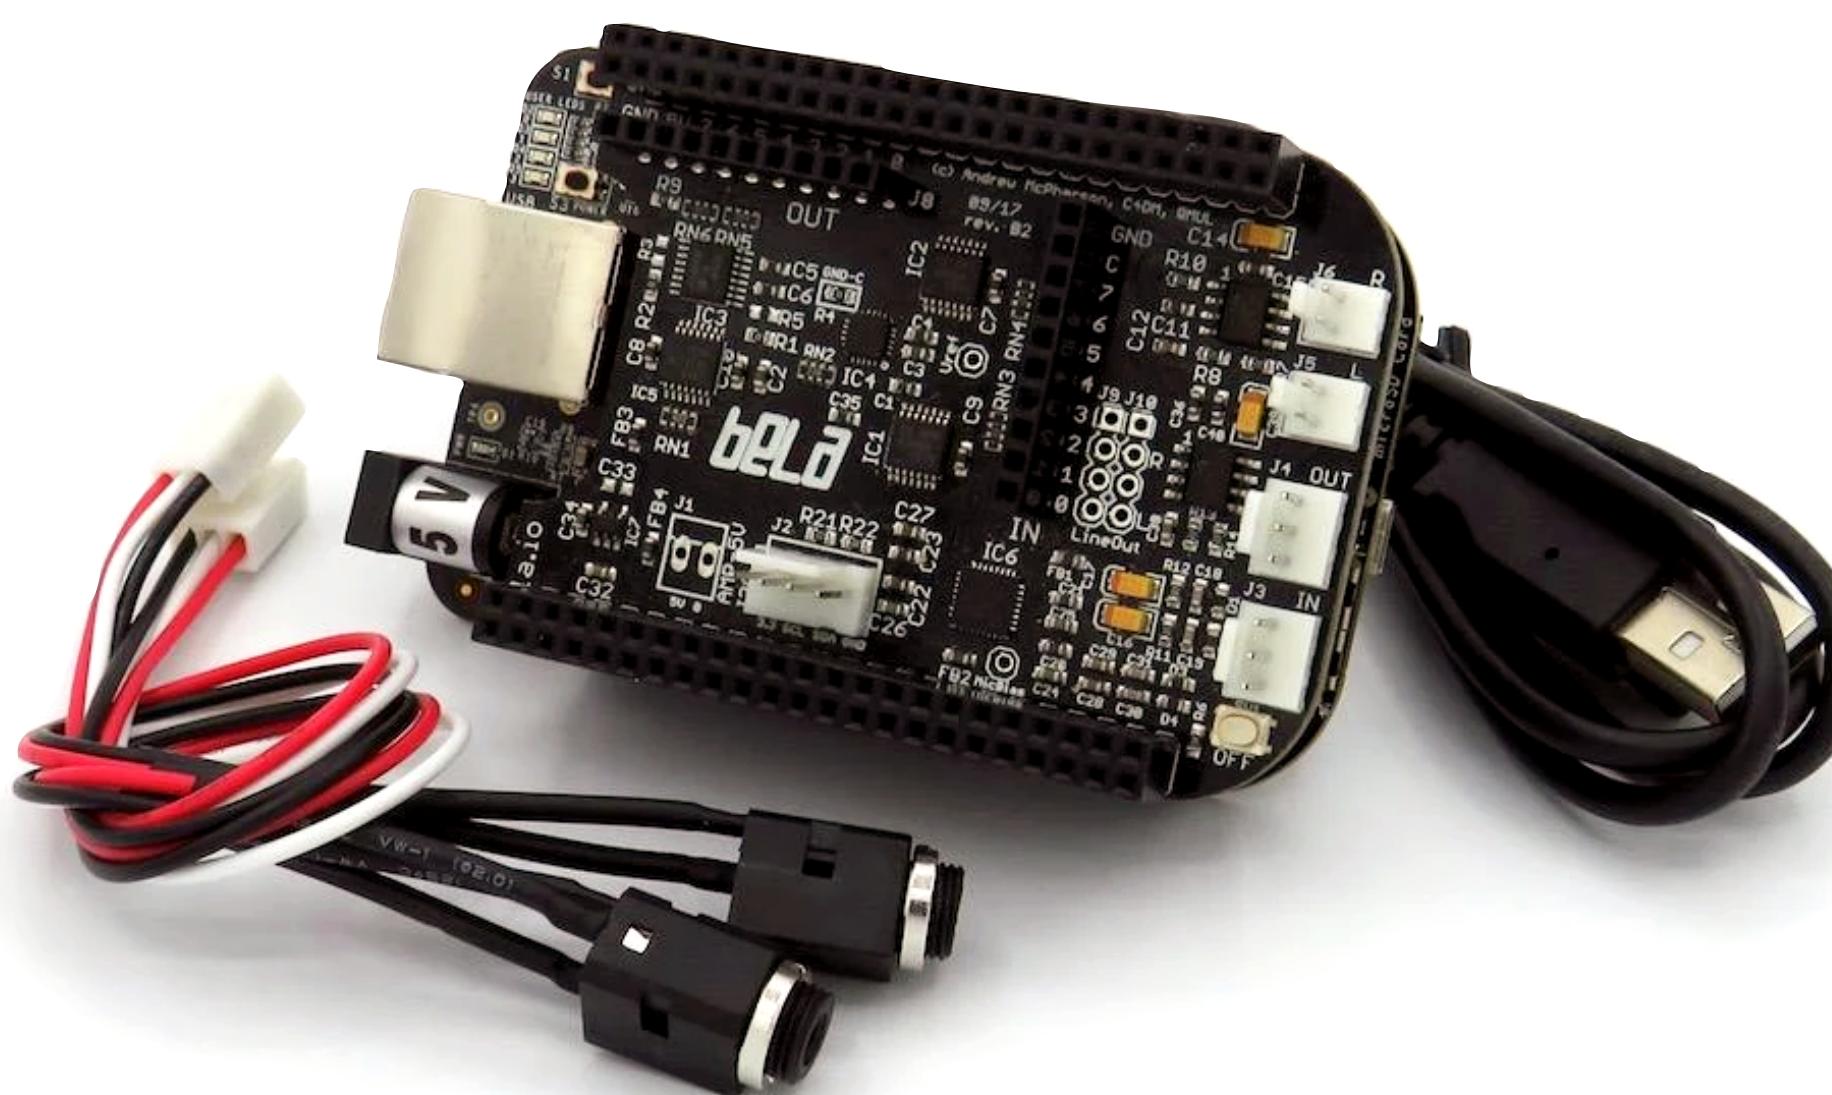
**What you'll make today:**

Wavetable oscillator

**Companion materials:**

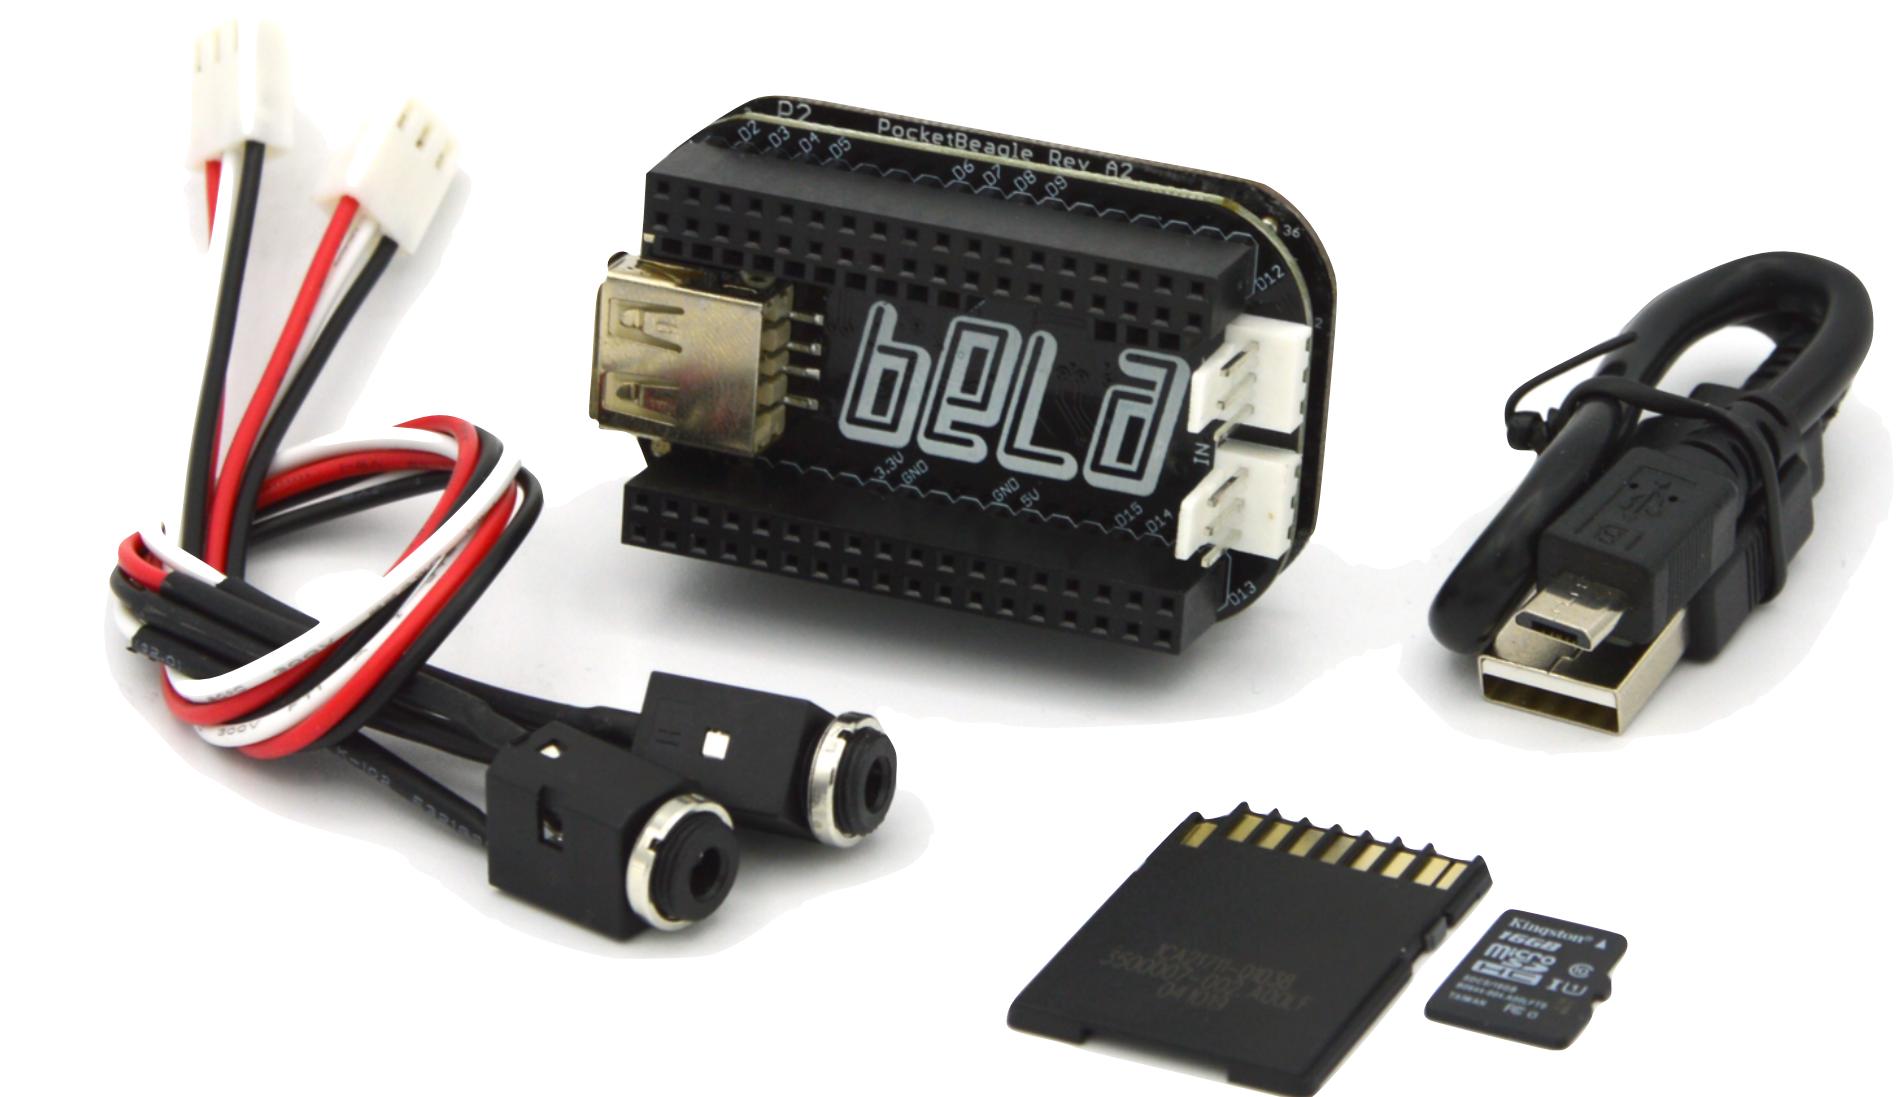
**[github.com/BelaPlatform/bela-online-course](https://github.com/BelaPlatform/bela-online-course)**

# What you'll need



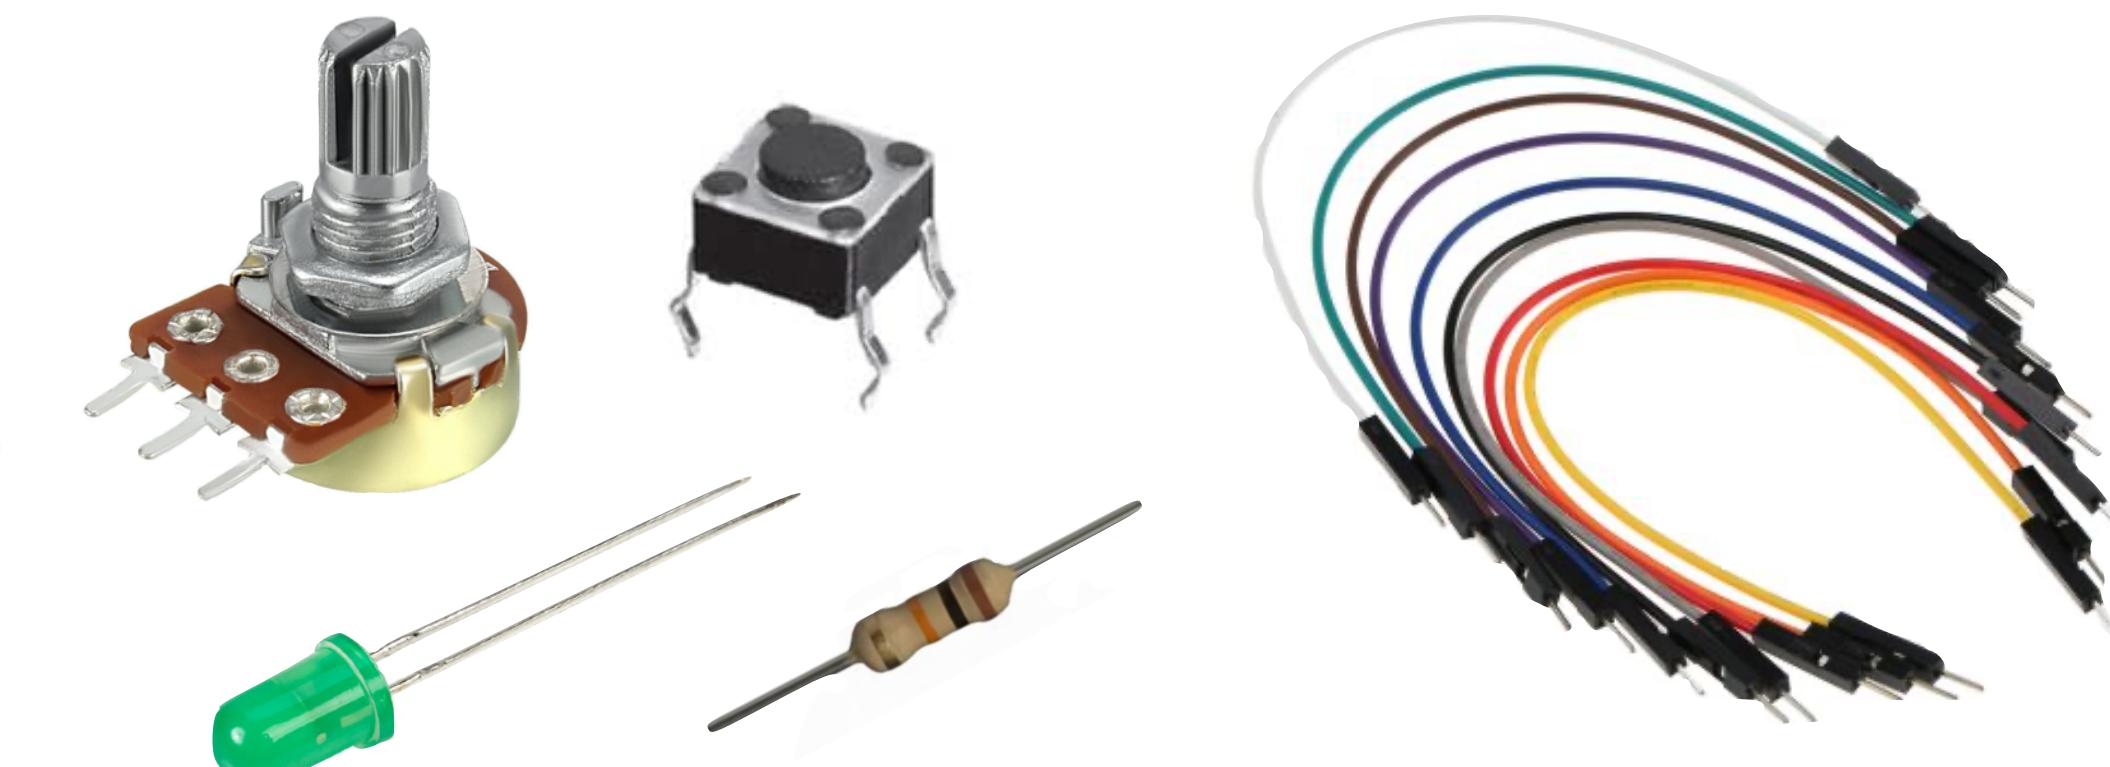
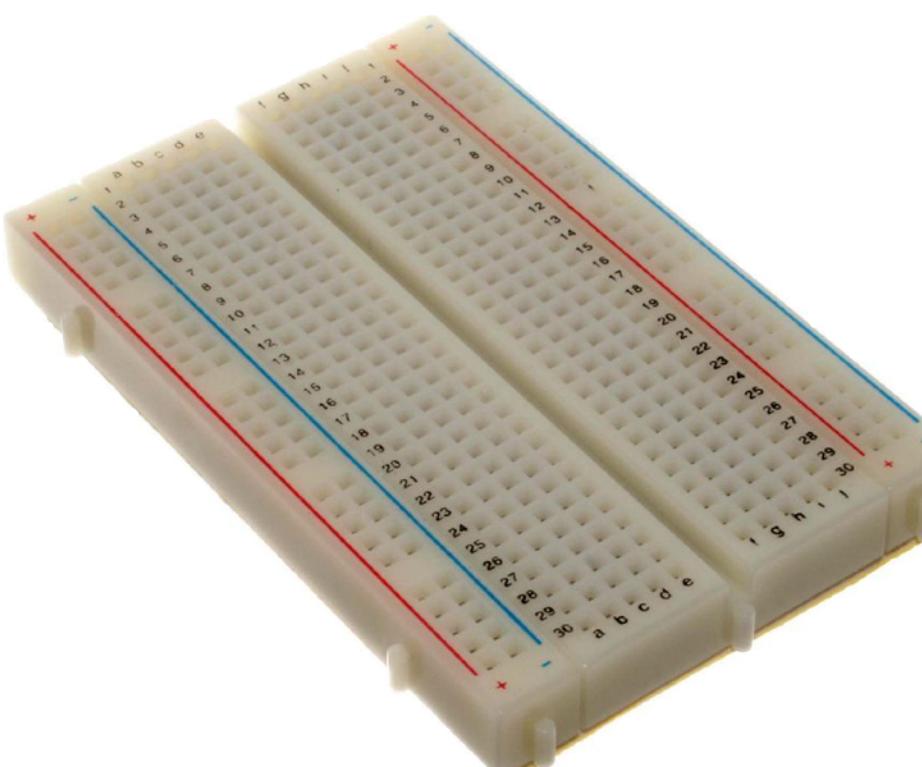
Bela Starter Kit

or



Bela Mini Starter Kit

Recommended  
for some lectures:



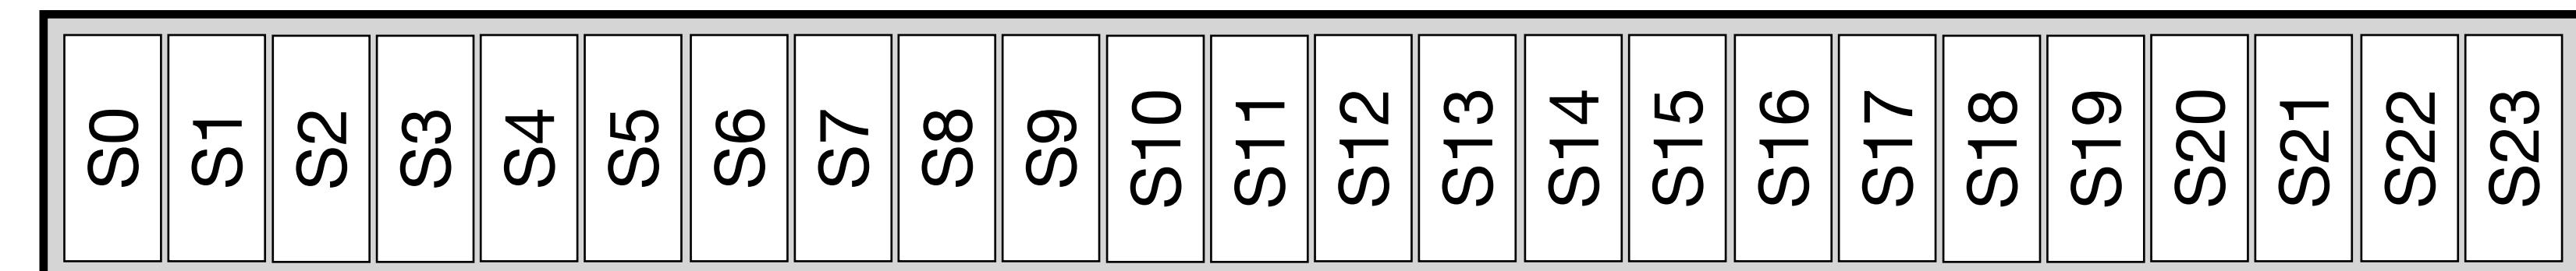
# The Bela C / C++ API

- Every Bela program has three main functions:
- `bool setup(BelaContext *context, void *userData)`
  - Runs **once at the beginning**, before audio starts
  - Data structure `context` holds info on channels, sample rates, block sizes
  - Return `true` if initialisation was successful (`false` stops the program)
- `void render(BelaContext *context, void *userData)`
  - The **audio callback** function
  - Called automatically by the Bela system for **each new block**
  - Where most of your code goes: process the samples and return as quickly as possible!
- `void cleanup(BelaContext *context, void *userData)`
  - Runs **once at the end**, when the program stops
  - Use to release any resources you allocated in `setup()`

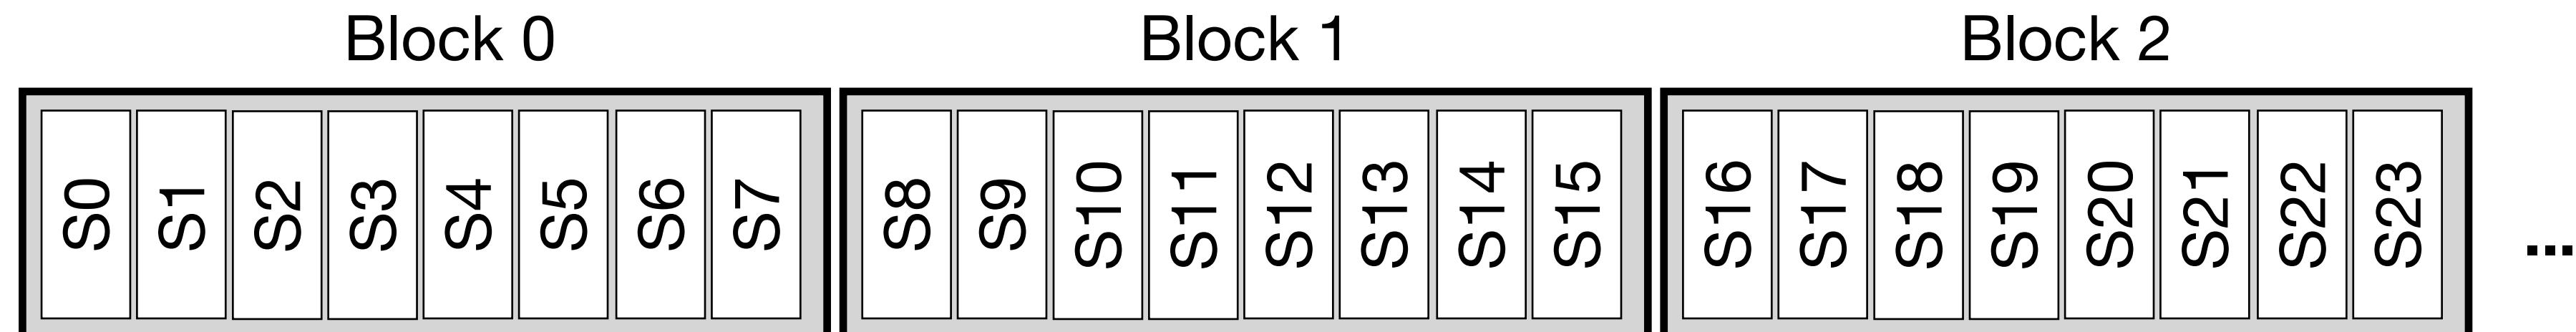
# Playing audio samples

- Suppose we want to play a clip of recorded audio
  - For example, something that we load from a WAV file
  - Let's not worry yet about how we load the sound from storage
- The sound, once loaded, will be stored in a **buffer (i.e. array)**

```
float buffer[24];
```



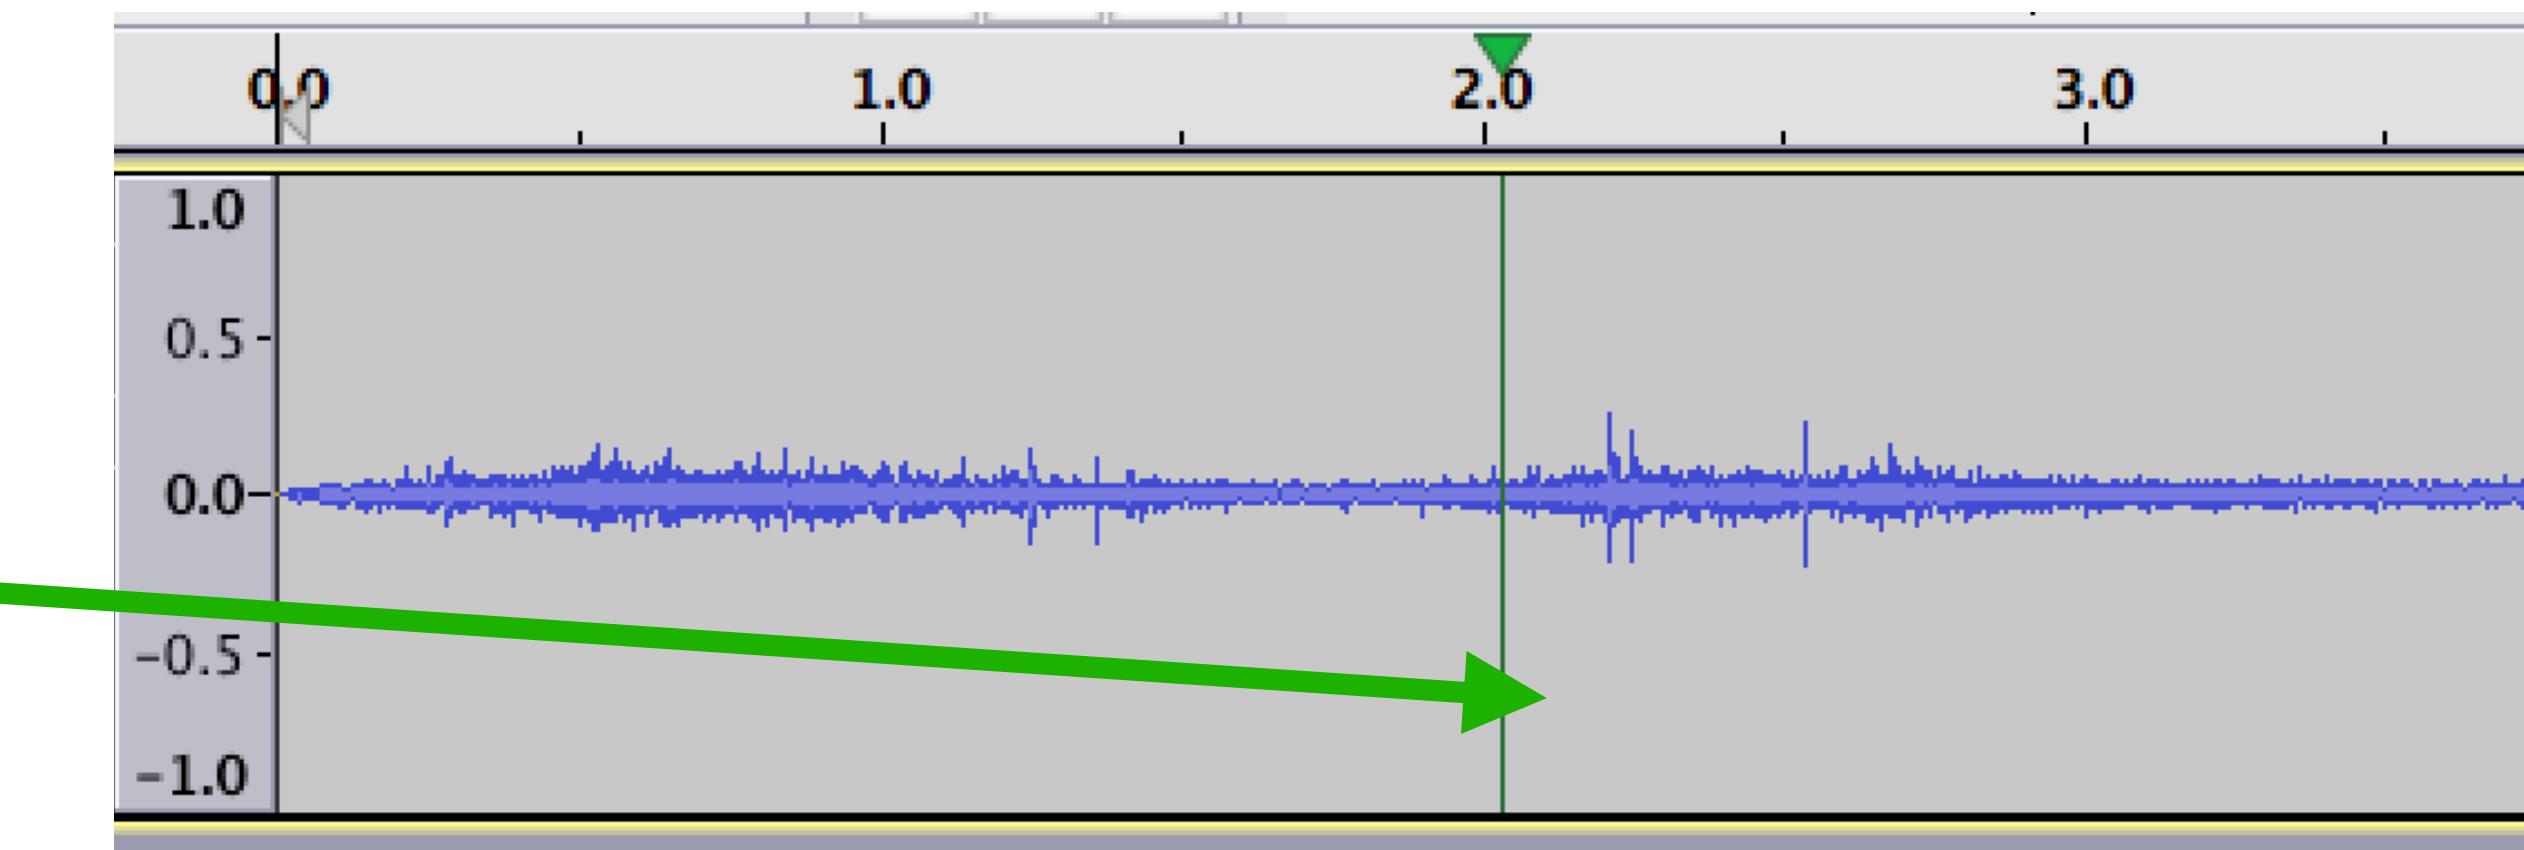
- The buffer is almost certainly **longer than the block size** of our real-time system
- We need to copy these samples **block-by-block** into our audio output



- Besides the buffer itself, **what do we need to keep track of** to play this sound in real time?

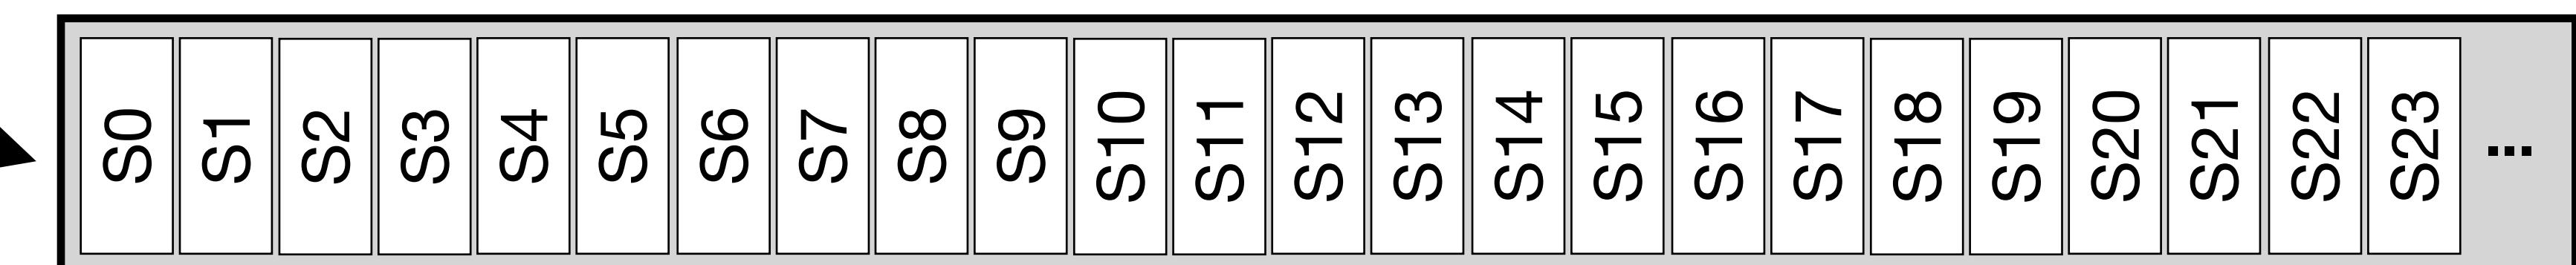
# Keeping track of position

- Audio editors show the play head:
  - The location in the buffer being played
  - As the sound plays, the play head moves forward at constant speed



- In DSP programming, the play head is a special case of a read pointer
  - A reference indicating which index in the buffer to play next
  - Just like the phase of an oscillator, we want to remember this from one block to the next
    - So it should be a global variable!

the buffer doesn't change as we play it



but the read pointer moves

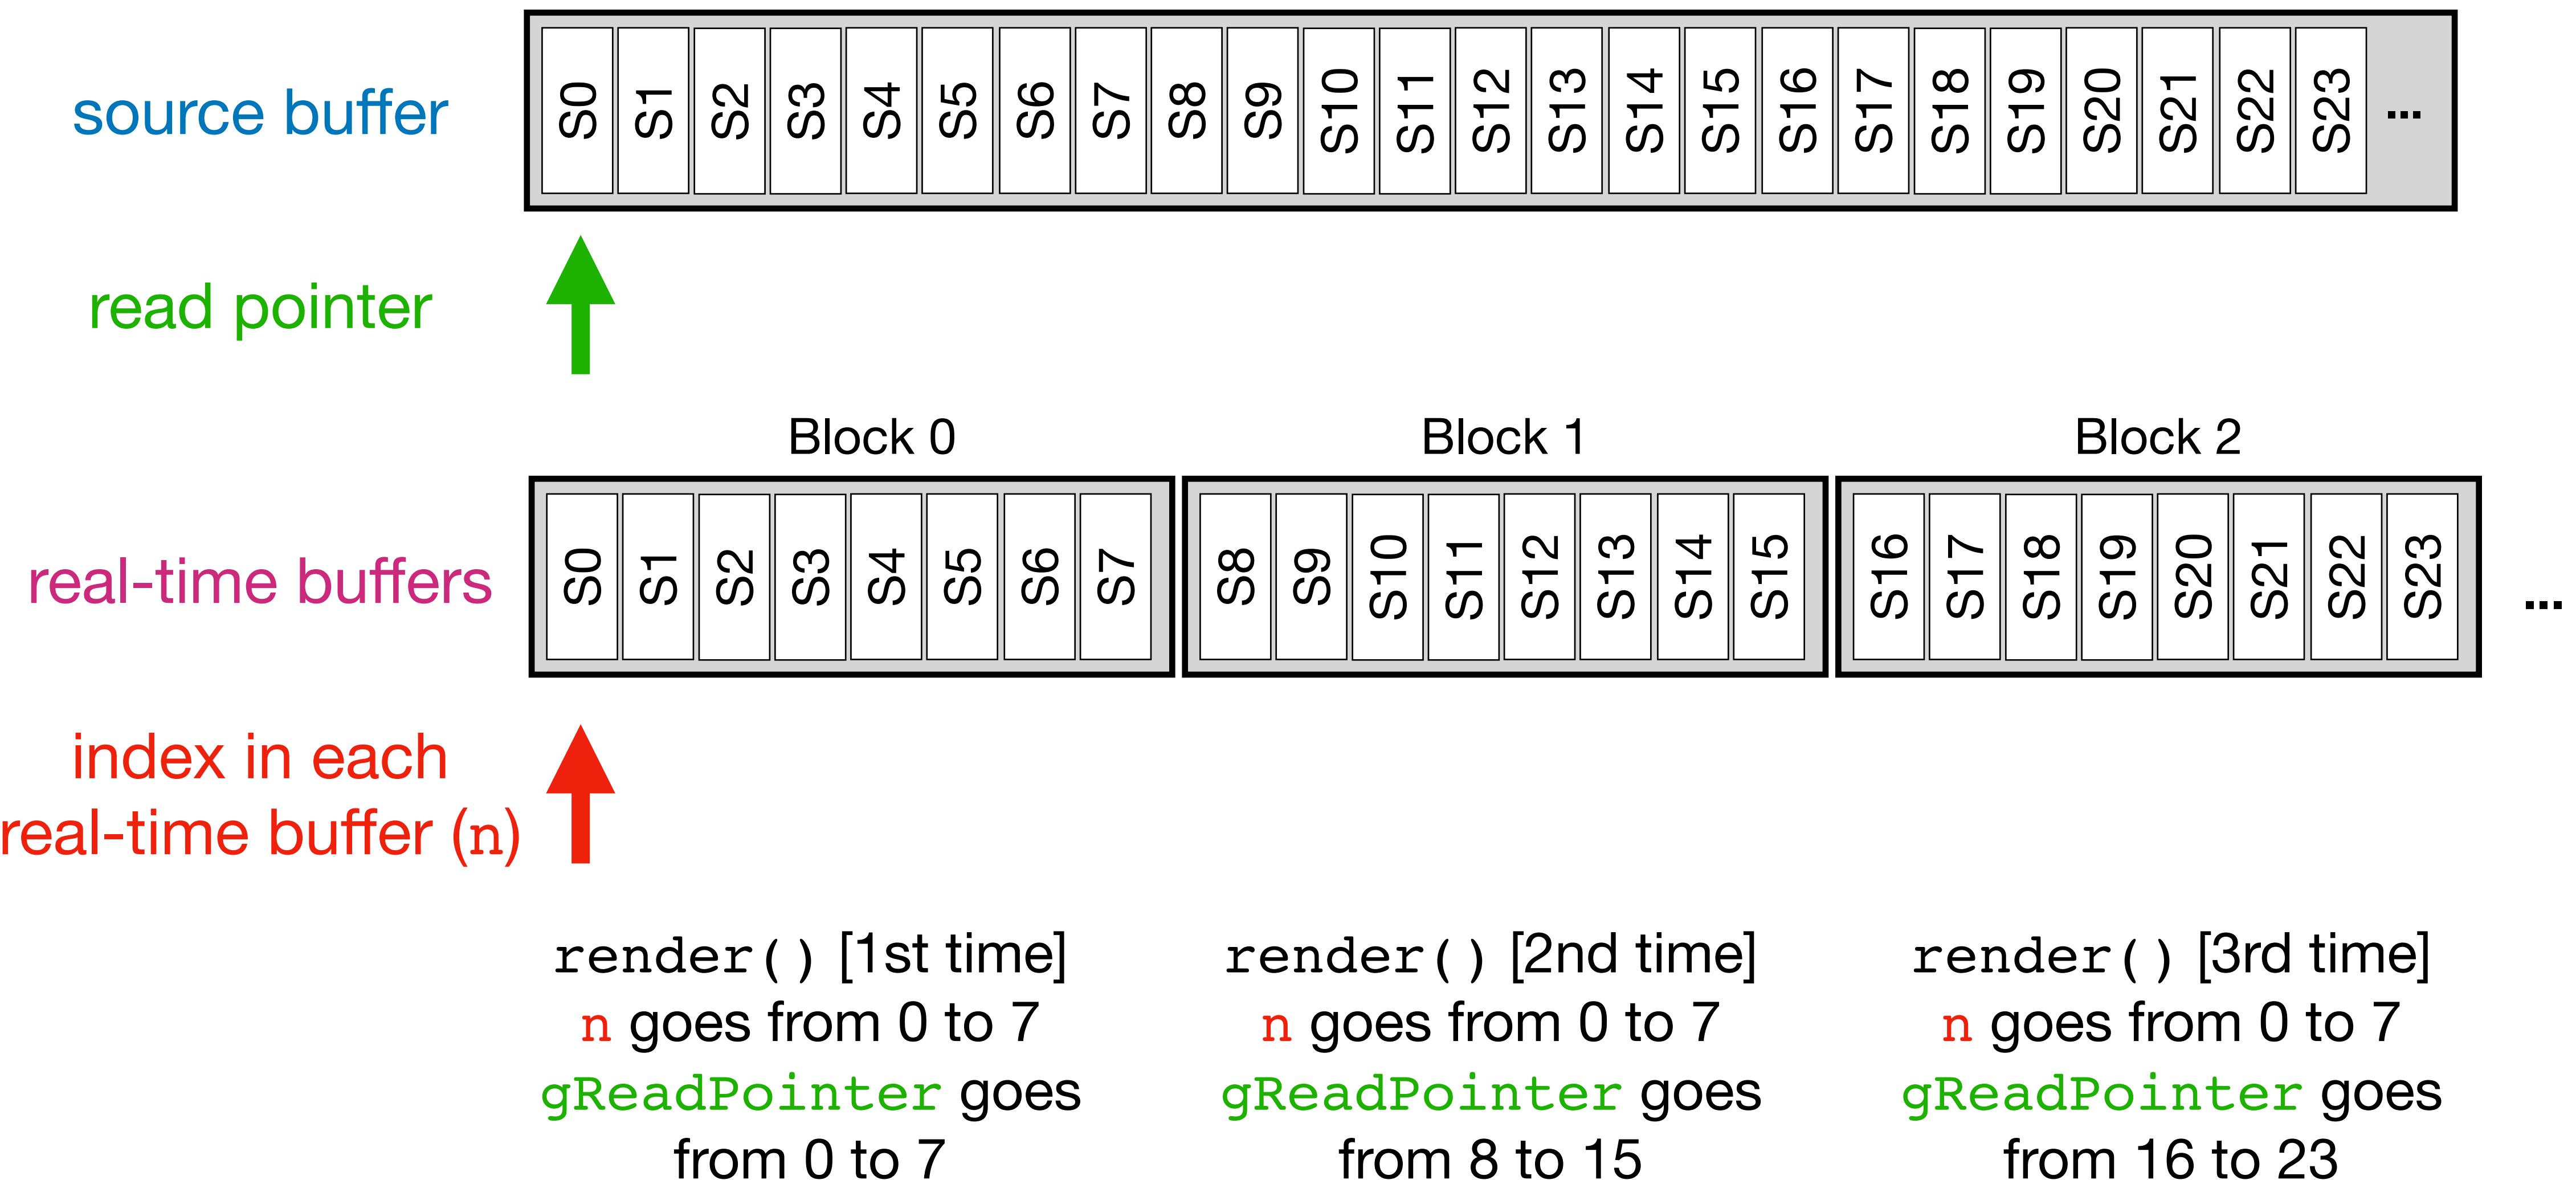
read pointer



# Fun with indexing

- One of the hardest parts of working with buffers can be keeping track of what each index means
- In this case, we've got two different kinds of buffers to think about:
  1. The recorded sound (let's call it the source buffer)
    - Only one buffer whose contents don't change
    - Length: number of samples in the source sound (possibly long)
  2. The buffer for each real-time audio block
    - A new buffer each time `render()` is called, accessed via `audiowrite()`
    - Length: block size of the real-time system (e.g. 16)
- Therefore, we need to keep track of two indexes:
  1. Where are we playing in the source buffer? (read pointer or play head)
  2. Where are we writing in the output buffer? (starts over from 0 each block)

# Fun with indexing



# Playing audio samples

## Global variables

keep track of source buffer, read pointer

```
std::vector<float> gSampleBuffer; // Buffer that holds the sound file
int gReadPointer = 0; // Position of the last frame we played
```

This runs once per sample within the block (default 16 times on Bela)

```
void render(BelaContext *context, void *userData)
{
    for(unsigned int n = 0; n < context->audioFrames; n++) {
        // Load a sample from the buffer which was loaded from the file
        float out = gSampleBuffer[gReadPointer];
        // Increment read pointer and reset to 0 when end of file is reached
        gReadPointer++;
        if(gReadPointer >= gSampleBuffer.size())
            gReadPointer = 0;
        for(unsigned int channel = 0; channel < context->audioOutChannels; channel++) {
            // Write the sample to every audio output channel
            audioWrite(context, n, channel, out);
        }
    }
}
```

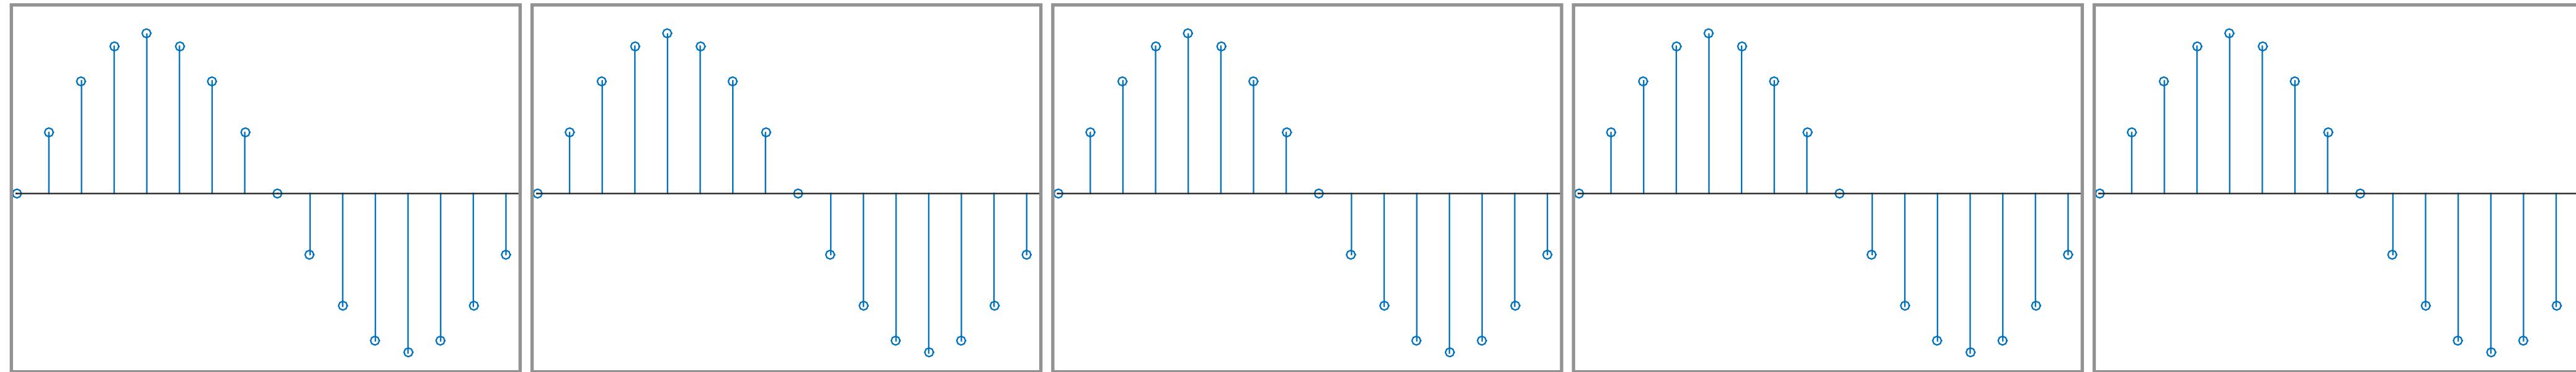
**gReadPointer** keeps track of where we are in the source buffer

**n** is used to specify which frame of the **output buffer** to write

update **gReadPointer**, whose value persists

# Uses of repeating buffers

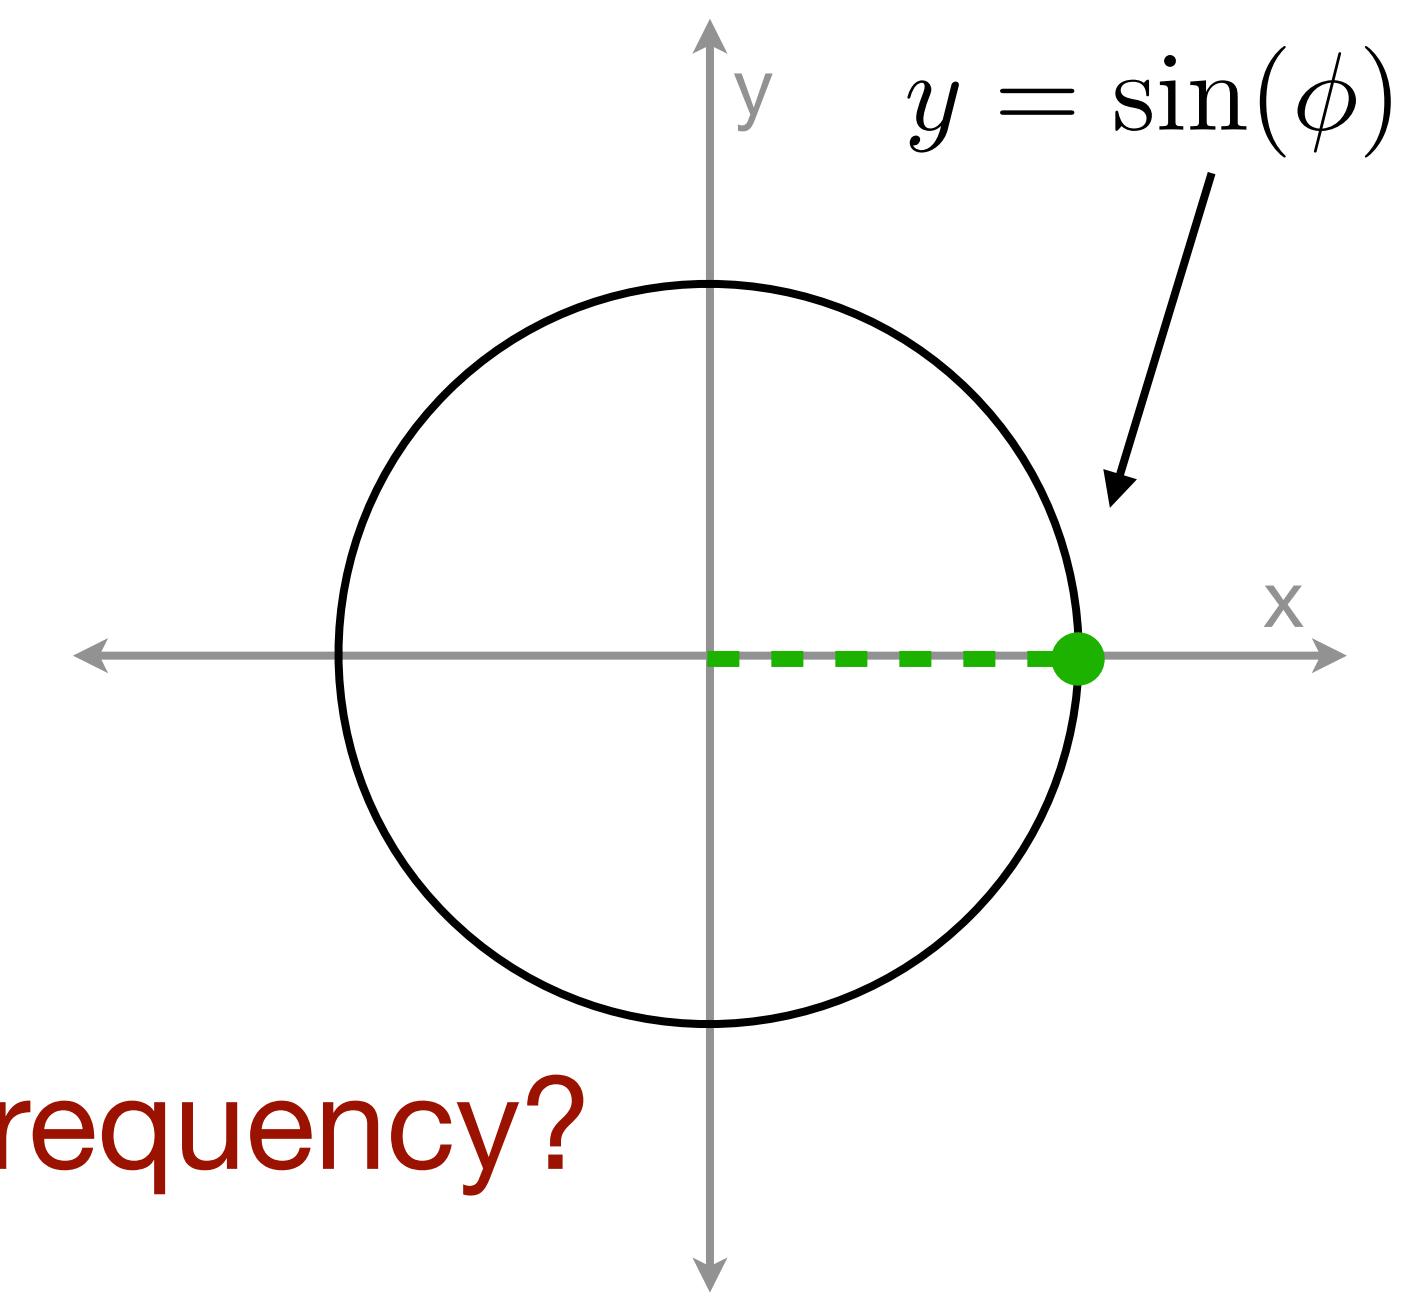
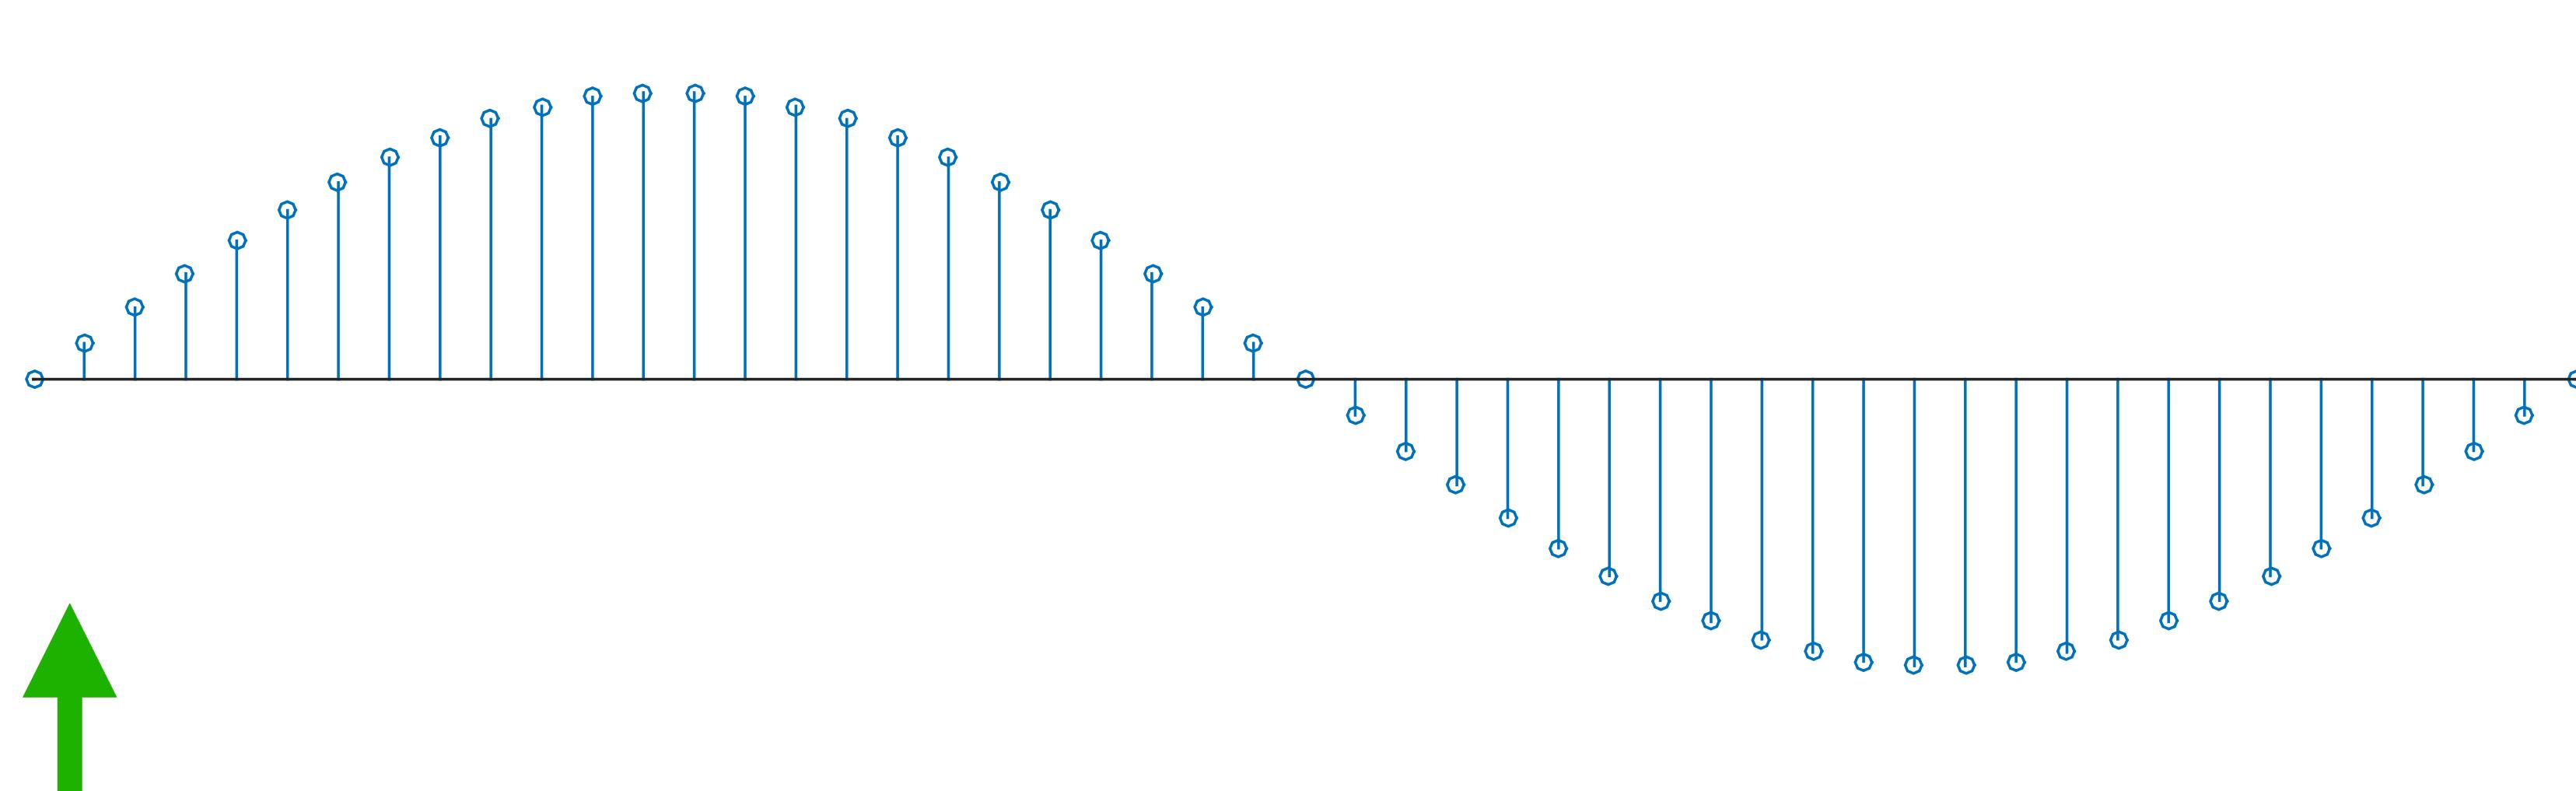
- There's another name for playing a buffer of samples over and over



- **Wavetable**
- This is another way of building an **oscillator**
  - Rather than calculating the output mathematically every sample, **play it from a buffer**
  - Another way of thinking about it: a **look-up table** of output values
- Advantages of wavetables over using functions like `sin()`:
  - **Less computation time** per sample, because we can calculate the table at the beginning
  - More **flexibility** on the waveform (which could be nearly any shape)

# Phase and frequency

- We need to keep track of where we are in the wavetable
  - i.e. keep track of the **read pointer** in the source buffer
  - We'll need to remember this between calls to `render()`, like before
- The read pointer is analogous to the **phase** in our previous oscillator:



- If the read pointer corresponds to phase, **what determines frequency?**
  - Frequency is the **derivative** (rate of change) of phase
  - Therefore, frequency is given by **rate of change** of the read pointer

# Wavetable task

- In our previous examples, the read pointer always moved at a rate of one sample per sample
  - It advanced one sample in the source buffer for each real-time output sample
- To implement a wavetable oscillator with adjustable frequency, we need to change how much `gReadPointer` increments each sample
  - Don't forget to `wrap` `gReadPointer` if it falls off the end of the buffer!
  - How do we calculate how far to move based on the frequency?
  - Hint: look at the calculation in the original `sin()` oscillator and think about what should replace  $2\pi$
- Task: using the `wavetable` example, implement a wavetable oscillator
  - Adjust the code so the frequency is correct at 220Hz

# Wavetable code

Previous approach:

$$x[n] = \sin(\phi) \quad \text{range } [0, 2\pi]$$
$$\phi = \phi + 2\pi f / f_s$$

New approach:

$$x[n] = \text{wavetable}[\phi] \quad \text{range } [0, l_{table}]$$
$$\phi = \phi + l_{table}f / f_s$$

```
int gReadPointer = 0; // Position of the last frame we played
```

```
void render(BelaContext *context, void *userData)
{
    for(unsigned int n = 0; n < context->audioFrames; n++) {
        // Load a sample from the buffer
        float out = gAmplitude * gWavetable[gReadPointer];
```

What is a possible problem here?

```
// Increment read pointer and wrap around when end of table is reached
gReadPointer += gWavetableLength * qFrequency / context->audioSampleRate;
while(gReadPointer >= gWavetableLength)
    gReadPointer -= gWavetableLength;
```

```
for(unsigned int channel = 0; channel < context->audioOutChannels; channel++) {
    // Write the sample to every audio output channel
    audioWrite(context, n, channel, out);
}
```

}

Minor point:  
why use `while()` and not `if()`?

# Fractional read pointer

- To play the wavetable at a particular frequency, we increment the read pointer by:  $\phi = \phi + l_{table}f/f_s$
- The increment might not be an integer
  - For  $l_{table} = 512, f = 220\text{Hz}, f_s = 44100\text{Hz}$ : increment is 2.55 samples
  - We need a float to hold the read pointer
  - If we use an int, the increment will be rounded down to 2 samples
    - Notice: float to int always rounds down (dropping the decimals)
    - What is the frequency in this case?
    - If  $\Delta\phi = 2$  is the increment, then  $f = \Delta\phi \cdot f_s/l_{table} = 172.2\text{Hz}$
- We can't use a float to index an array
  - Doesn't make sense in code: `gSampleBuffer[1.5]` has no meaning
  - Doesn't make sense mathematically:  $x[1.5]$  (fractional index in discrete time) is undefined

# Fractional read pointer

- What should we do instead?
- We can keep the fraction all the way up to the point we read the array
  - At that point, round down and convert to an int
  - Or alternatively, round to the nearest int
  - But only do this to access the array; don't change gReadPointer itself
- Try this now!
  - Hint: `floorf( )` rounds a float downward

# Fractional read pointer

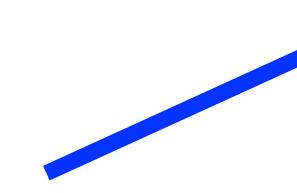
```
float gReadPointer = 0;           // Position of the last frame we played

void render(BelaContext *context, void *userData)
{
    for(unsigned int n = 0; n < context->audioFrames; n++) {
        // Load a sample from the buffer
        int indexBelow = floorf(gReadPointer);
        float out = gAmplitude * gSampleBuffer[indexBelow];

        // Increment read pointer and wrap around when end of table is reached
        gReadPointer += gSampleBufferLength * gFrequency / context->audioSampleRate;
        while(gReadPointer >= gSampleBufferLength)
            gReadPointer -= gSampleBufferLength;

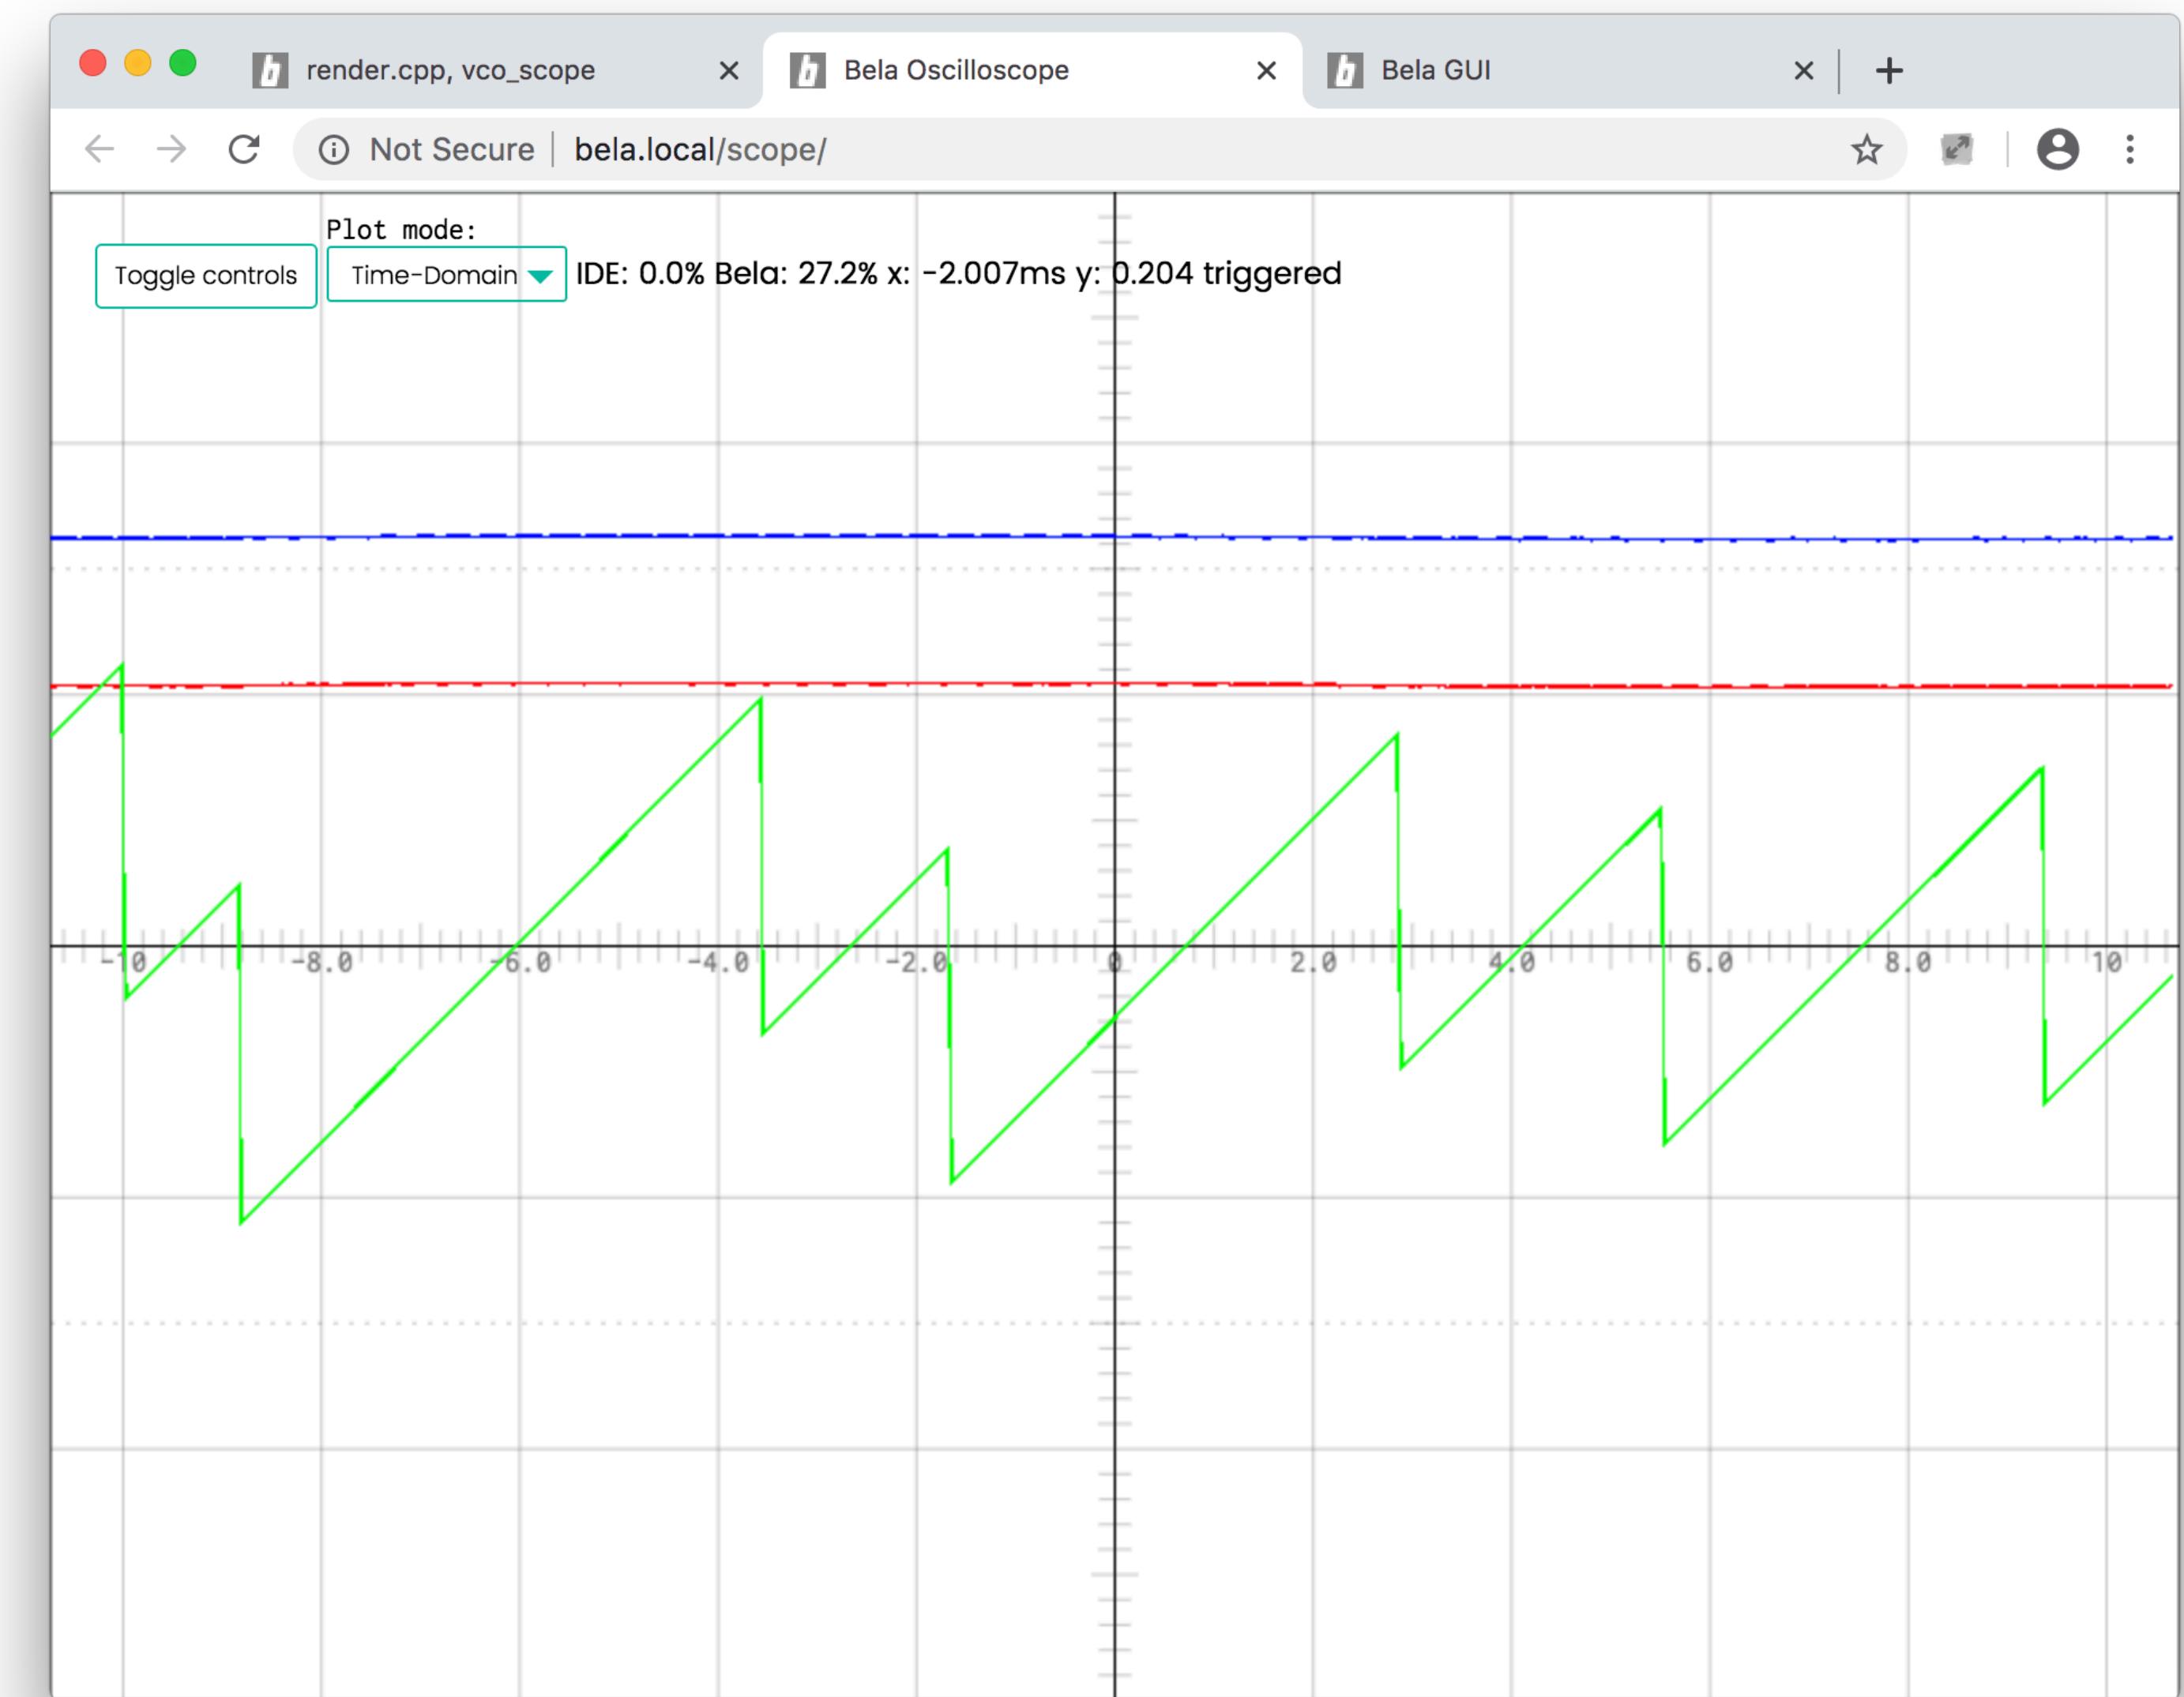
        for(unsigned int channel = 0; channel < context->audioOutChannels; channel++) {
            // Write the sample to every audio output channel
            audioWrite(context, n, channel, out);
        }
    }
}
```

We still need to convert the result to an int



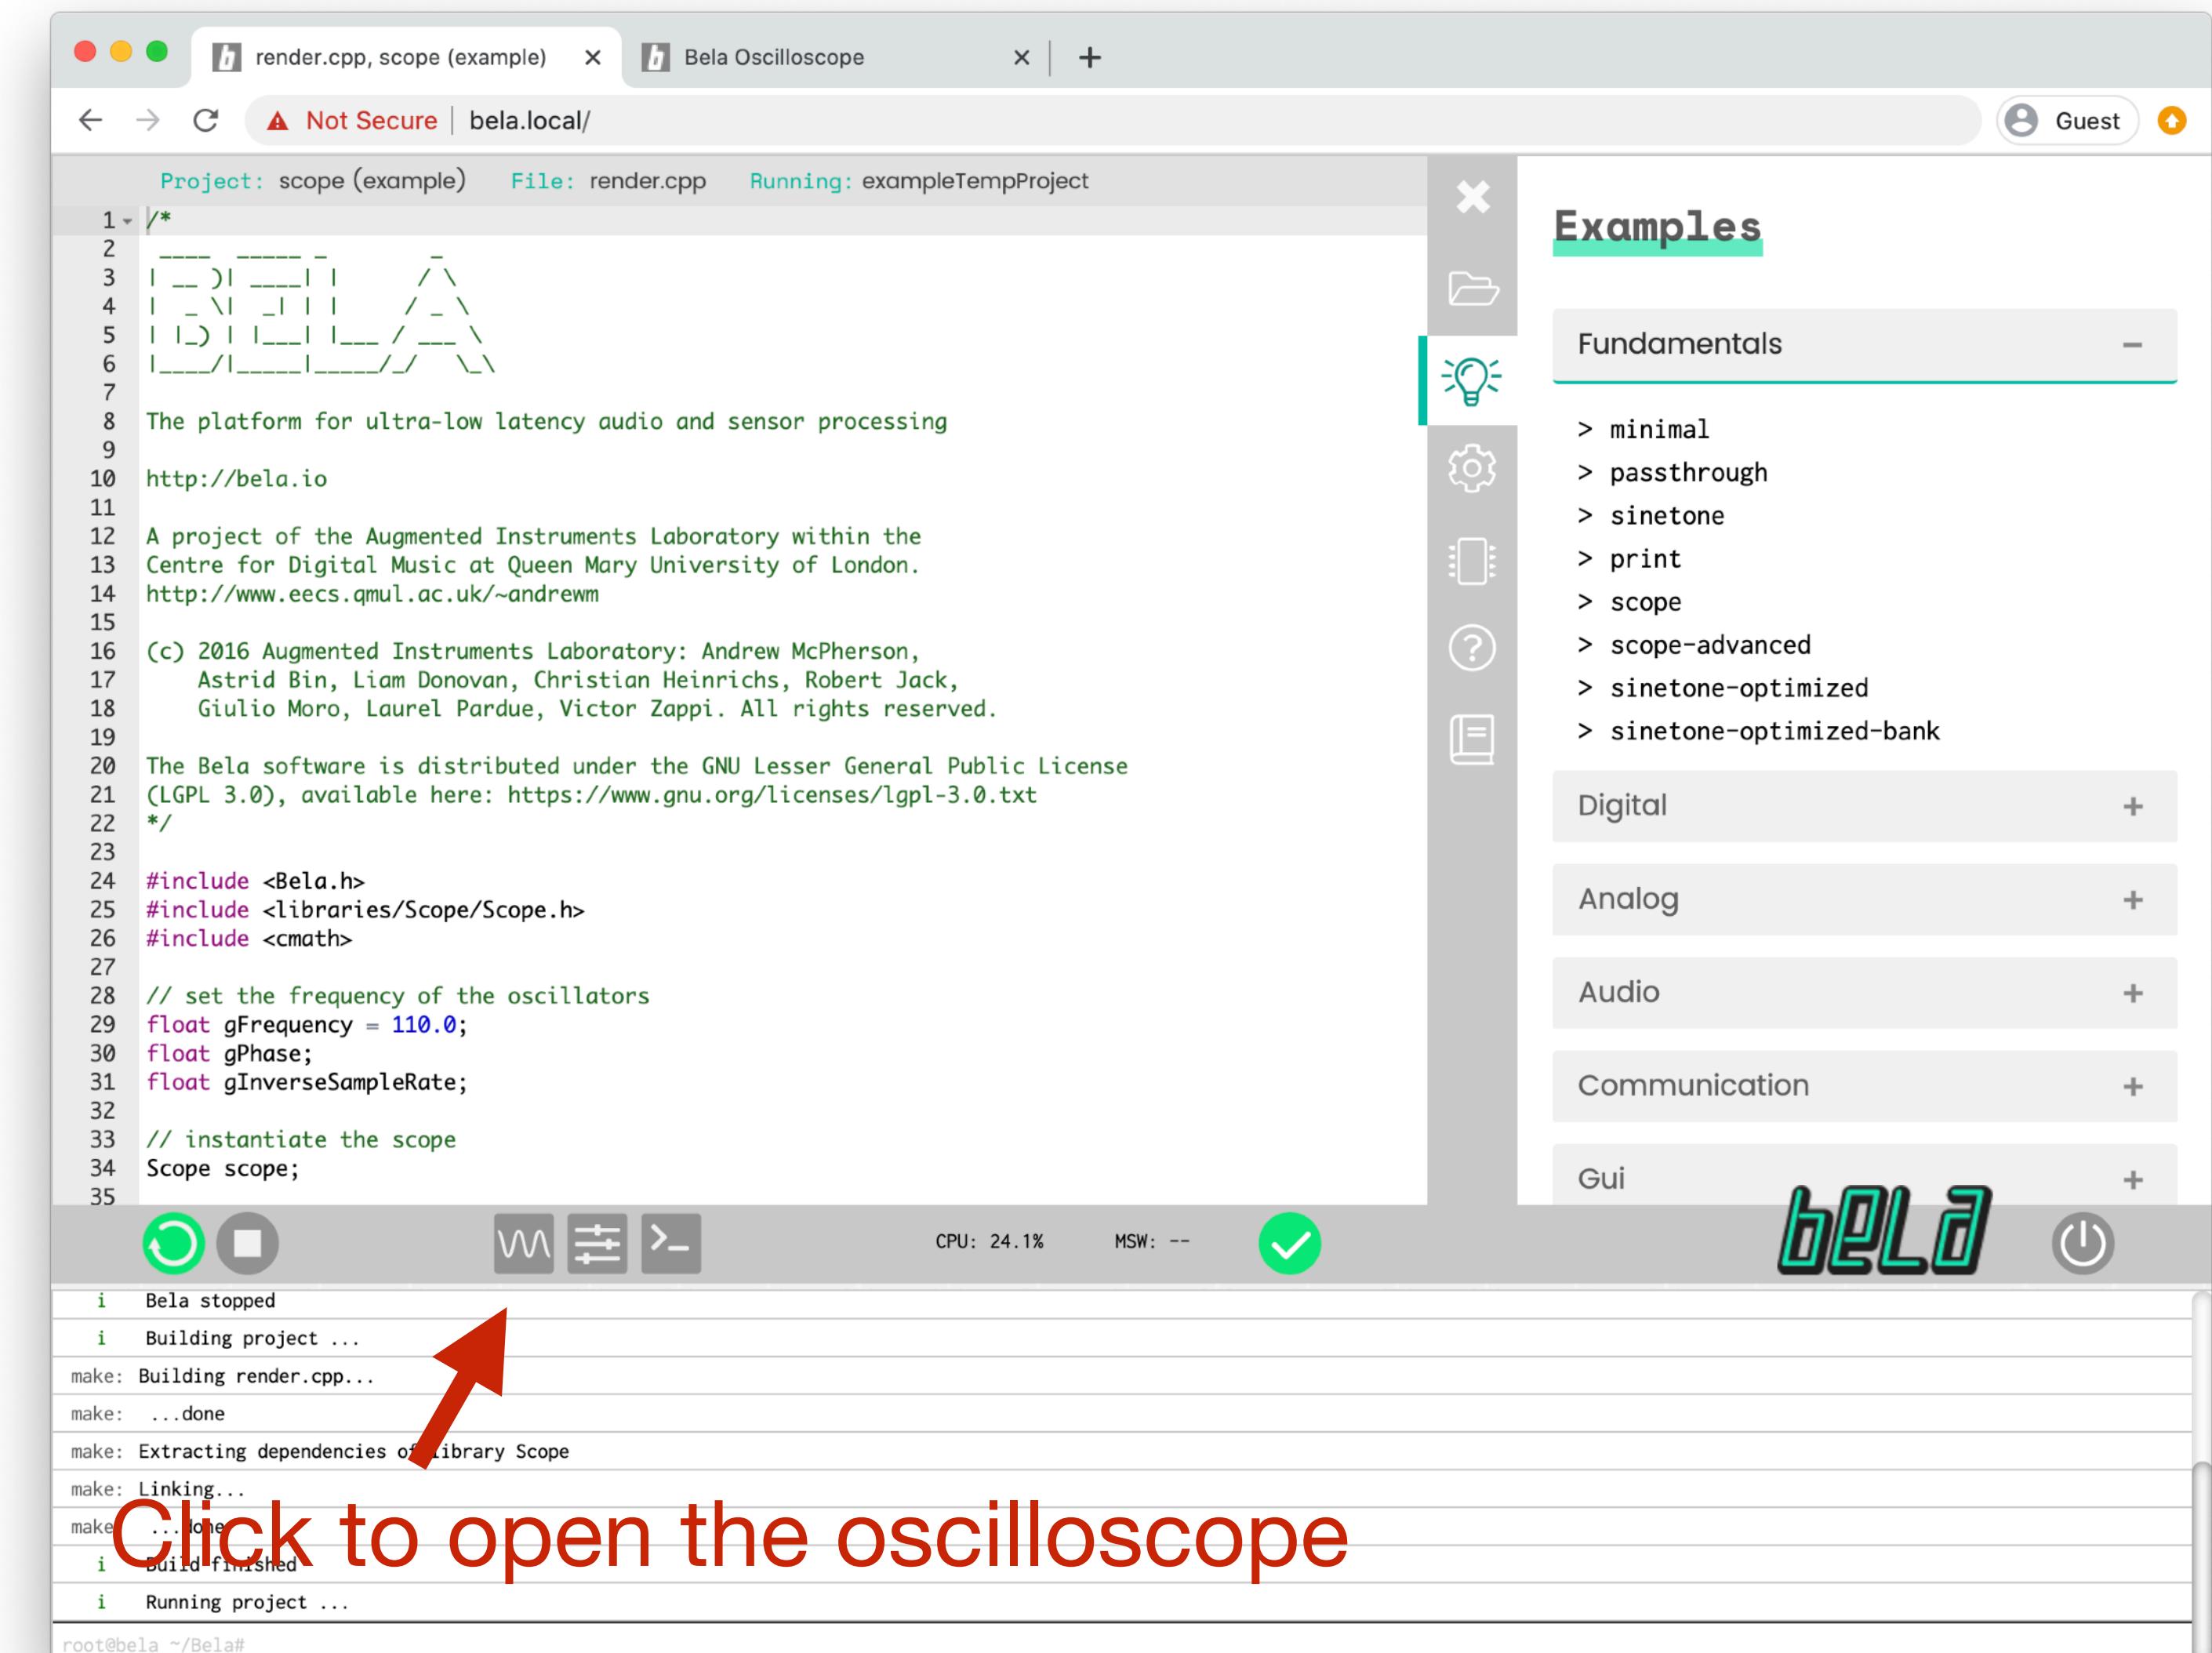
# Oscilloscope

- Bela features an oscilloscope that runs in the browser
  - ▶ As the name “oscillo-scope” implies, a way of viewing oscillations (signals)
- Typical use of oscilloscope is an X-Y plot
  - ▶ Time on the X-axis
  - ▶ Amplitude on the Y-axis
- The scope can plot any signal made or used by your program
  - ▶ Not just I/O signals!



# Launching the oscilloscope

- In the Examples tab of the IDE, open the project **Fundamentals/scope**
- Click the oscilloscope button to launch the scope in a new browser tab
  - ▶ You'll see the oscillator output in the scope



The screenshot shows the Bela IDE interface. The main window displays the code for the 'scope' example. The code includes comments about the platform being used for ultra-low latency audio and sensor processing, credits to the Augmented Instruments Laboratory at Queen Mary University of London, and the GNU Lesser General Public License (LGPL 3.0). The code also includes includes for Bela.h, libraries/Scope/Scope.h, and cmath, and defines variables for frequency, phase, and inverse sample rate, along with instantiating a Scope object.

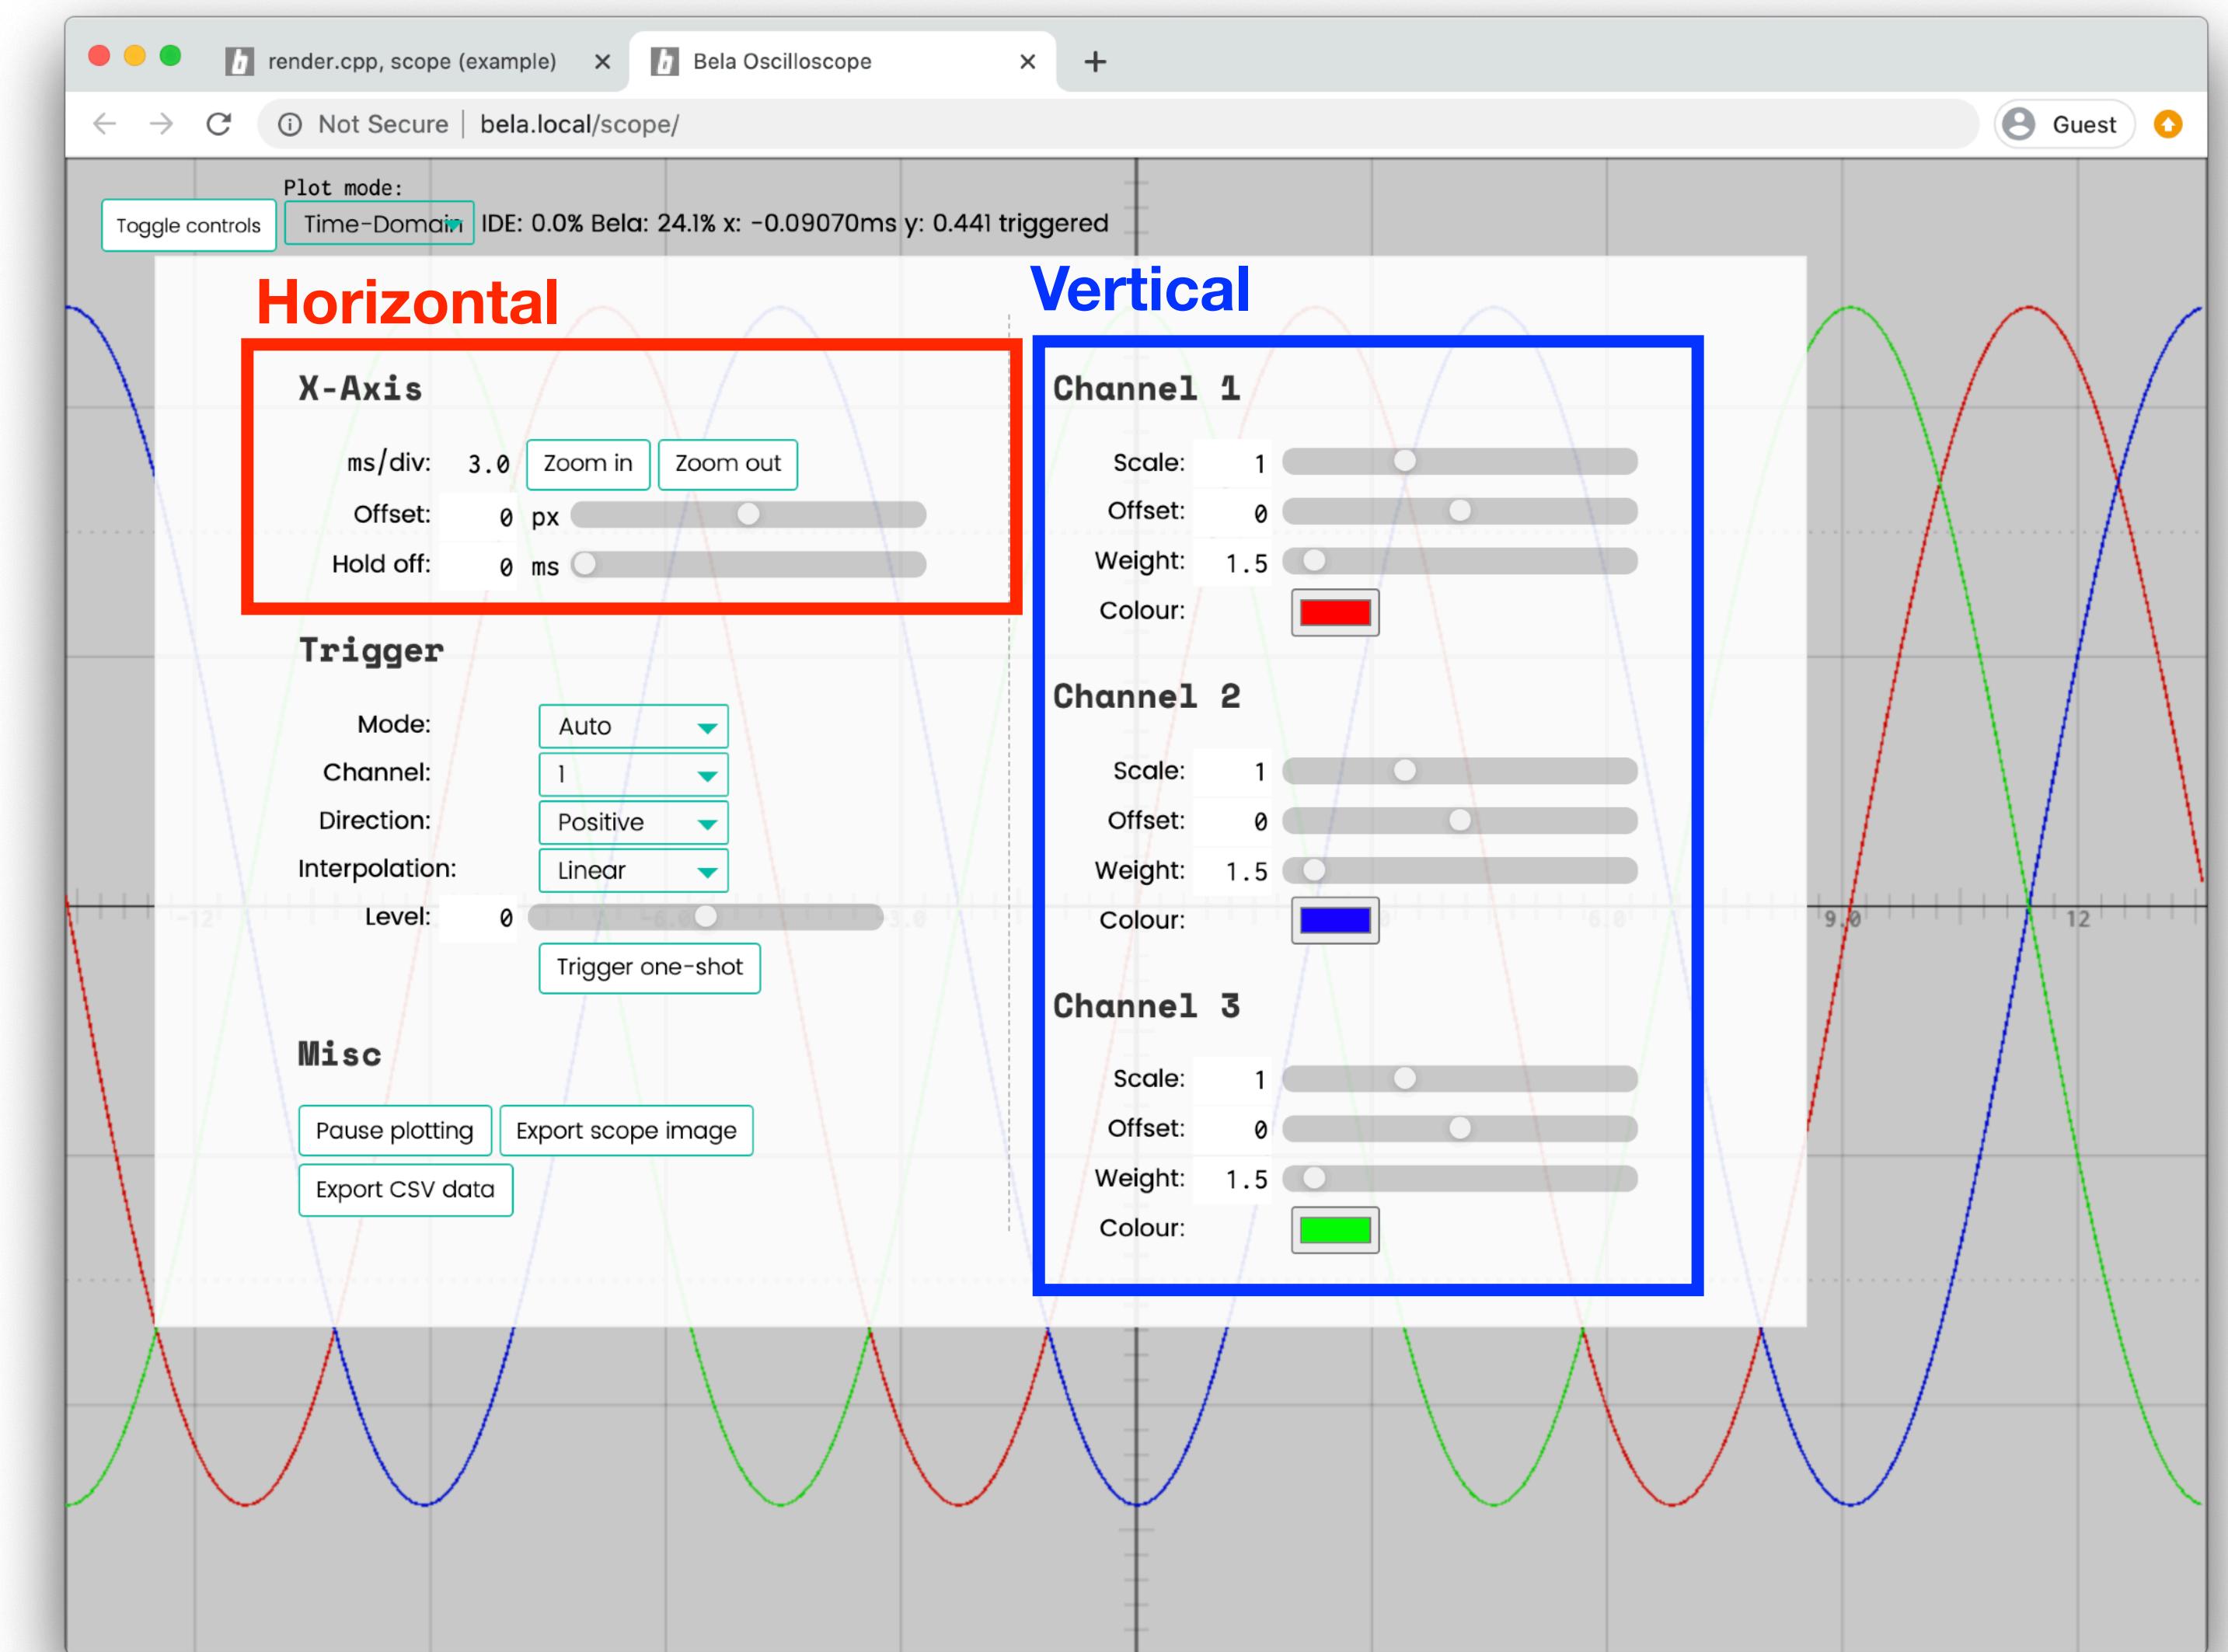
The right side of the interface features a sidebar titled 'Examples' with a tree view. Under the 'Fundamentals' category, several examples are listed: minimal, passthrough, sinetone, print, scope, scope-advanced, sinetone-optimized, and sinetone-optimized-bank. Below this, categories for Digital, Analog, Audio, Communication, and Gui are shown, each with a '+' sign to expand.

At the bottom of the interface, there is a toolbar with various icons, including a power button icon. A red arrow points to the 'scope' button in the toolbar.

A red annotation with the text "Click to open the oscilloscope" is overlaid on the bottom left of the screenshot, pointing towards the 'scope' button in the toolbar.

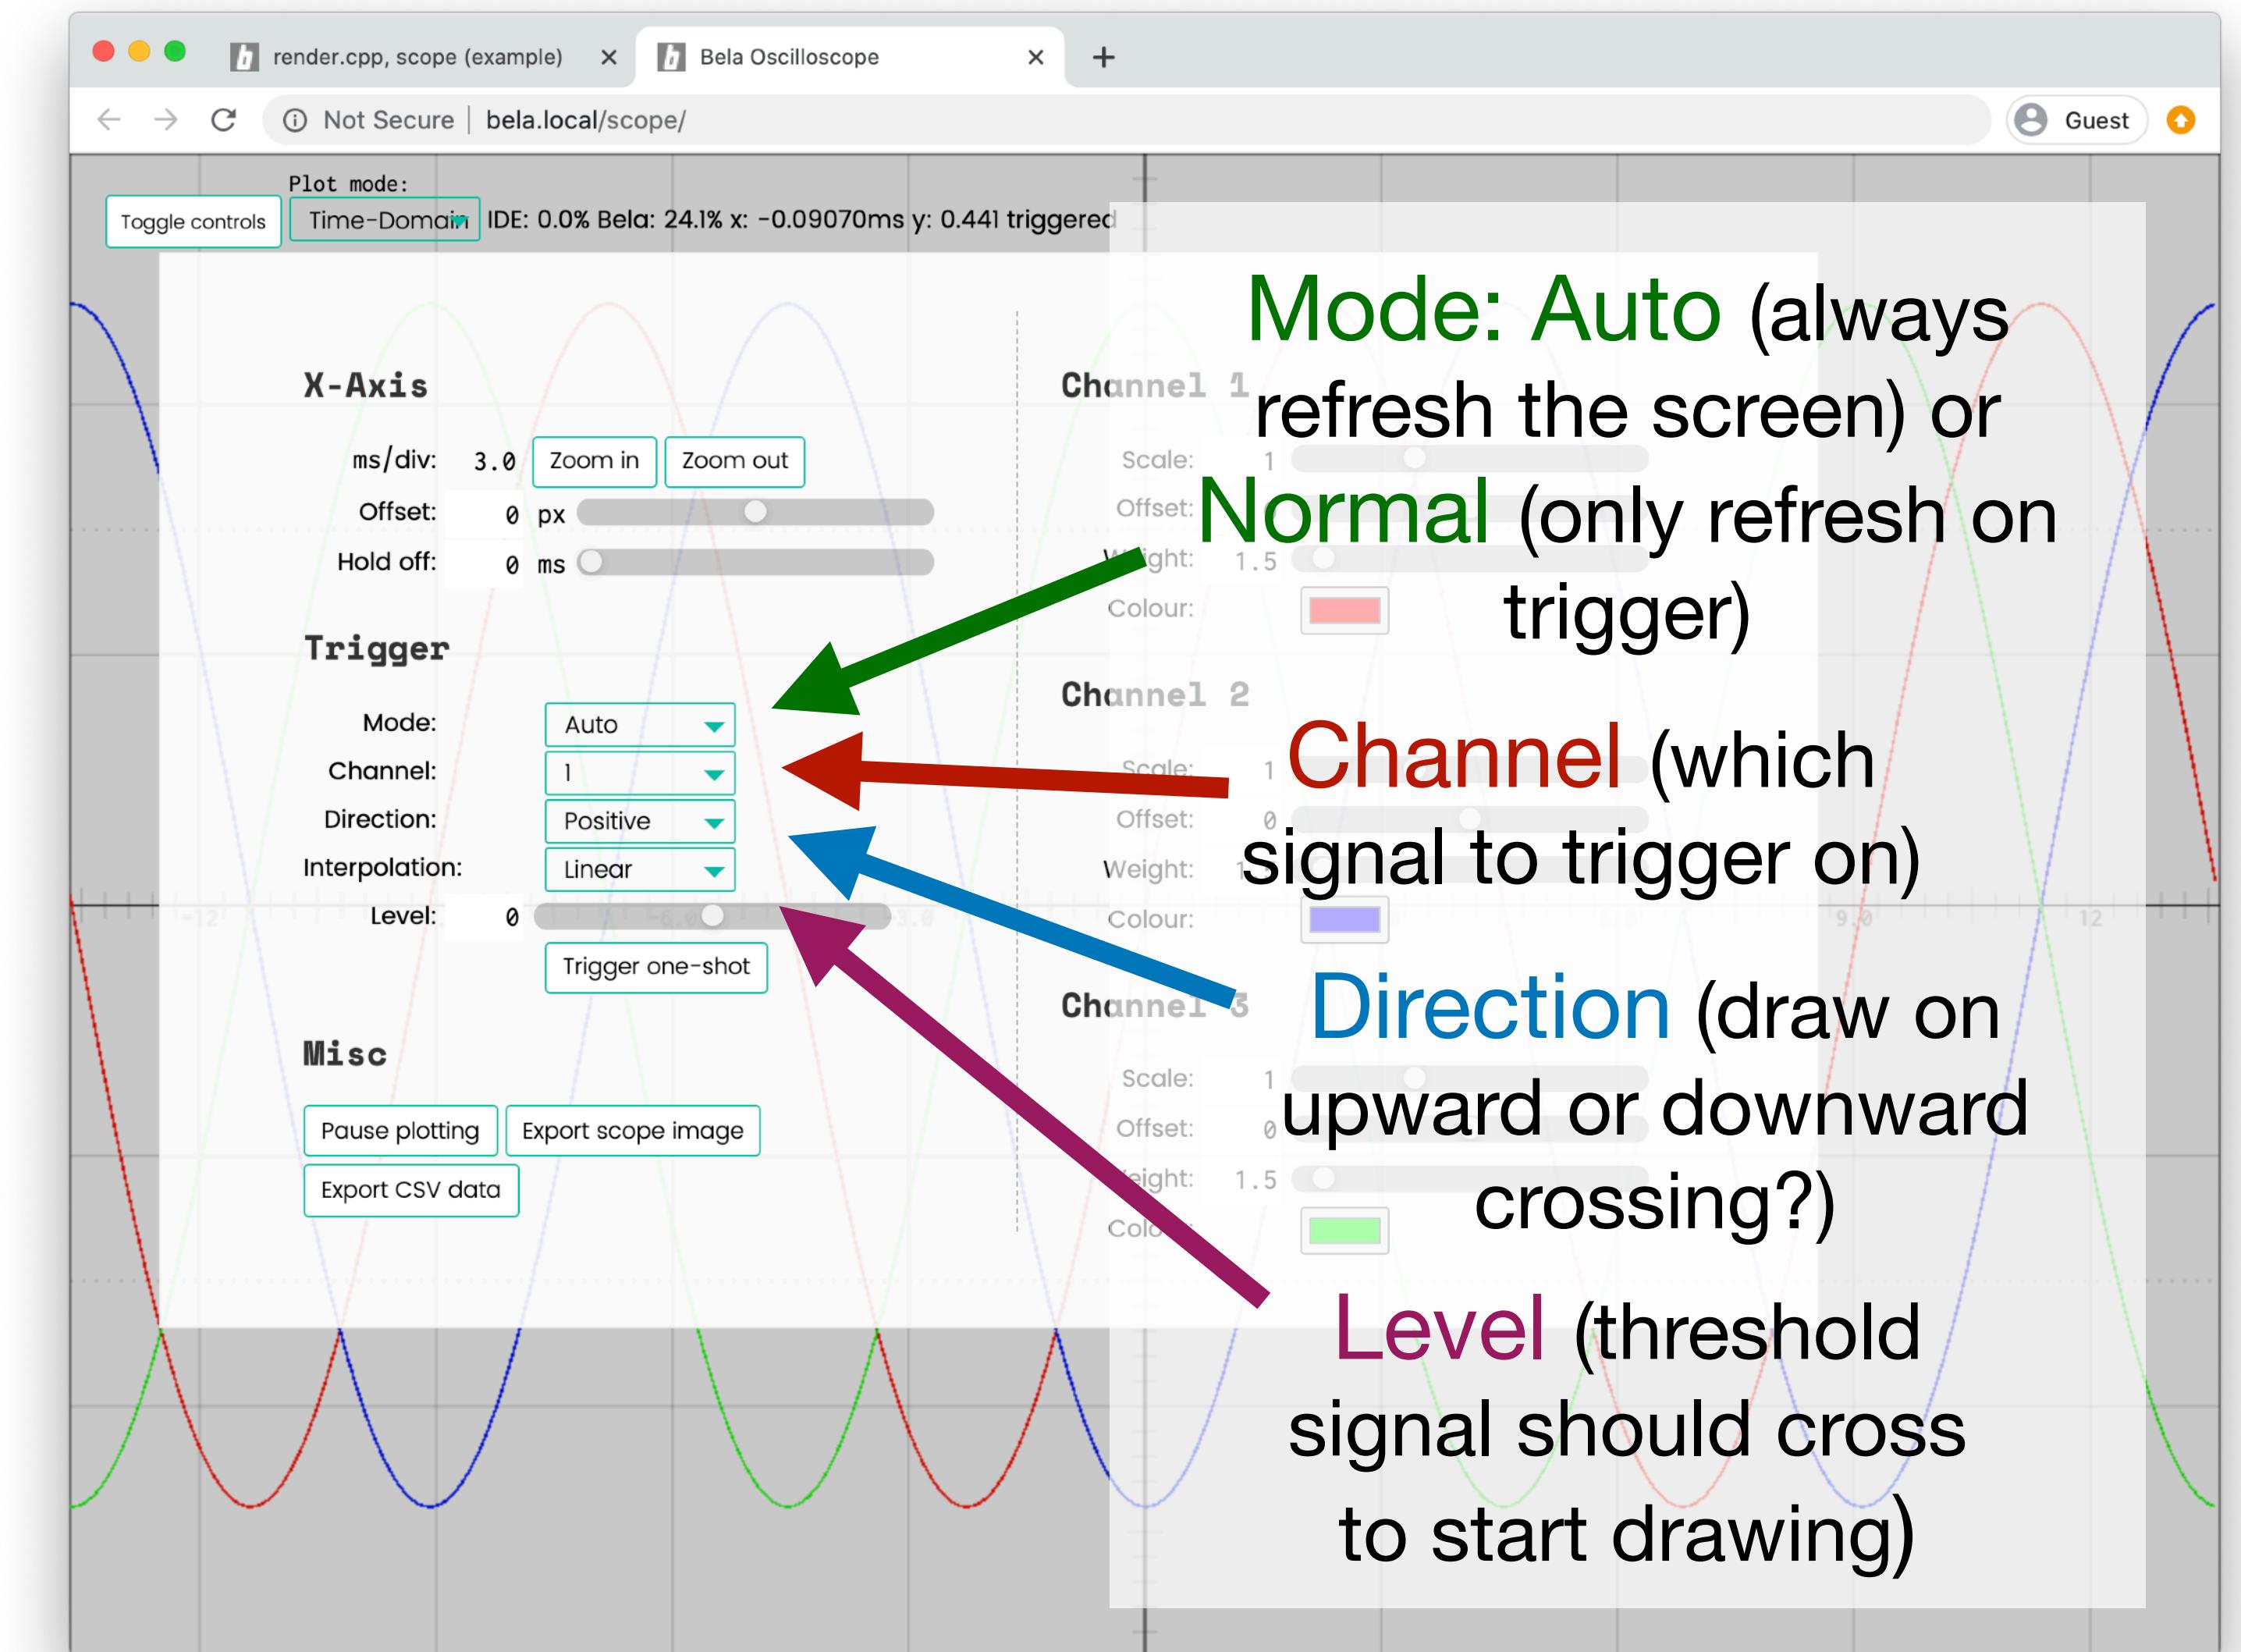
# Oscilloscope controls: scaling

- The most important controls on any scope are the **scales**:
  - Horizontal (time) scale
  - Vertical (signal / voltage) scales
- On Bela, we can adjust:
  - Separate vertical scales for each channel (i.e. each signal)
  - Only one horizontal scale for all channels (**why?**)
- Remember default ranges:
  - Audio I/O is -1 to 1 full scale
  - Analog I/O is 0 to 1
  - Your own signals might have other ranges!



# Oscilloscope controls: trigger

- Oscilloscopes show the signal in brief snapshots
  - One screenful corresponds to a limited window of time
  - To show the signal in real time, the scope **regularly refreshes the screen** with new time windows
- To keep the picture still, what do we need to do?
  - Start drawing at the **same phase** in the waveform each time
  - This is what the trigger does: **sets where and how the screen starts drawing**



# Oscilloscope code

- To use the scope, you need to include it in your code and tell it which signals to display
- Steps:
  1. Include Scope.h
  2. Make a global variable of type Scope
  3. Set the number of channels and the sample rate in setup()
  4. Send the signals in render()

```
#include <Bela.h>
#include <libraries/Scope/Scope.h>

// Oscilloscope
Scope gScope;

bool setup(BelaContext *context, void *userData)
{
    // Set up the browser-based oscilloscope
    gScope.setup(3, context->audioSampleRate);

    return true;
}

void render(BelaContext *context, void *userData)
{
    for(unsigned int n = 0; n < context->audioFrames; n++) {
        // [calculate signal0, signal1, signal2 here]
        // Log signals to the oscilloscope
        gScope.log(signal0, signal1, signal2);
    }
}
```

# Oscilloscope code

- Use the method `Scope::log()` to send signals to the scope
- Two forms of `log()`:
  - Variable arguments (each channel in its own argument)
  - Or pass an array of `float`
- It's important that you call `log()` at the right rate
  - Call it once per frame at the sample rate you specified at initialisation
  - Need to send the values for all signals at the same time

```
#include <Bela.h>
#include <libraries/Scope/Scope.h>

// Oscilloscope
Scope gScope;

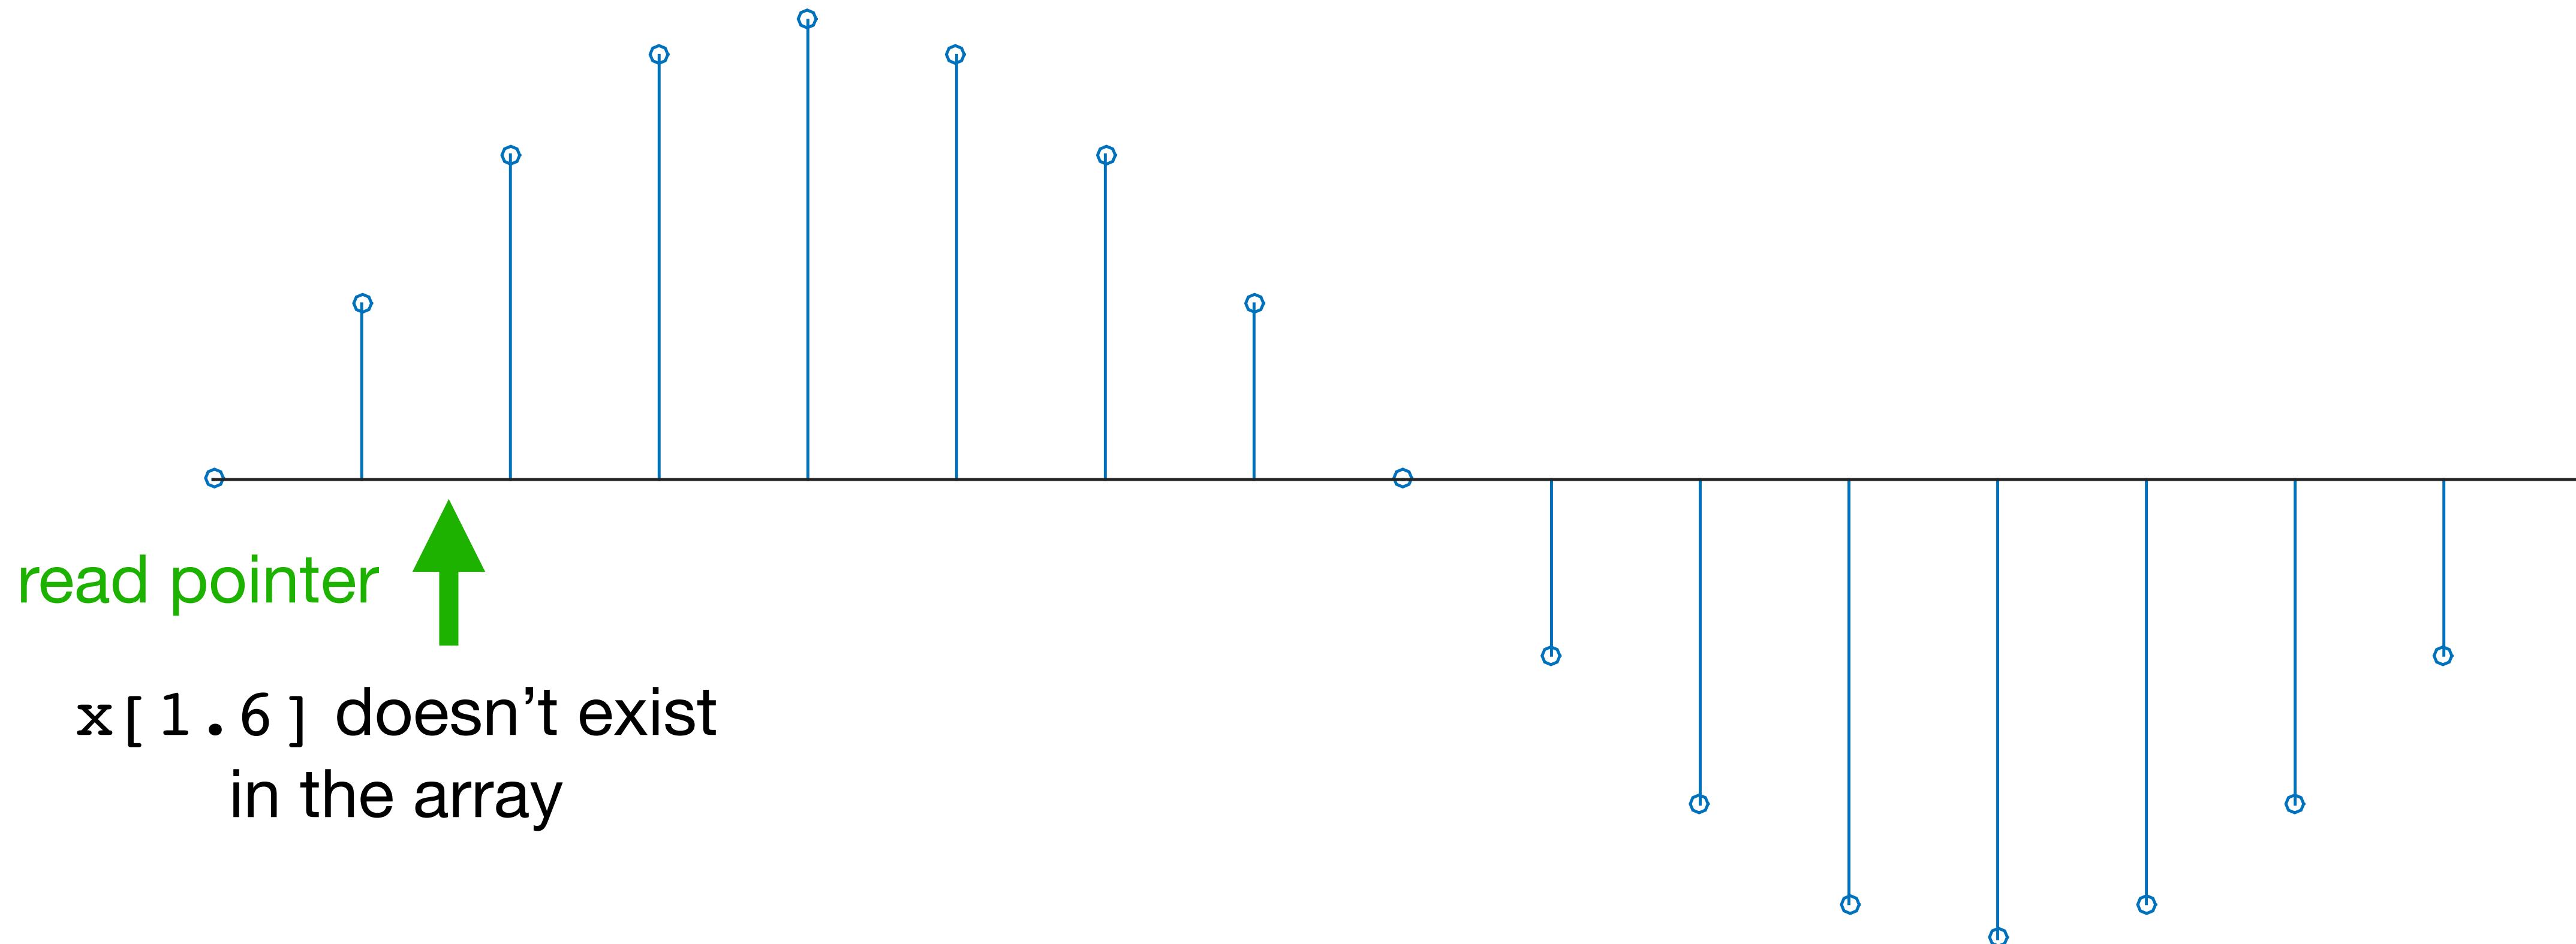
bool setup(BelaContext *context, void *userData)
{
    // Set up the browser-based oscilloscope
    gScope.setup(3, context->audioSampleRate);

    return true;
}

void render(BelaContext *context, void *userData)
{
    for(unsigned int n = 0; n < context->audioFrames; n++) {
        // [calculate signal0, signal1, signal2 here]
        // Log signals to the oscilloscope
        gScope.log(signal0, signal1, signal2);
    }
}
```

# A better fractional index?

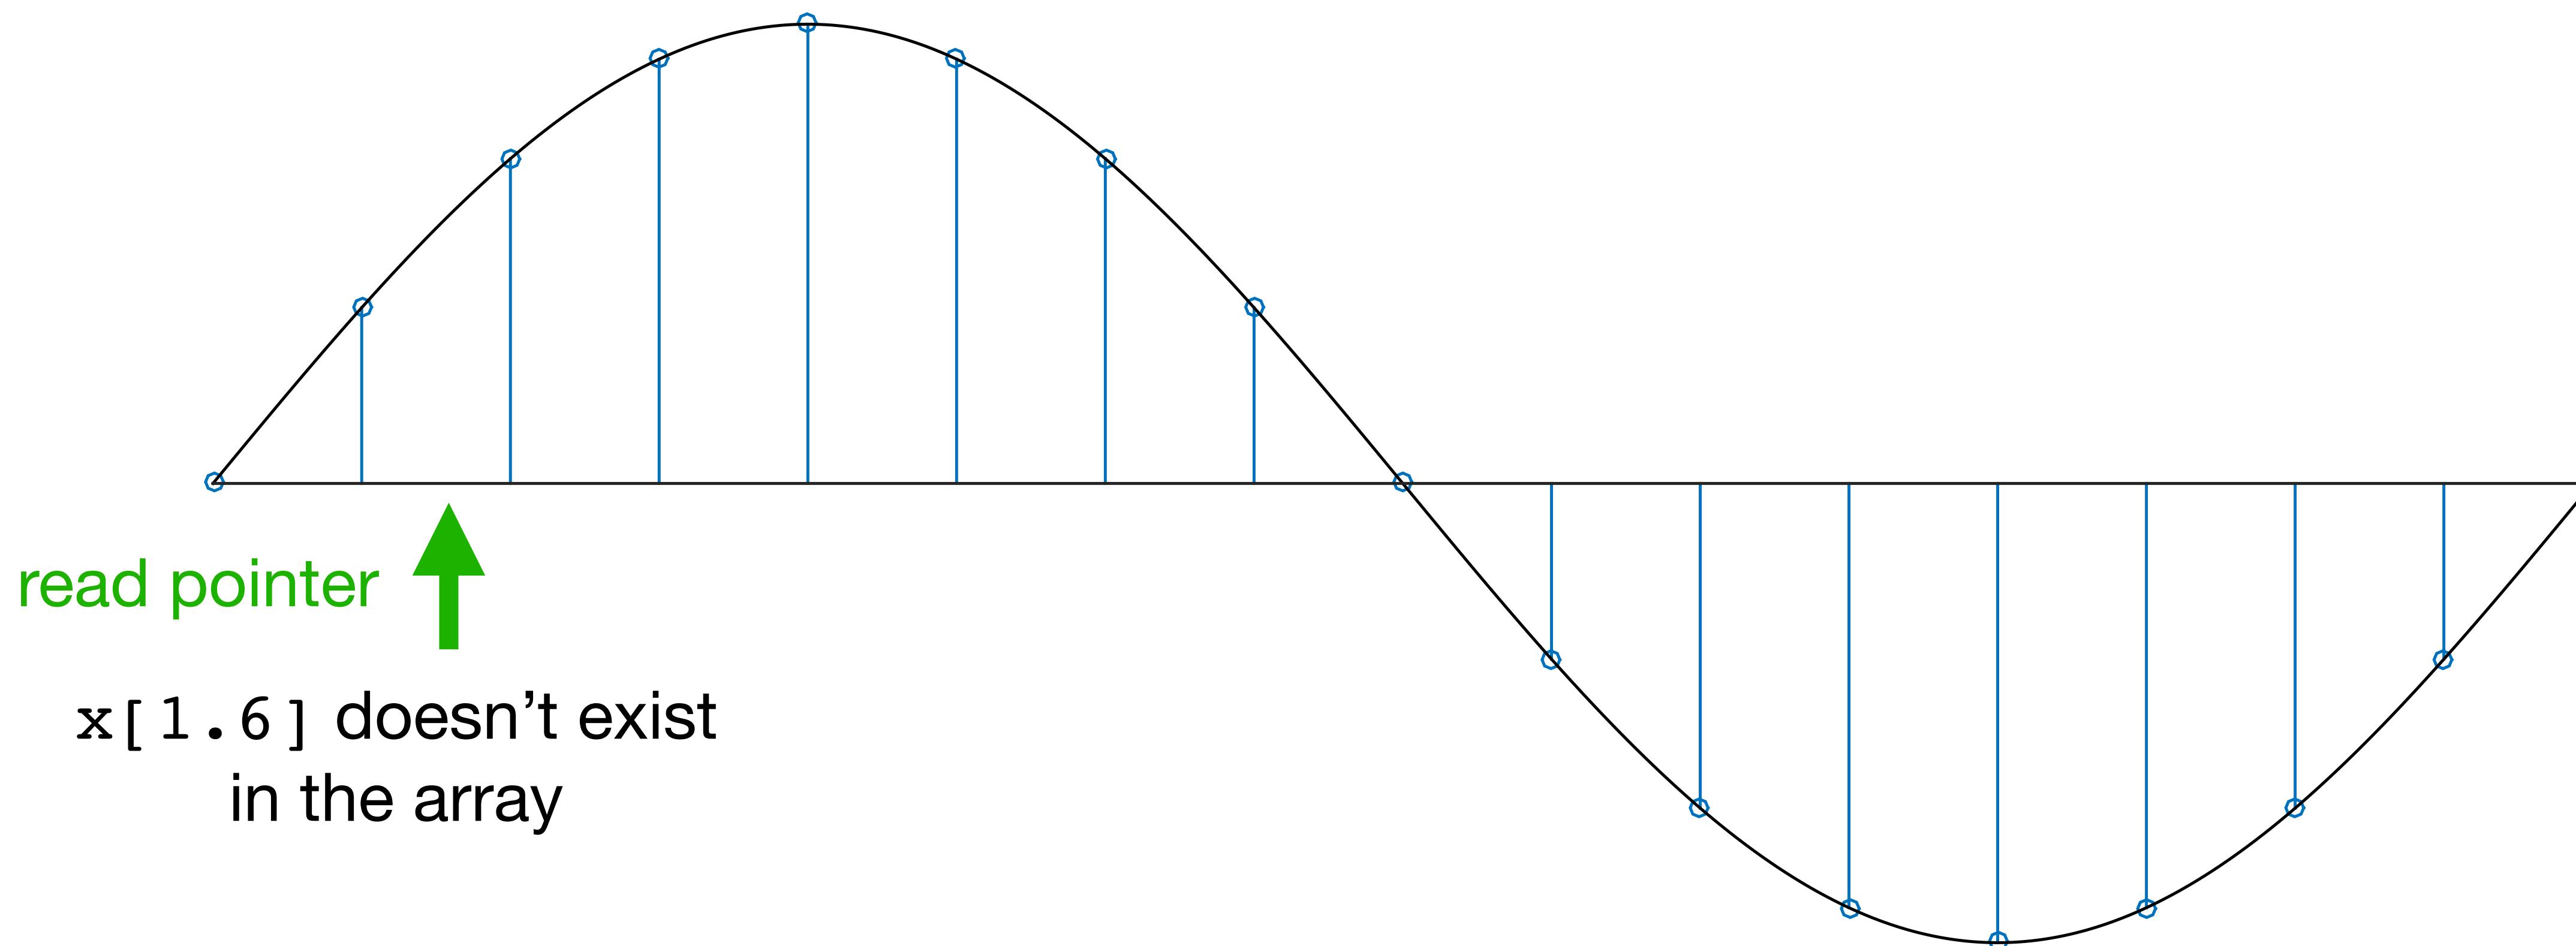
- We want to read fractional array values that, strictly speaking, don't exist...



# A better fractional index?

- We want to read fractional array values that, strictly speaking, don't exist...

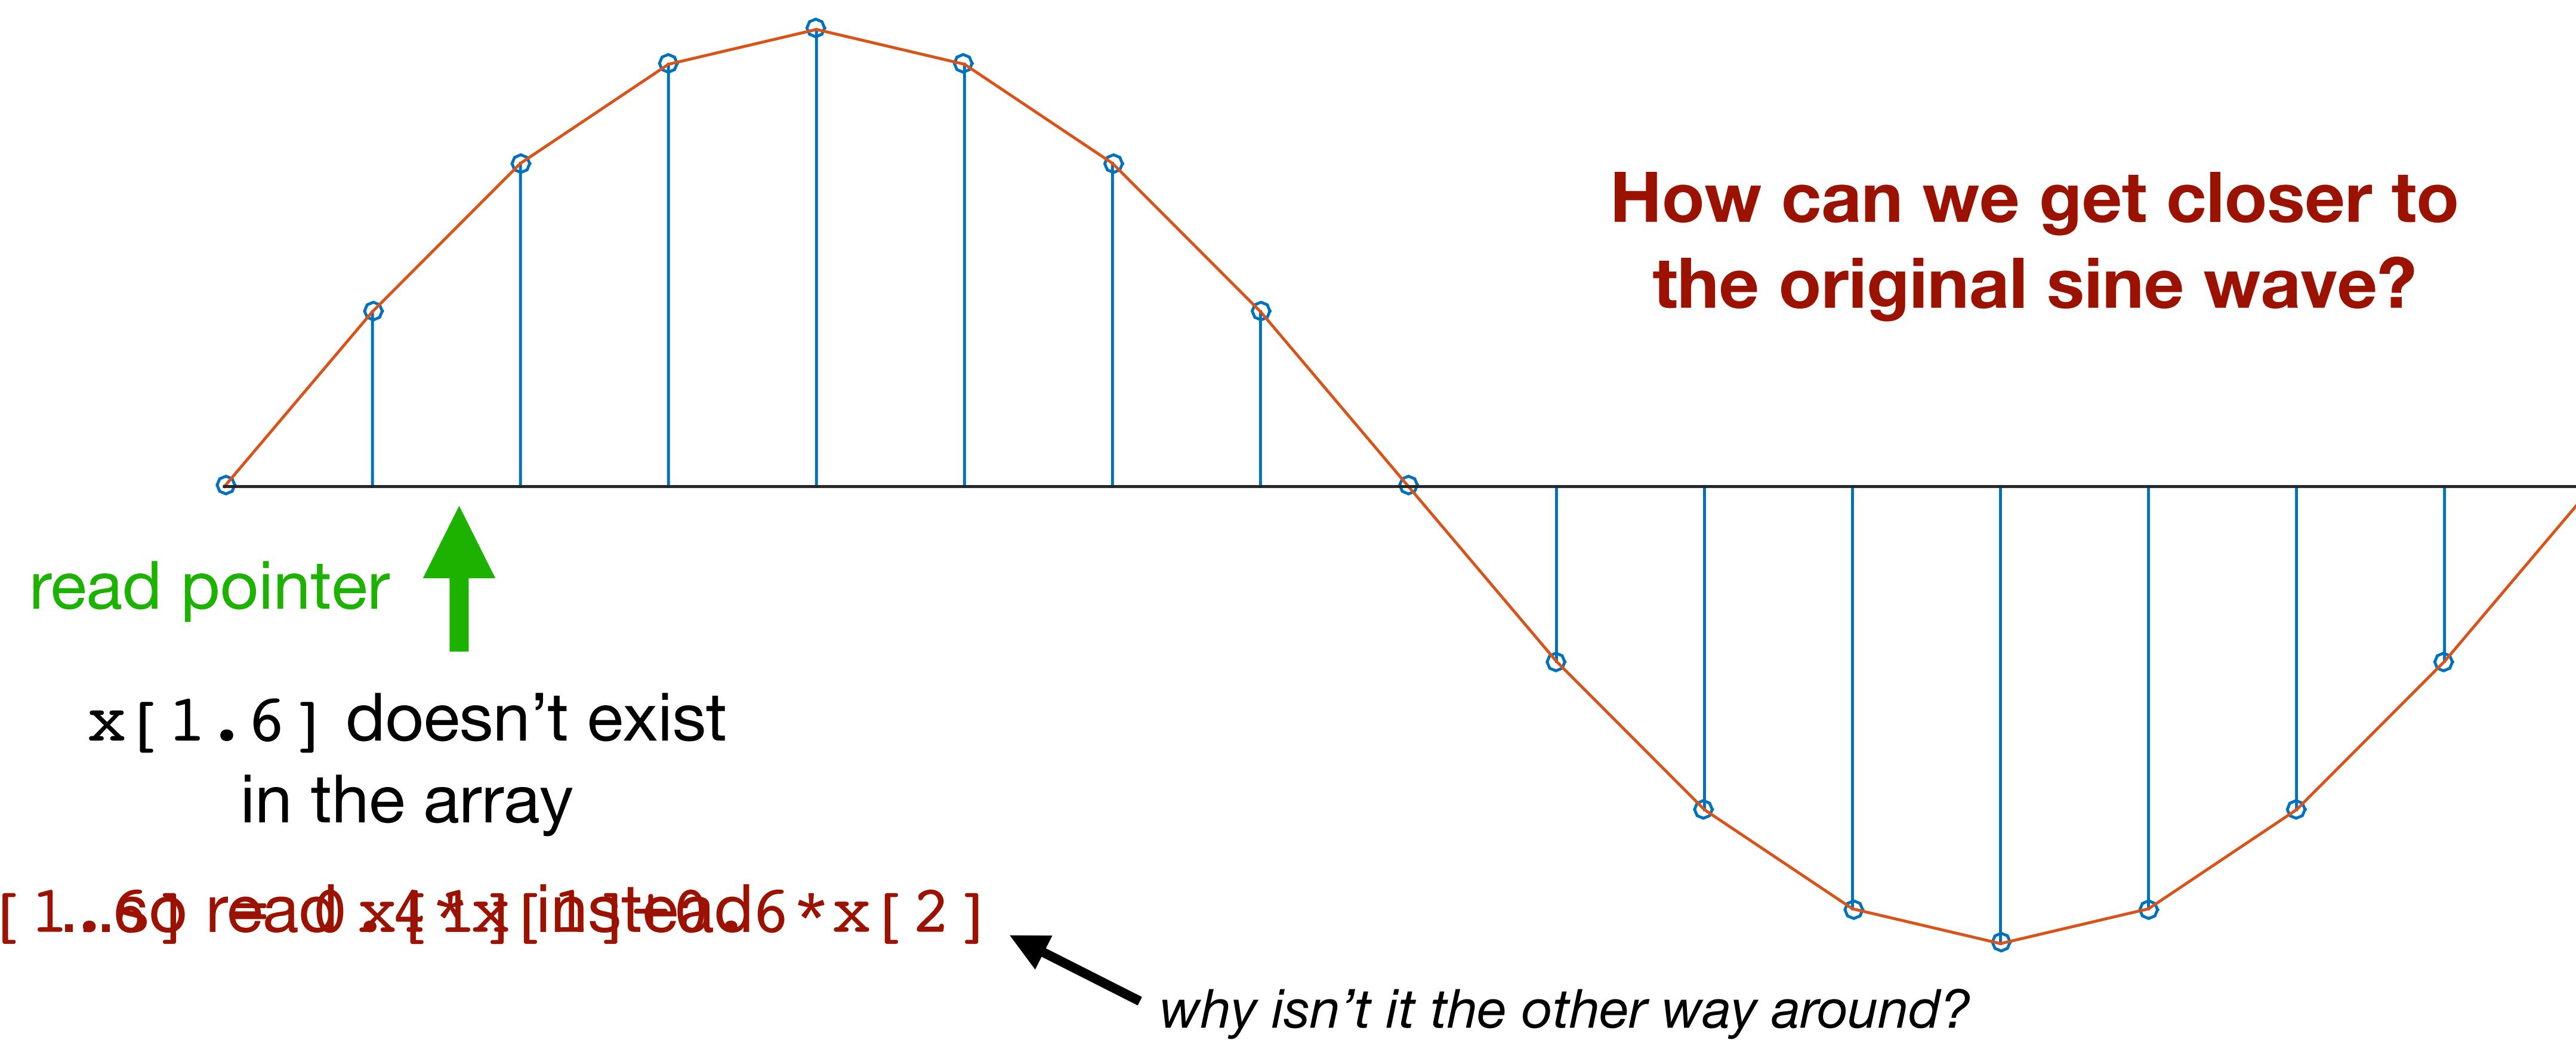
**What we wish we could read:**



# A better fractional index?

- We want to read fractional array values that, strictly speaking, don't exist...

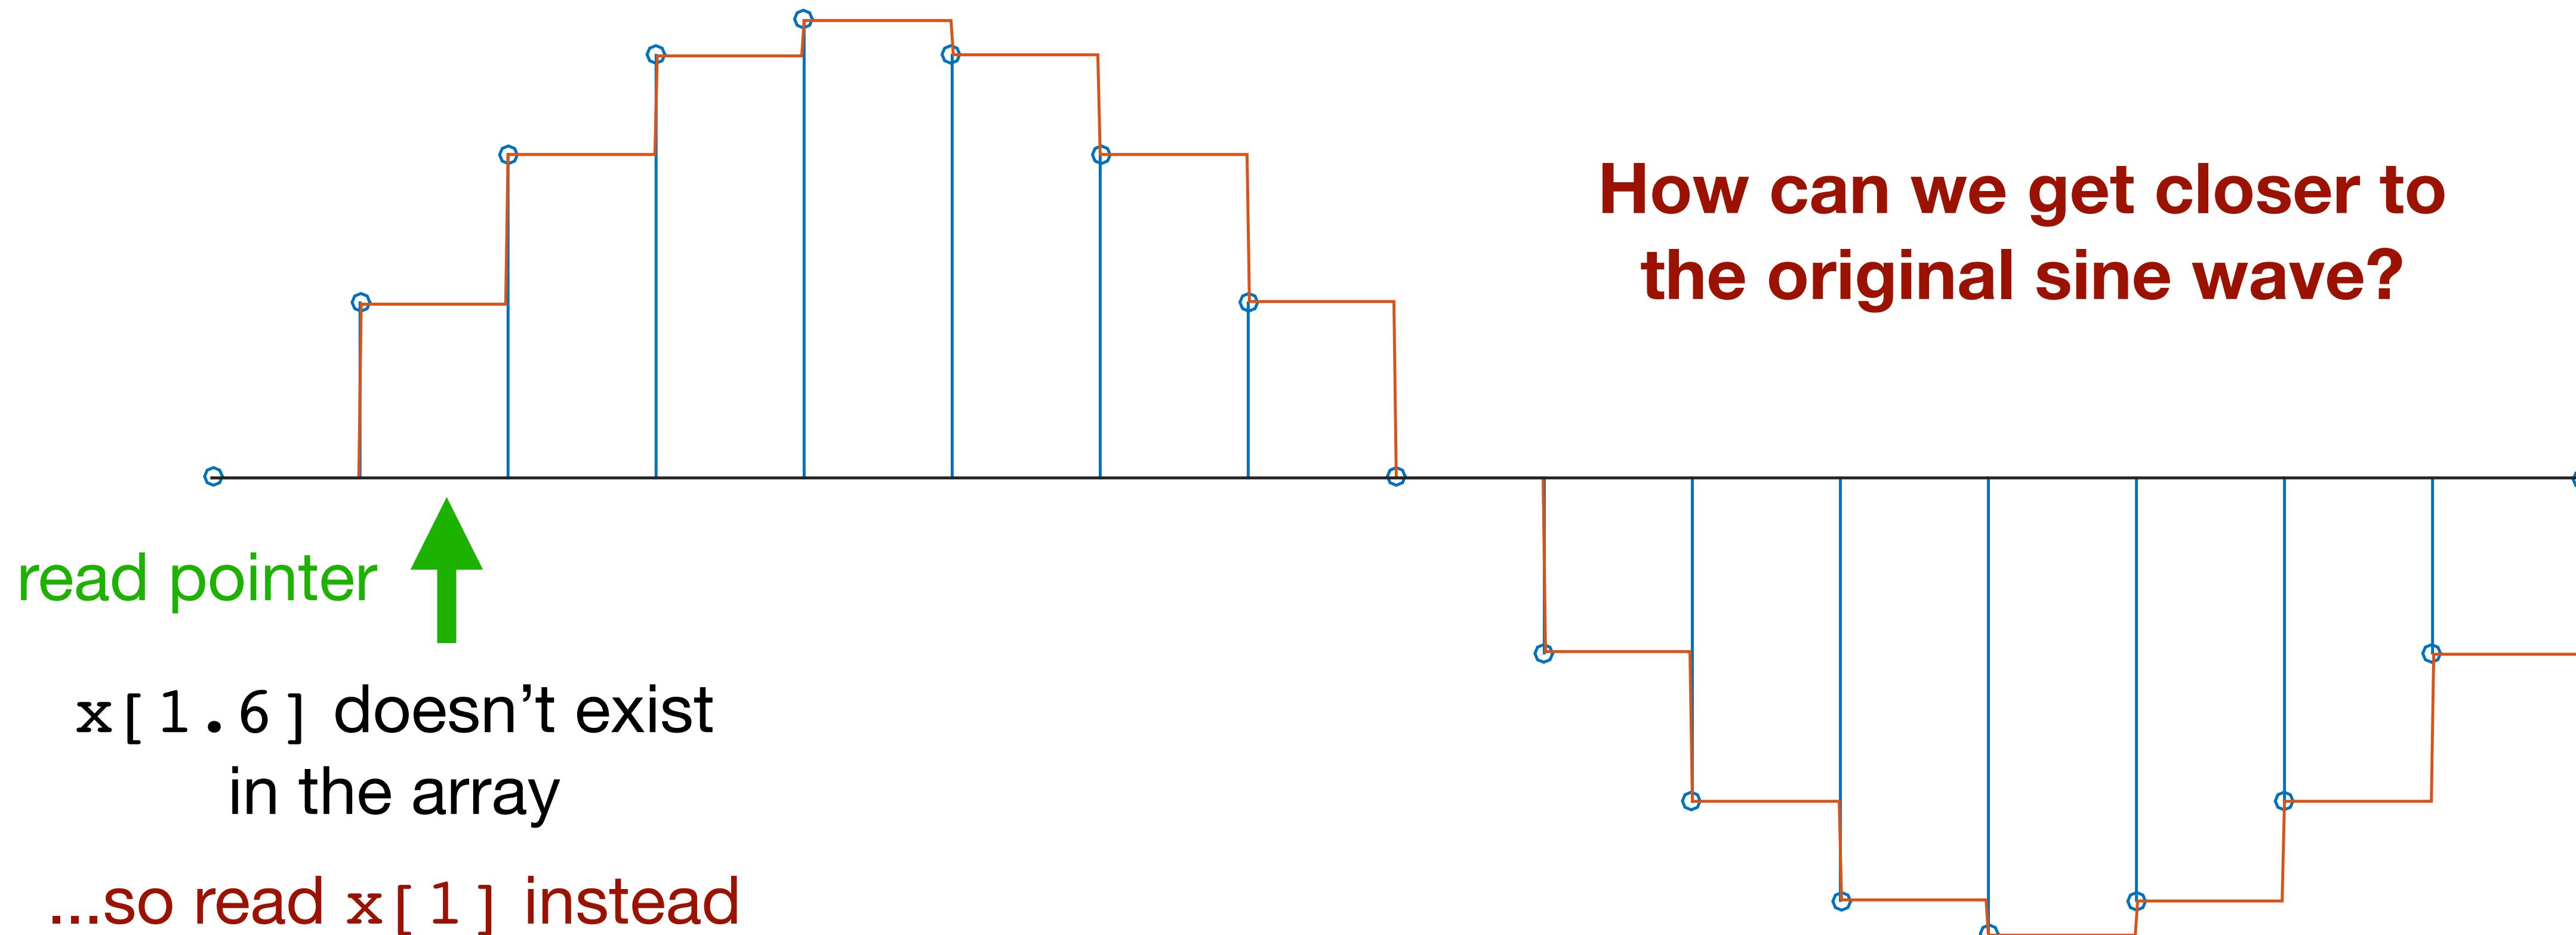
Take a weighted average of the neighbouring samples



# A better fractional index?

- We want to read fractional array values that, strictly speaking, don't exist...

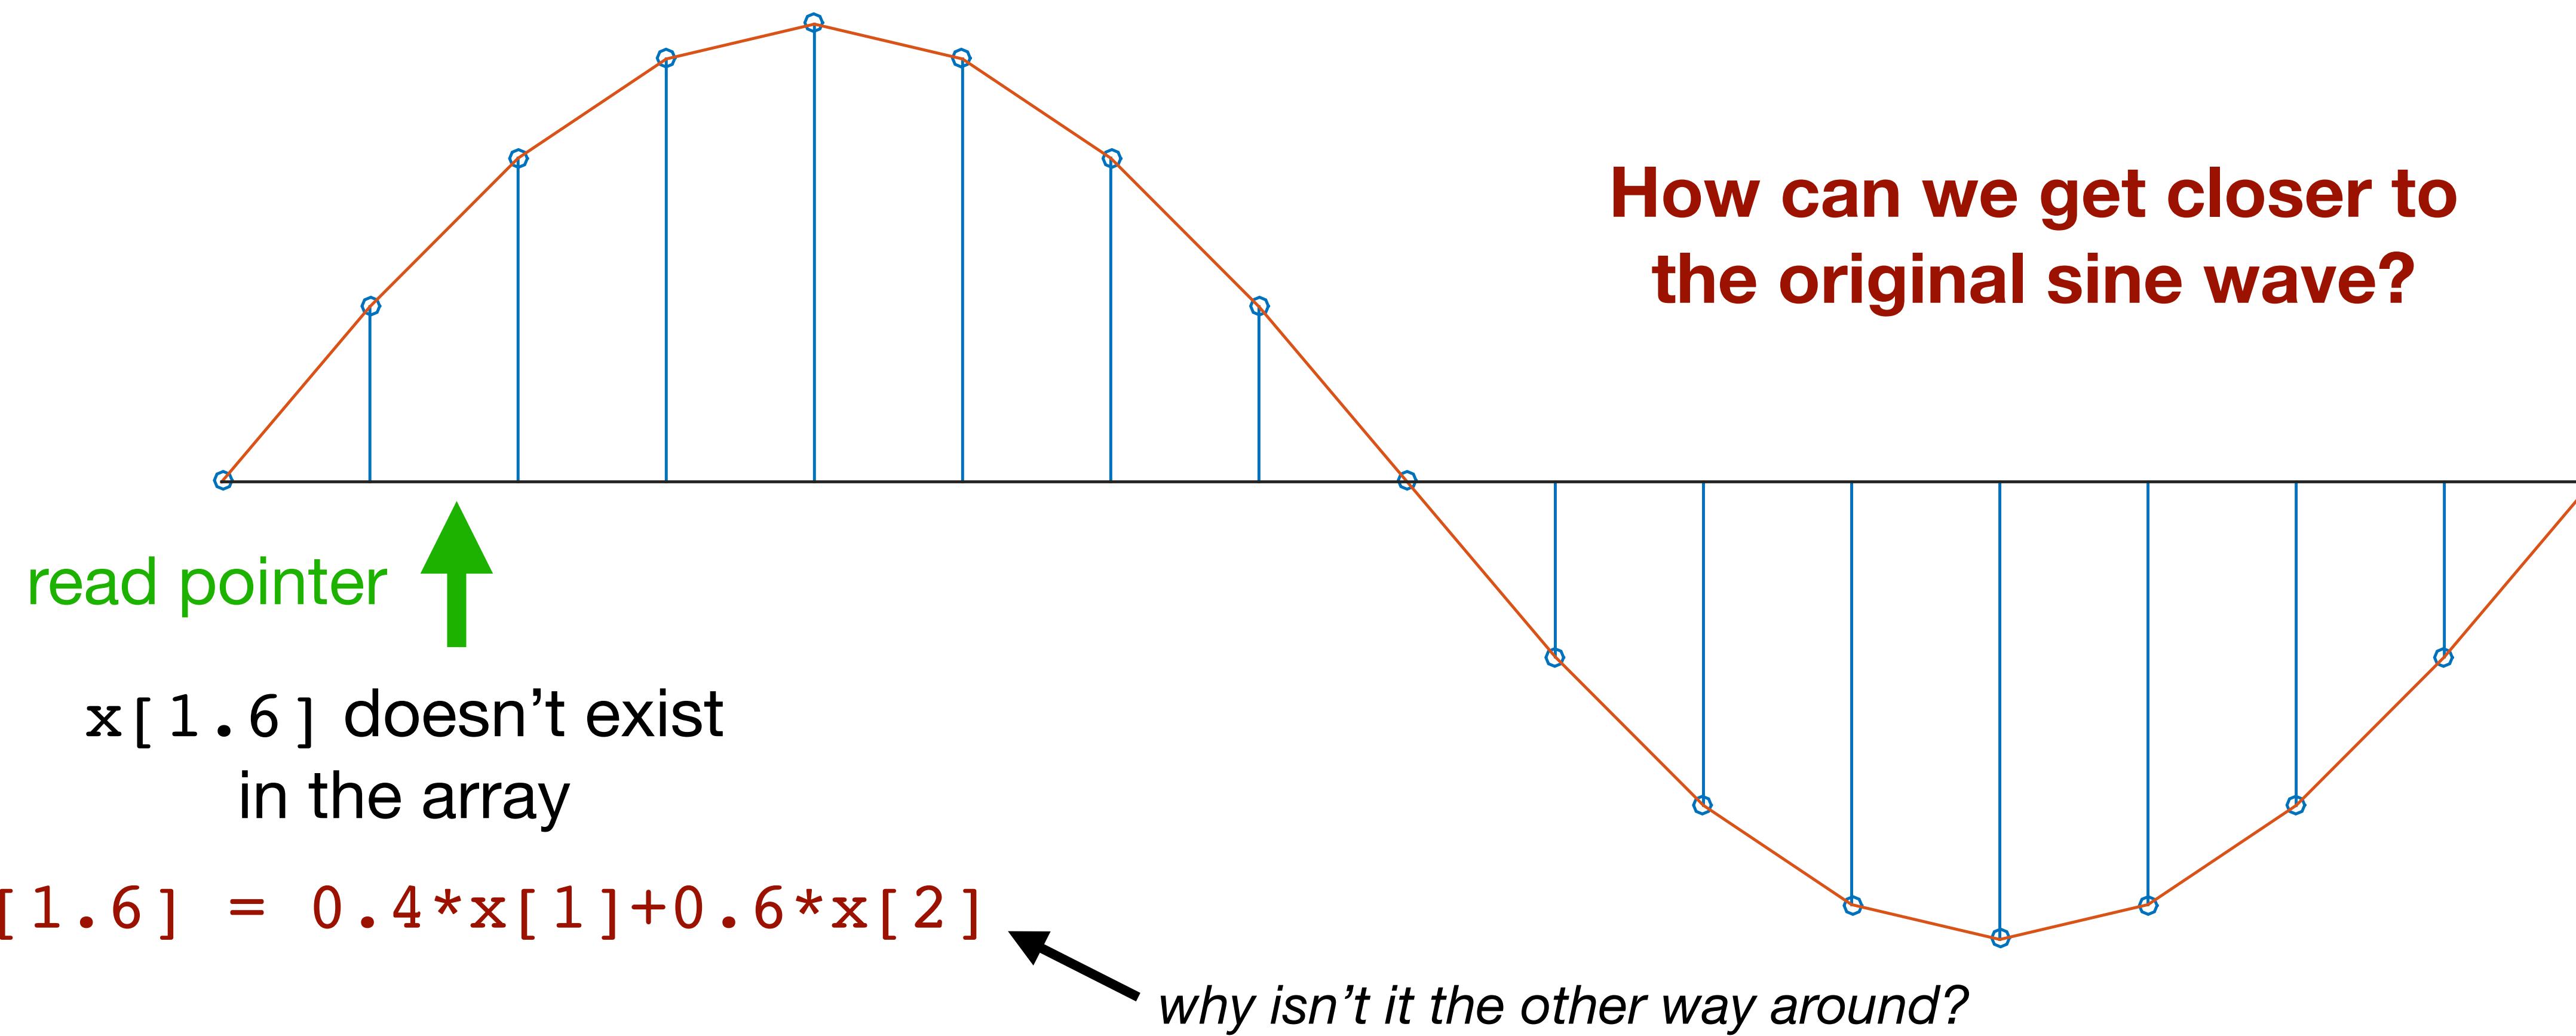
**What our code is *actually* reading:**



# A better fractional index?

- We want to read fractional array values that, strictly speaking, don't exist...

**Take a weighted average of the neighbouring samples**



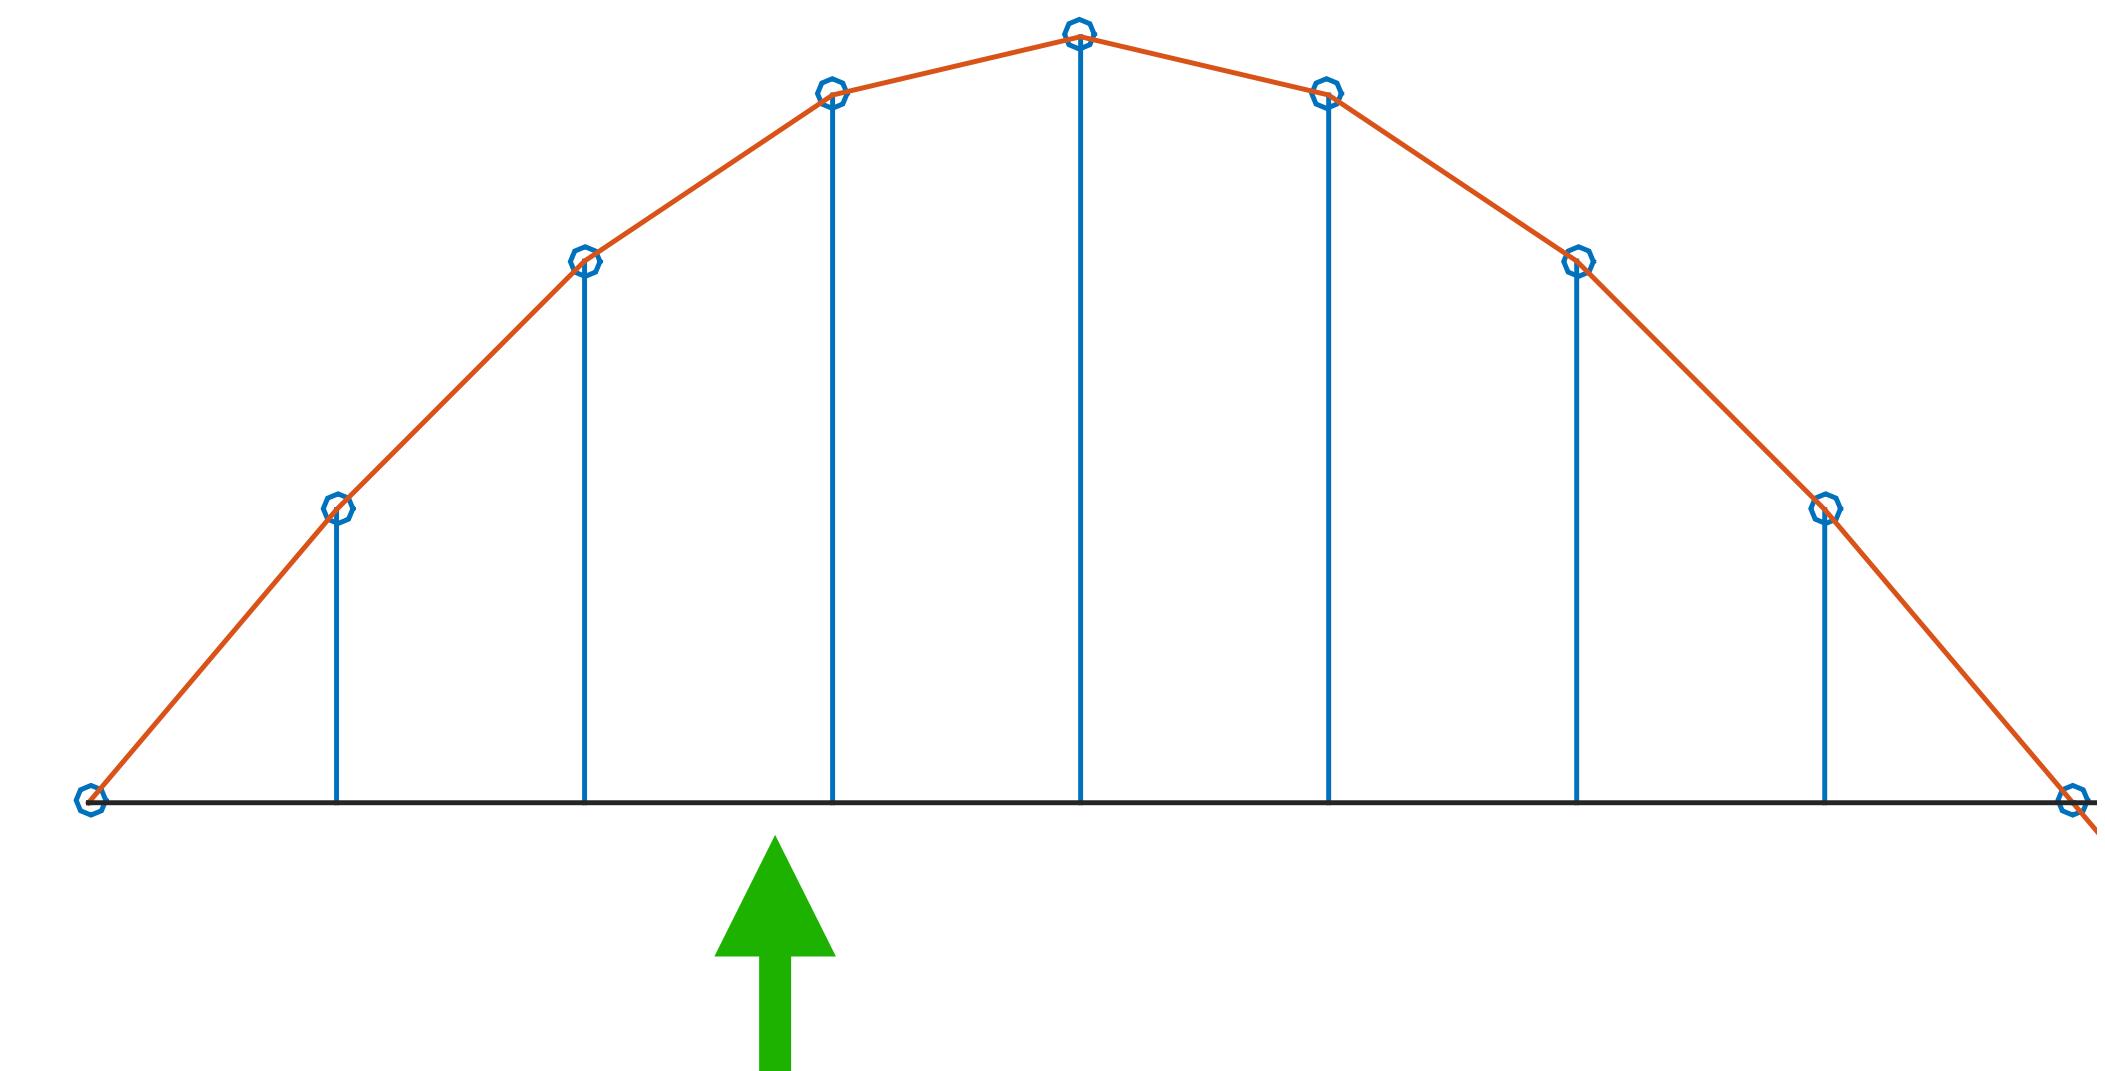
# Linear interpolation

- Linear interpolation means taking a weighted average of neighbouring samples
  - The visual metaphor is drawing a straight line between the samples
- It is not a fully accurate reconstruction of an underlying continuous signal
  - Linear interpolation, like nearest-neighbour interpolation, will introduce distortion
  - A perfect mathematical interpolation needs sinc functions, which are not causal
  - There are more accurate approximations such as cubic interpolation
  - See Reiss & McPherson, *Audio Effects: Theory Implementation and Application*, p. 26
- To implement linear interpolation, we need to know:
  - The samples on either side of the fractional value
  - The fractional component of the index, to determine the weight of each sample

# Linear interpolation

- Suppose  $x$  is our signal,  $n$  is the integer part of the index and  $f$  is the fractional part:  $0 \leq f < 1$ . Then we have:

$$x(n + f) = (1 - f)x[n] + fx[n + 1]$$



- In a wavetable, what happens if  $n+1$  is off the edge of the table?
  - e.g.  $x(15.1)$  in a 16-sample table
  - Valid indices in the array are 0 to 15
  - Wrap around to the beginning:**  $x(15.1) = 0.9x[15] + 0.1x[0]$
- Task:** implement linear interpolation in the `wavetable` example

$$x(2.7) = 0.3x[2] + 0.7x[3]$$

# Linear interpolation code

```
// The pointer will take a fractional index. Look for the sample on
// either side which are indices we can actually read into the buffer.
// If we get to the end of the buffer, wrap around to 0.
int indexBelow = floorf(gReadPointer);
int indexAbove = indexBelow + 1;
if(indexAbove >= gWavetableLength)
    indexAbove = 0;

// For linear interpolation, we need to decide how much to weigh each
// sample. The closer the fractional part of the index is to 0, the
// more weight we give to the "below" sample. The closer the fractional
// part is to 1, the more weight we give to the "above" sample.
float fractionAbove = gReadPointer - indexBelow;
float fractionBelow = 1.0 - fractionAbove;

// Calculate the weighted average of the "below" and "above" samples
float out = gAmplitude * (fractionBelow * gWavetable[indexBelow] +
                           fractionAbove * gWavetable[indexAbove]);

// Increment read pointer and reset to 0 when end of table is reached
gReadPointer += gWavetableLength * gFrequency / context->audioSampleRate;
while(gReadPointer >= gWavetableLength)
    gReadPointer -= gWavetableLength;
```

**n**, the integer below  $n+f$

**n+1**

**f**, the fractional component

our formula:

$$x(n+f) = (1-f)x[n] + fx[n+1]$$

# Populating the wavetable

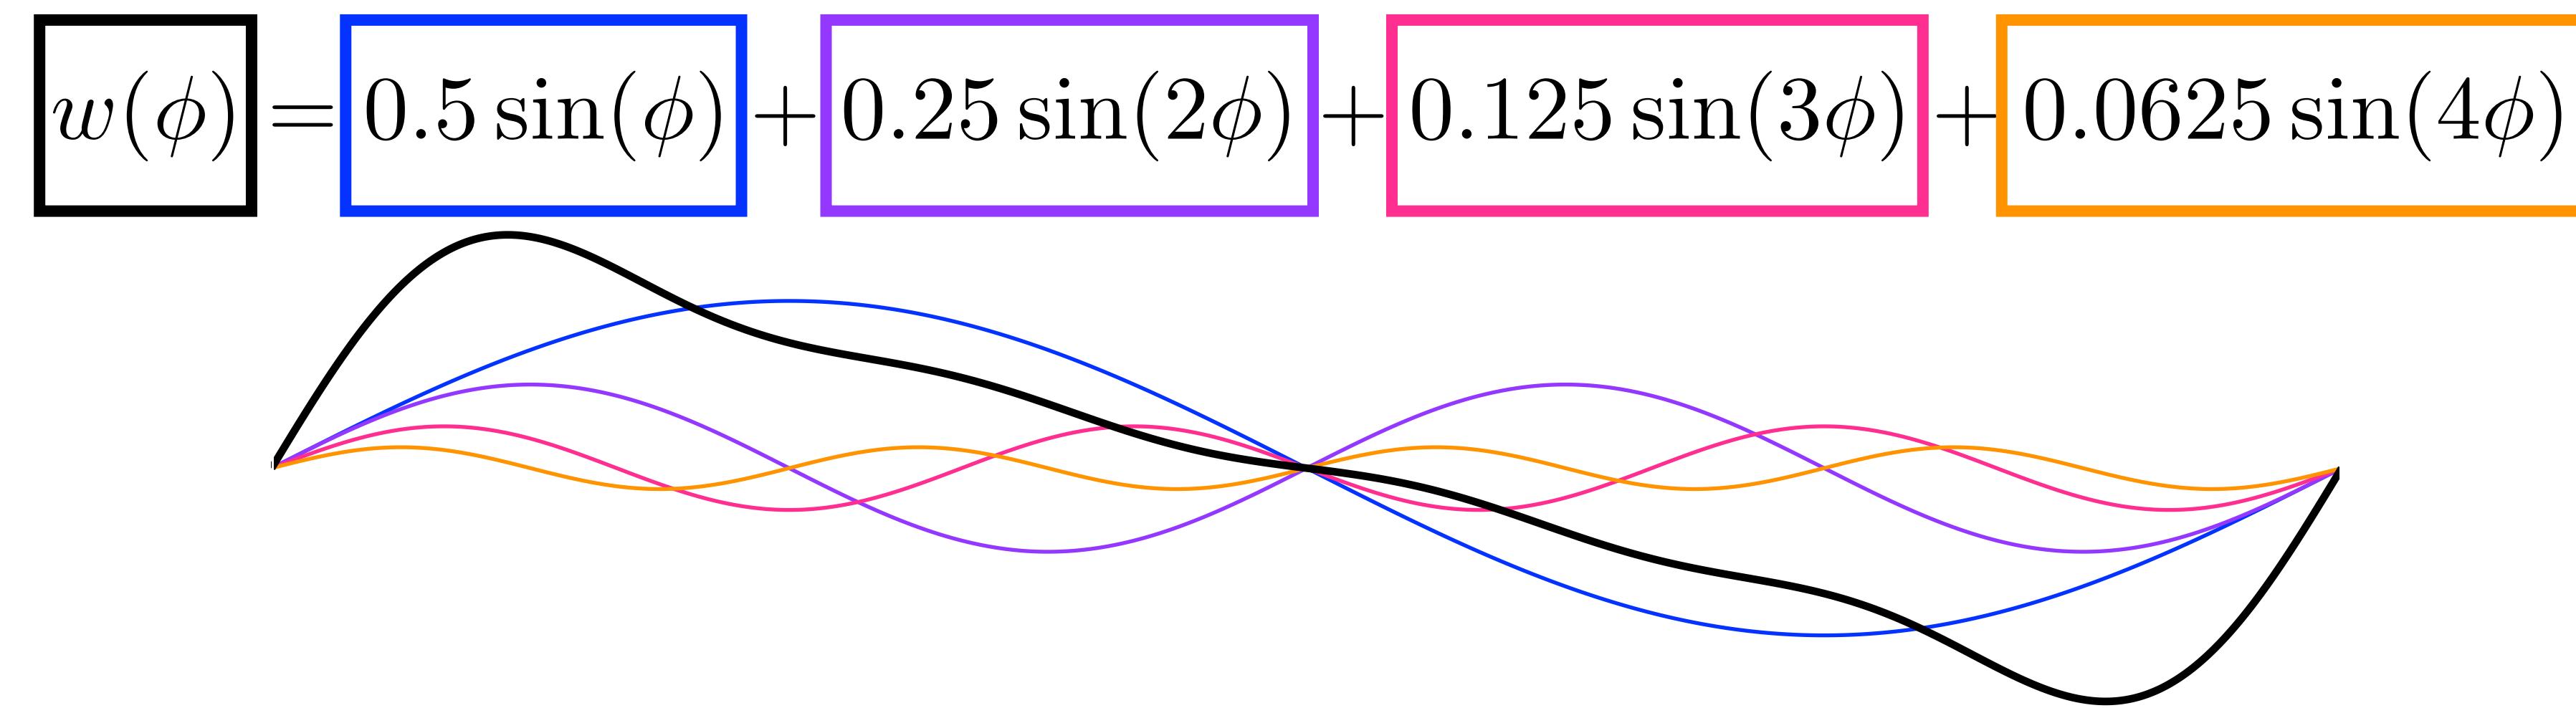
- How can we generate the contents of the wavetable?
  - Load samples from a file
  - Generate mathematically
- Suppose the wavetable has length N
  - Then the range of array indexes goes from 0 to N-1
  - We can populate its contents with a `for()` loop inside of `setup()`:

```
for(unsigned int n = 0; n < gWavetableLength; n++) {  
    gWavetable[n] = [REDACTED]  
}
```

- **Task:** change `wavetable` to implement a `sine` rather than a triangle wave
  - Remember, the range 0 to N-1 needs to correspond to one complete oscillation
    - i.e. an index of N would correspond to a phase of  $2\pi$  ( $2 \cdot 0 * M_PI$ )

# More complex waveforms

- One way of building more complex waveforms is to add several **harmonics** of a **fundamental frequency**
  - This process is known as **additive synthesis**, and is a whole topic of its own!
  - In general, you can sum together **any phase** of **integer multiples** of the fundamental frequency



- **Task:** in the **wavetable** project, implement the above equation in `setup()`
  - What is the relationship between  $\phi$  and n in the `for()` loop?

# Keep in touch!

Social media:

**@BelaPlatform**

**forum.bela.io**

**blog.bela.io**

More resources and contact info at:

**learn.bela.io/resources**