

C++ Real-Time Audio Programming with Bela

Dr Andrew McPherson

Centre for Digital Music
School of Electronic Engineering and Computer Science
Queen Mary University of London

Founder and Director, Bela

Course topics

Programming topics

- Working in real time
- Buffers and arrays
- Parameter control
- Classes and objects
- Analog and digital I/O
- Filtering
- Circular buffers
- Timing in real time
- State machines
- MIDI
- Block-based processing
- Threads
- Fixed point arithmetic
- ARM assembly language



Music/audio topics

- Oscillators
- Samples
- Wavetables
- Filters
- Control voltages
- Gates and triggers
- Delays and delay-based effects
- Metronomes and clocks
- Envelopes
- ADSR
- MIDI
- Additive synthesis
- Phase vocoders
- Impulse reverb

Today

Lecture 4: Parameter control

What you'll learn today:

Using the Bela GUI to control projects from the browser
Linear and logarithmic scales

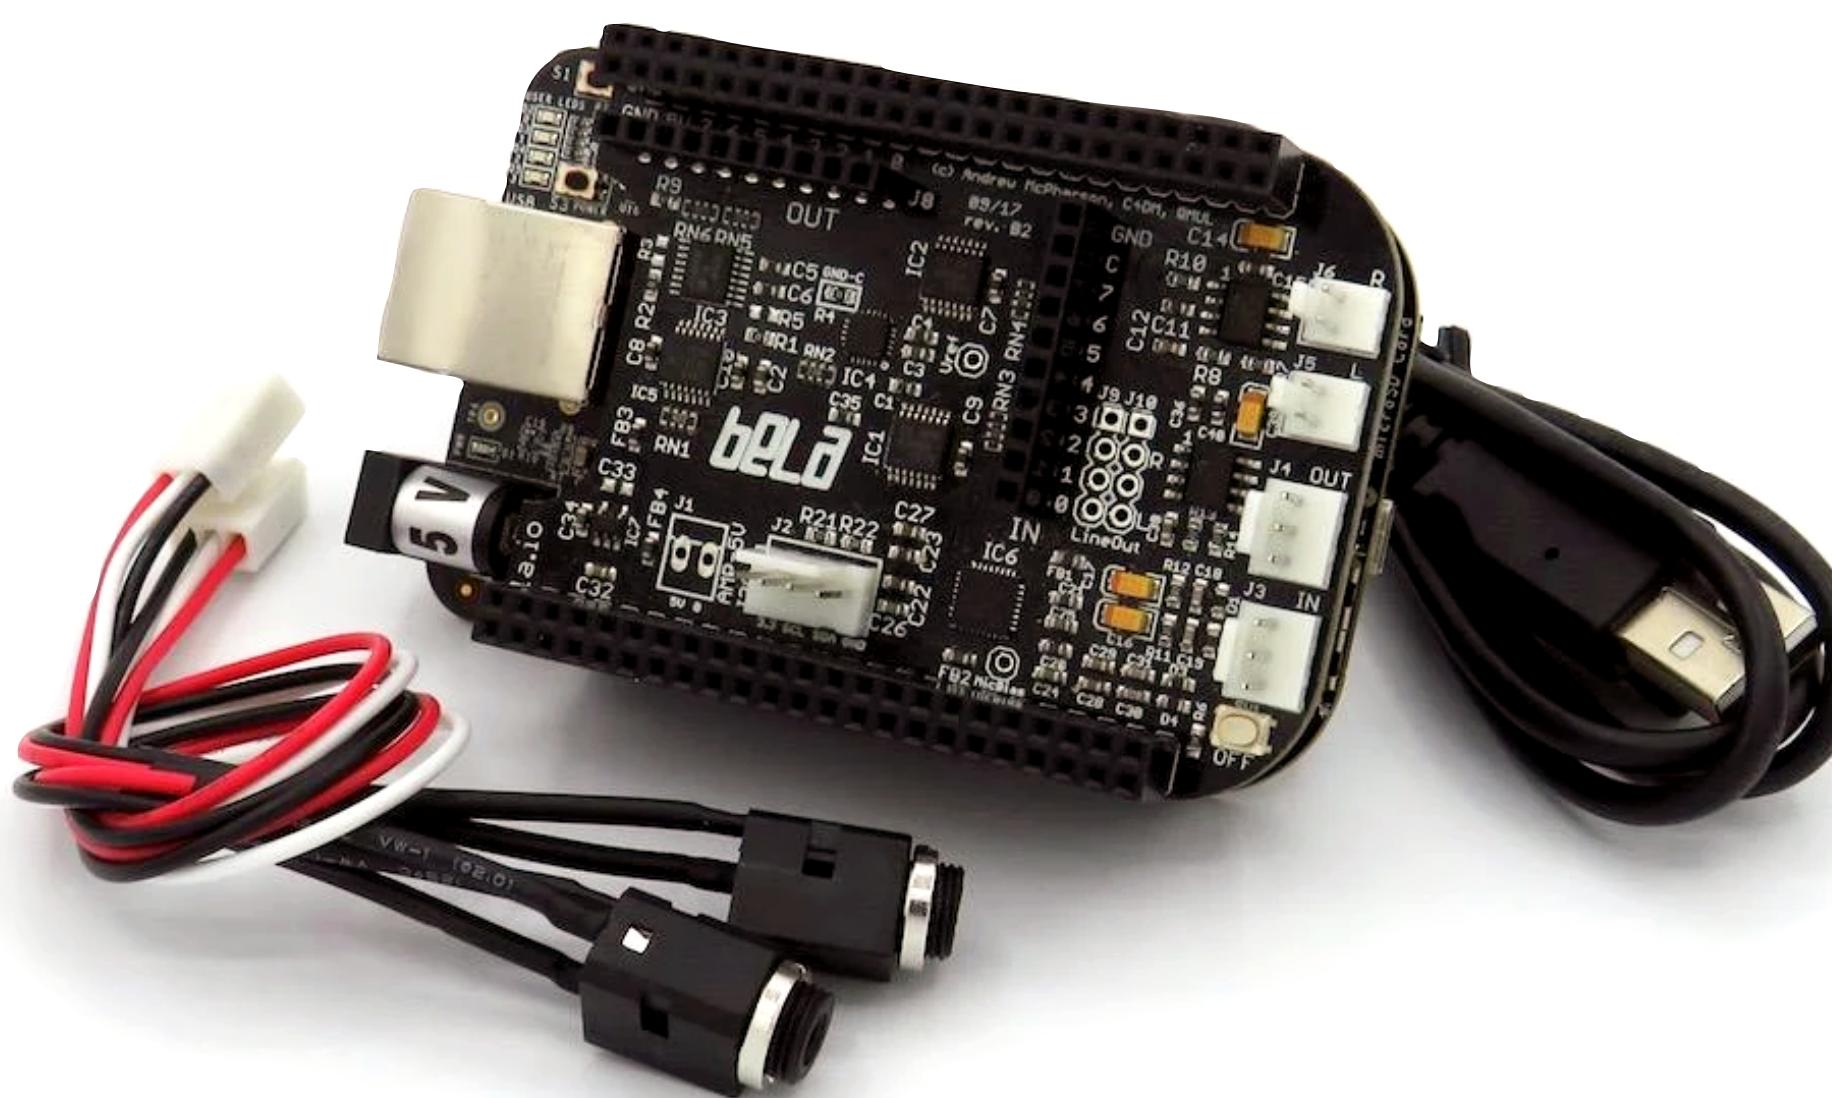
What you'll make today:

Multiple oscillators with detuning

Companion materials:

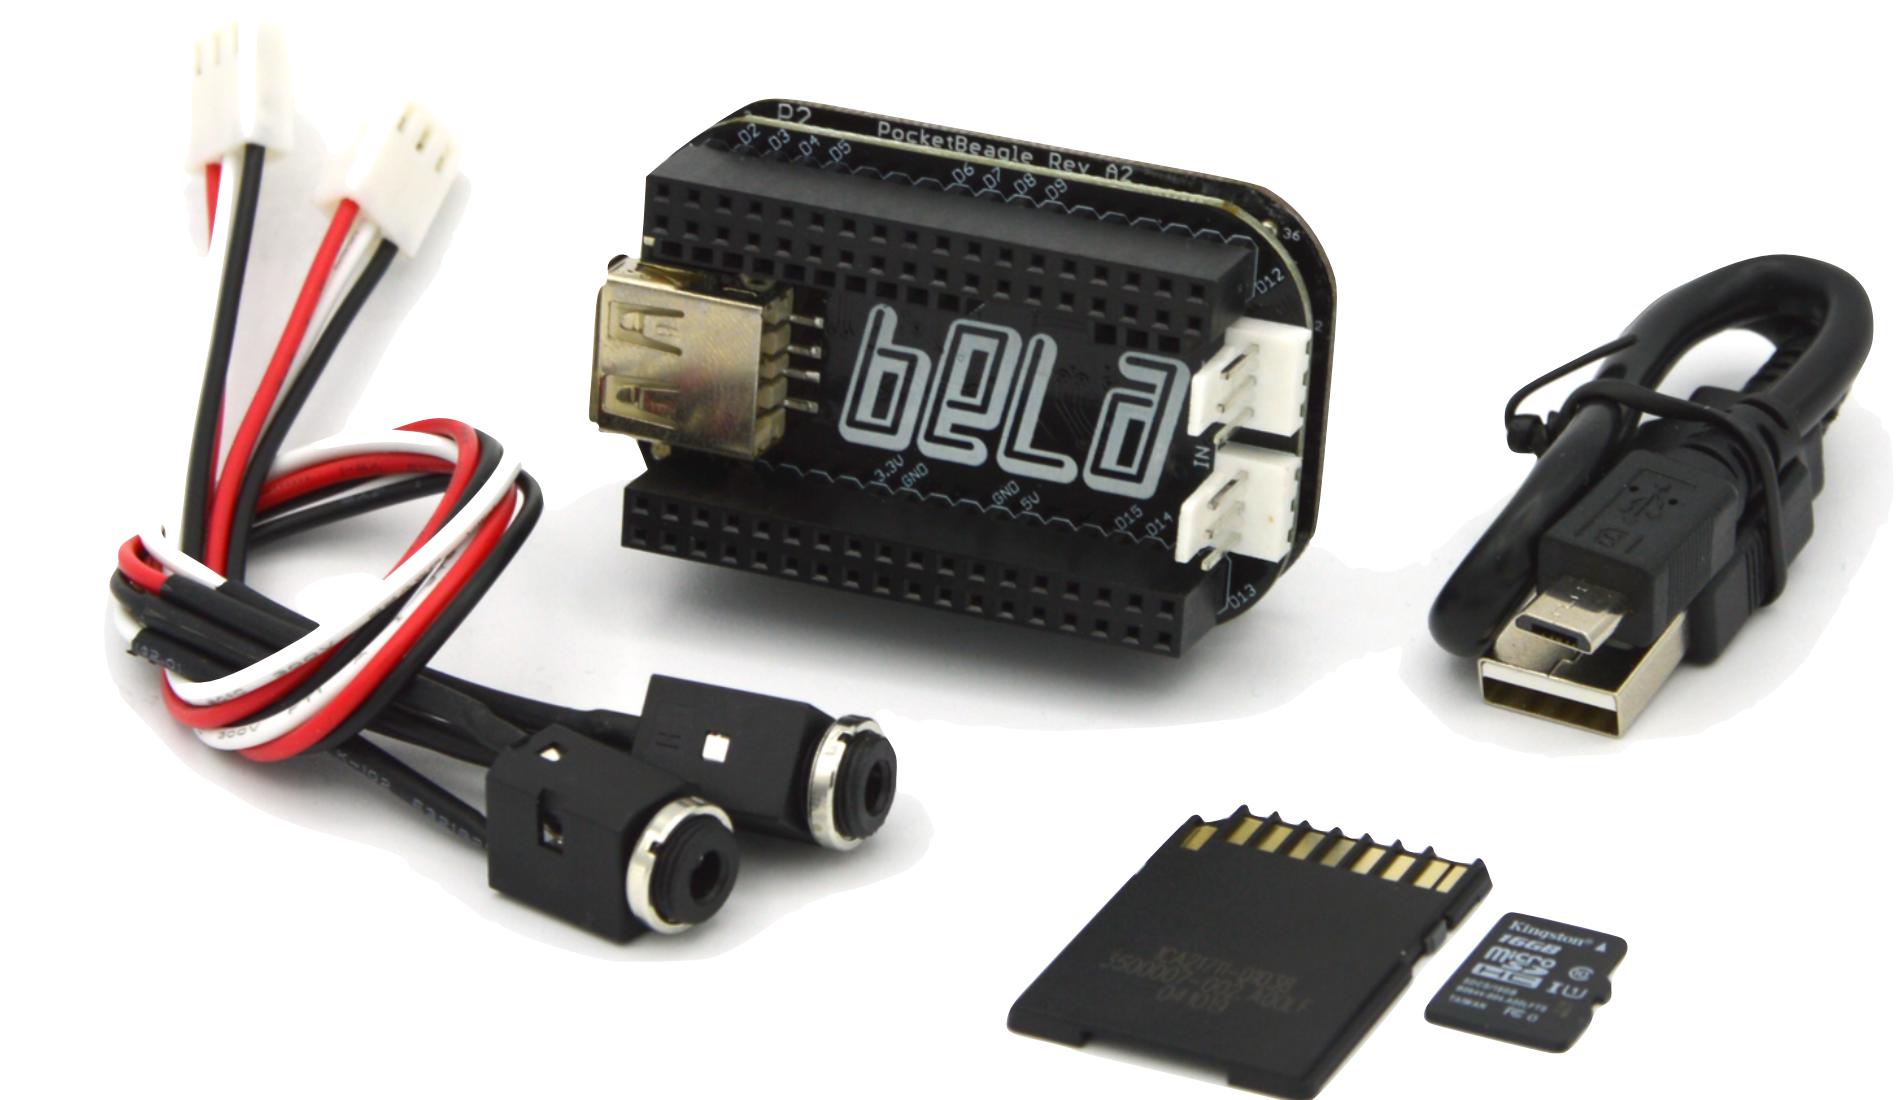
github.com/BelaPlatform/bela-online-course

What you'll need



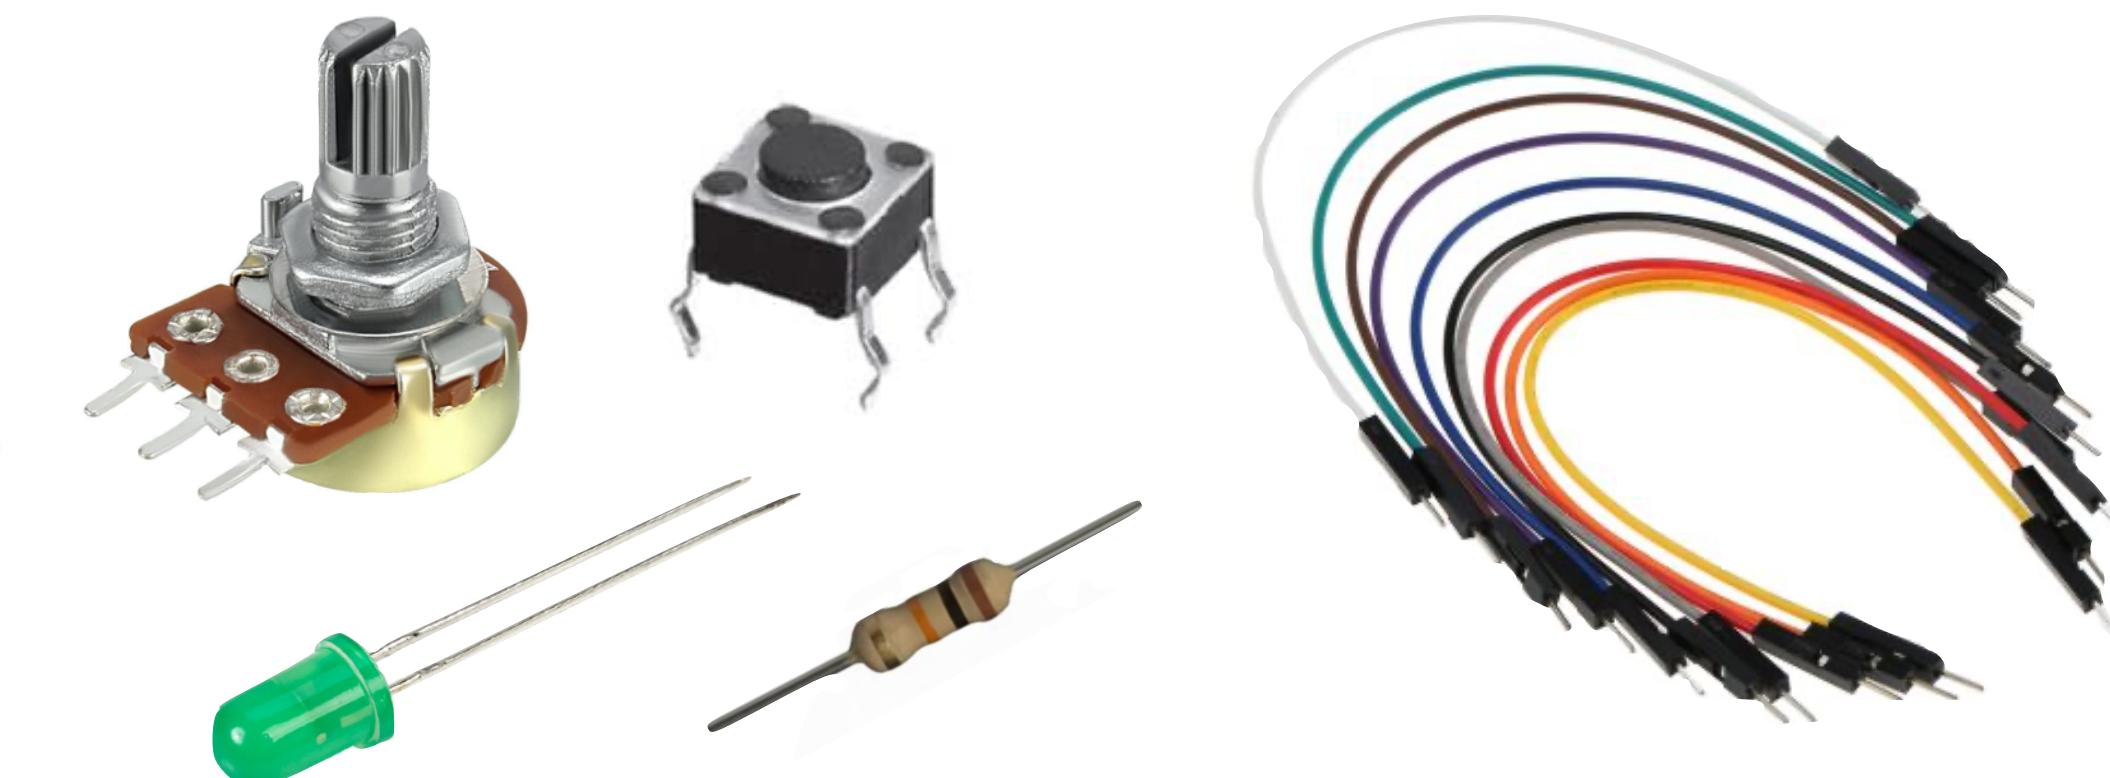
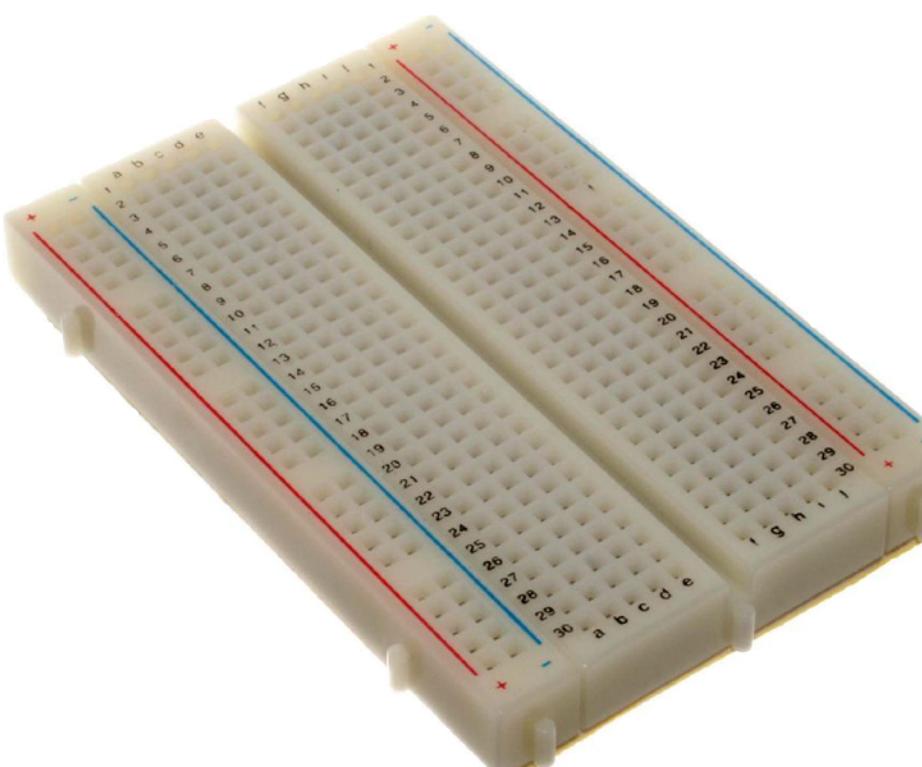
Bela Starter Kit

or



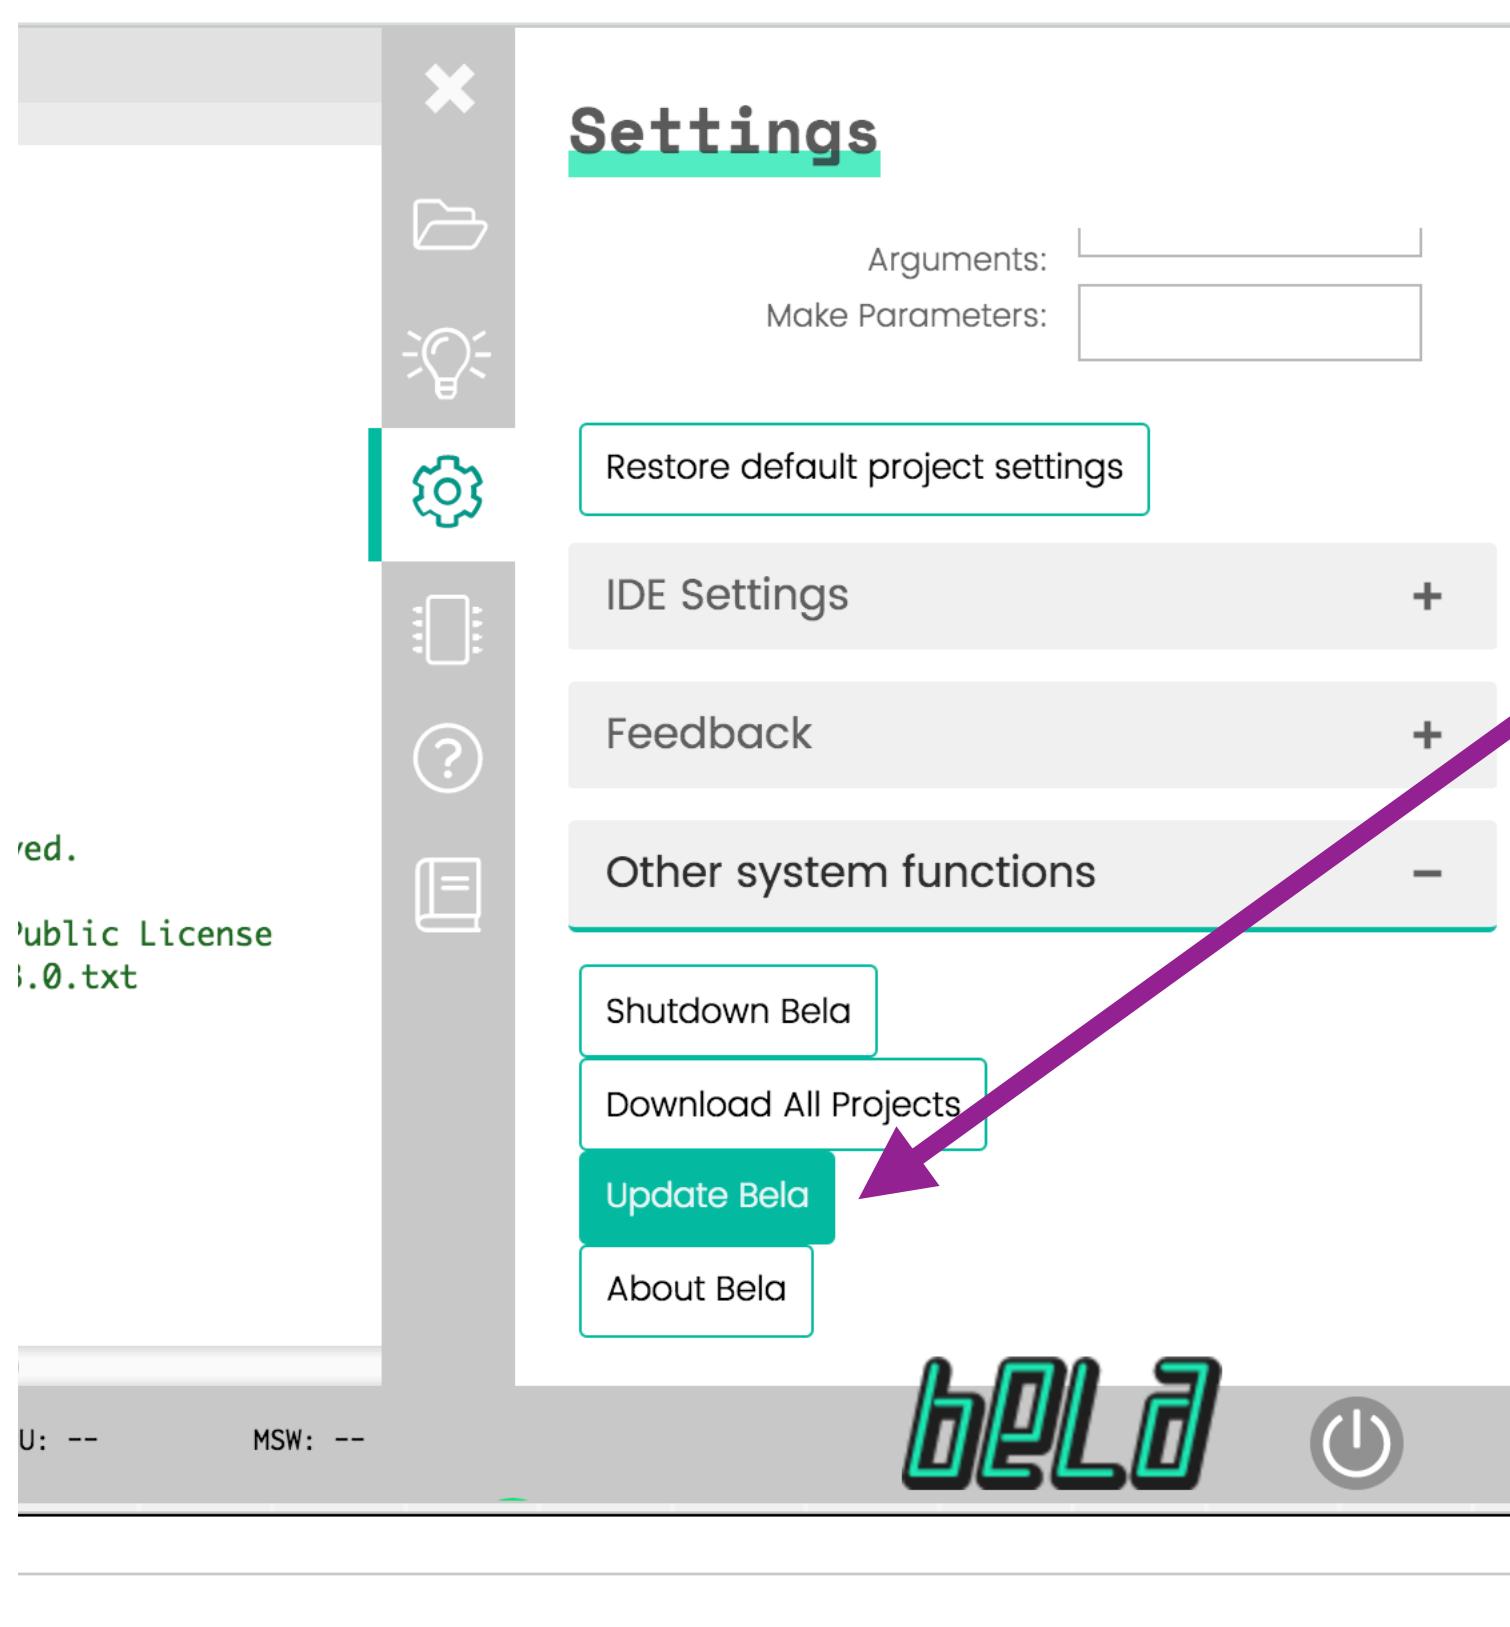
Bela Mini Starter Kit

Recommended
for some lectures:

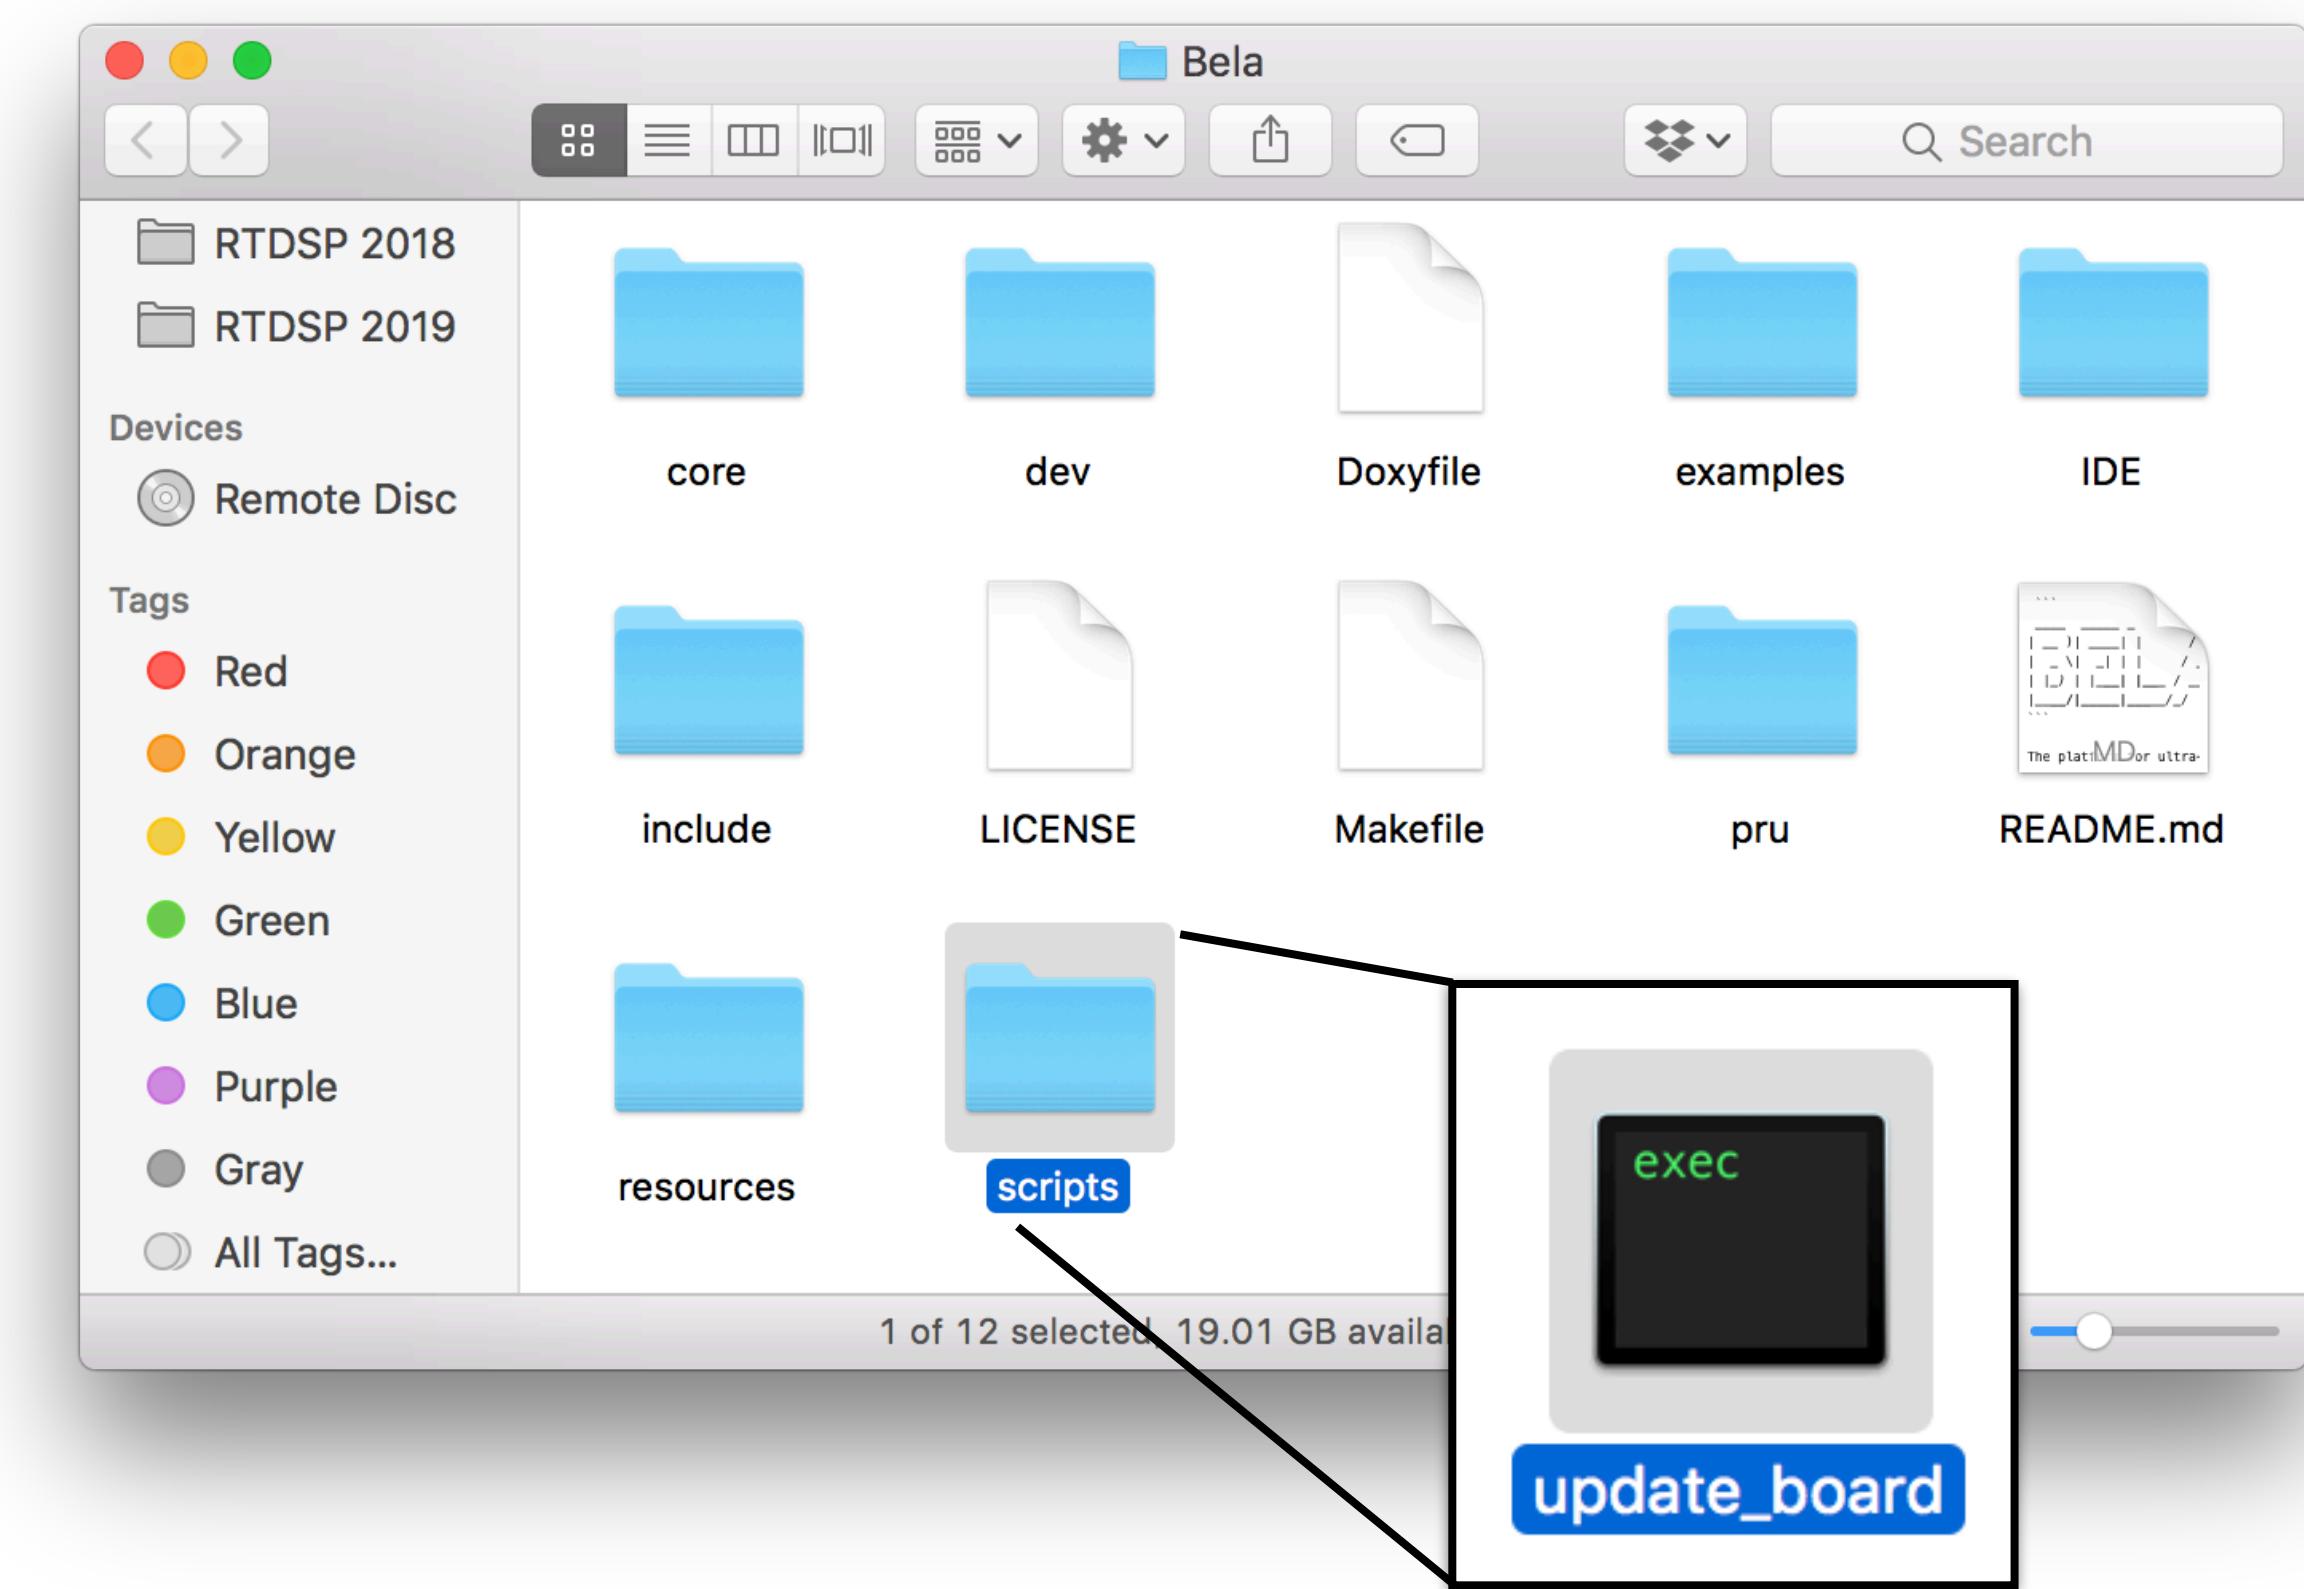


Updating Bela

- Download the latest Bela core software from: learn.bela.io/update
- Two ways to update the board:
 - By IDE (recommended)



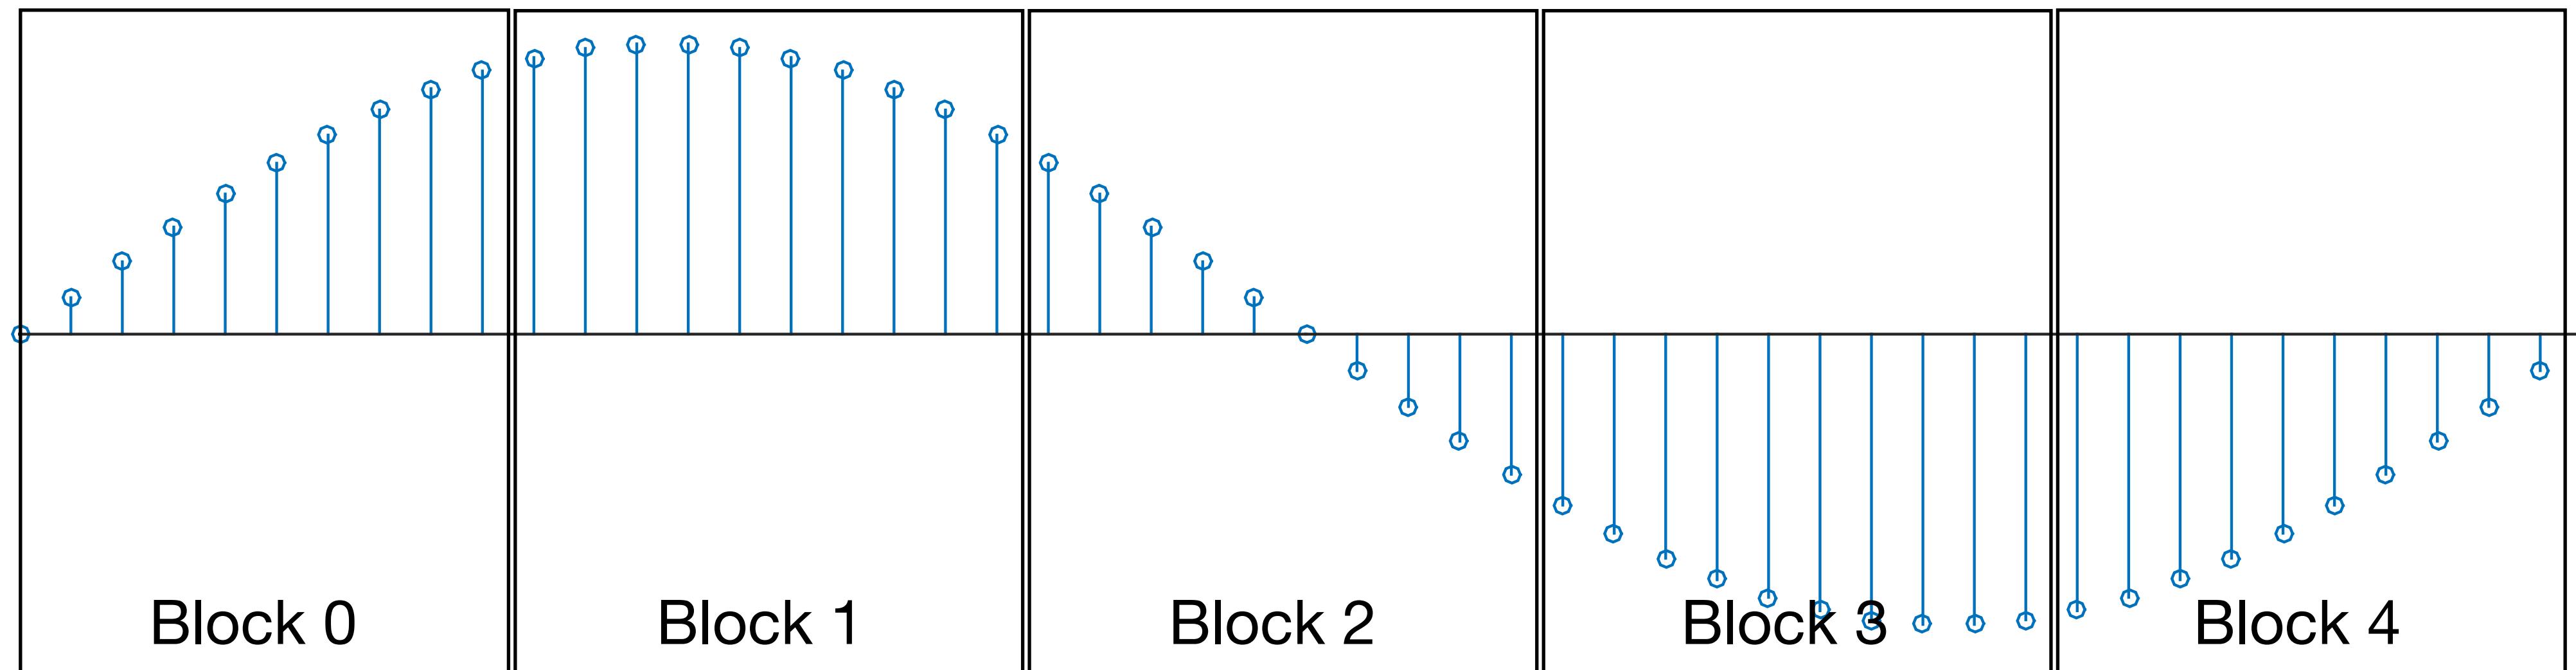
Update
Bela



By Script (Mac and Linux)
unzip the archive

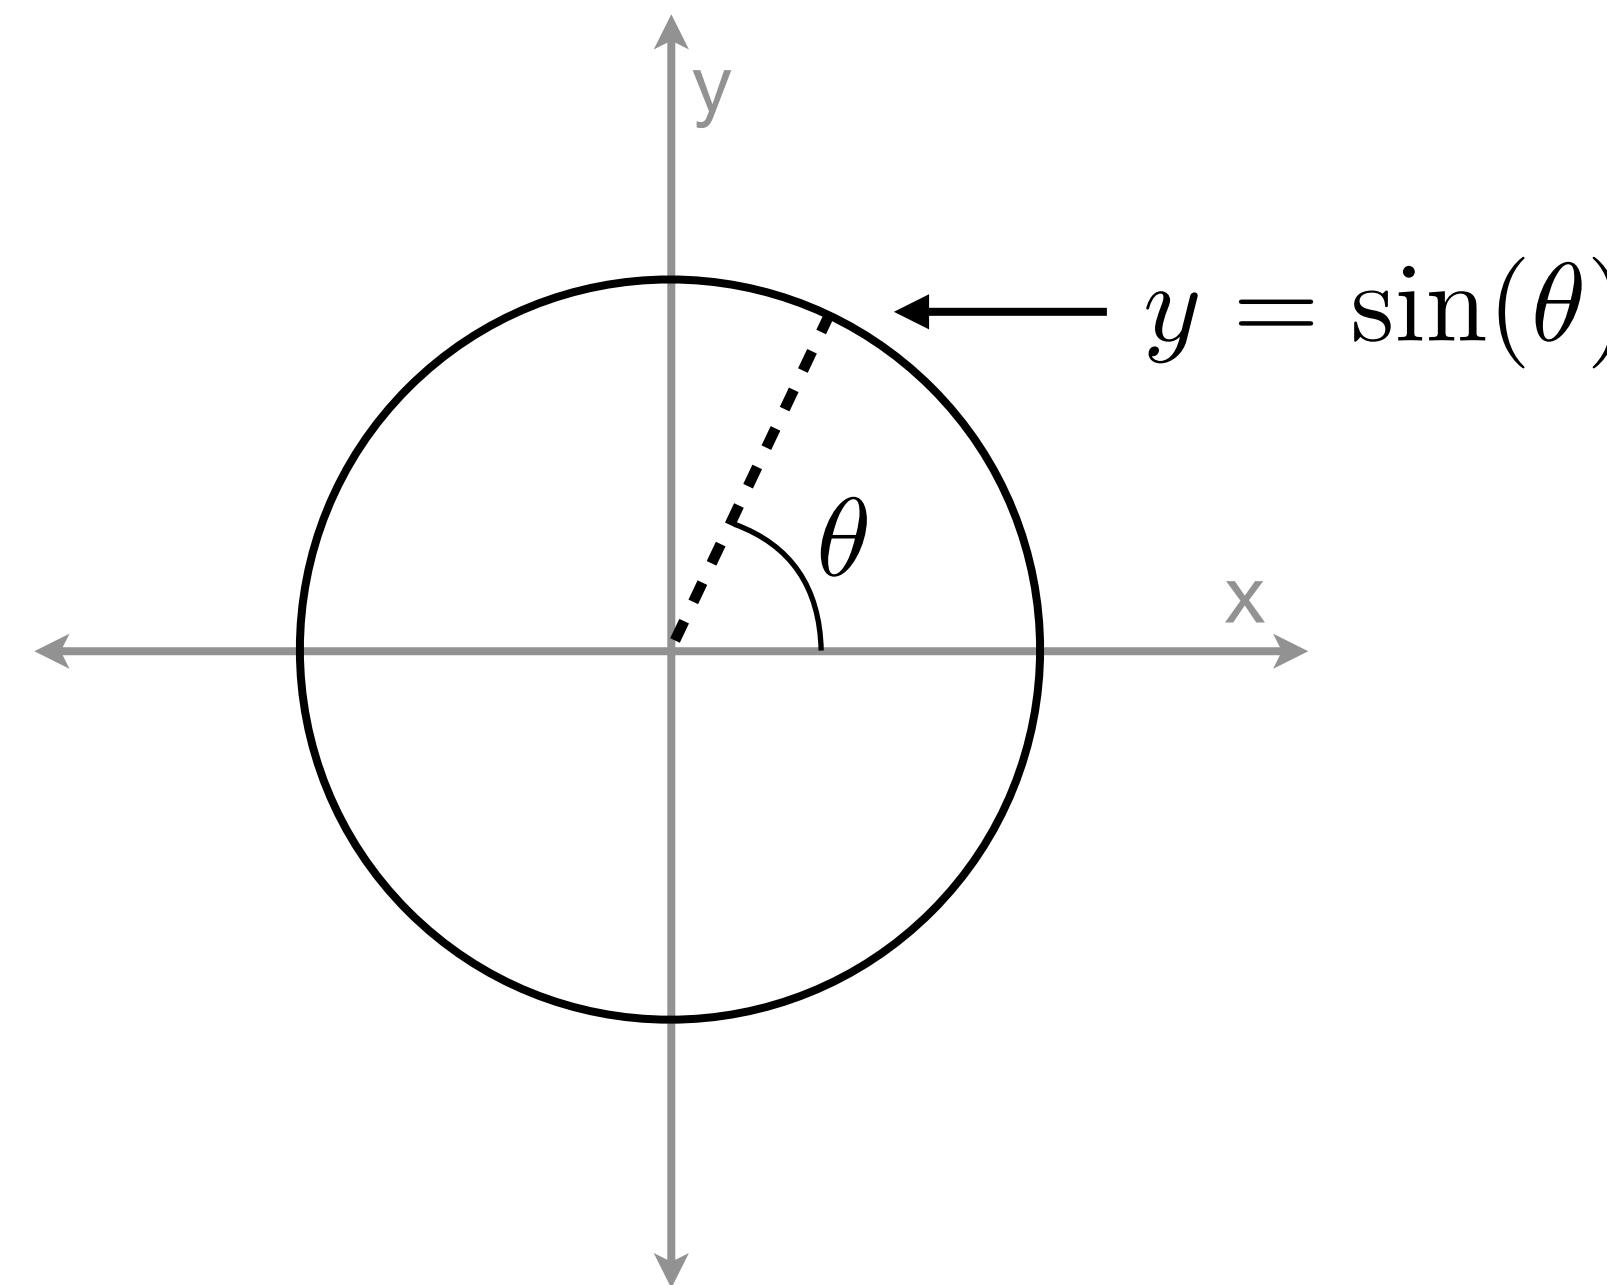
Real-time sine oscillator

- In Lecture 1, we made a real-time sine oscillator that generated one block at a time:



Oscillator frequency and phase

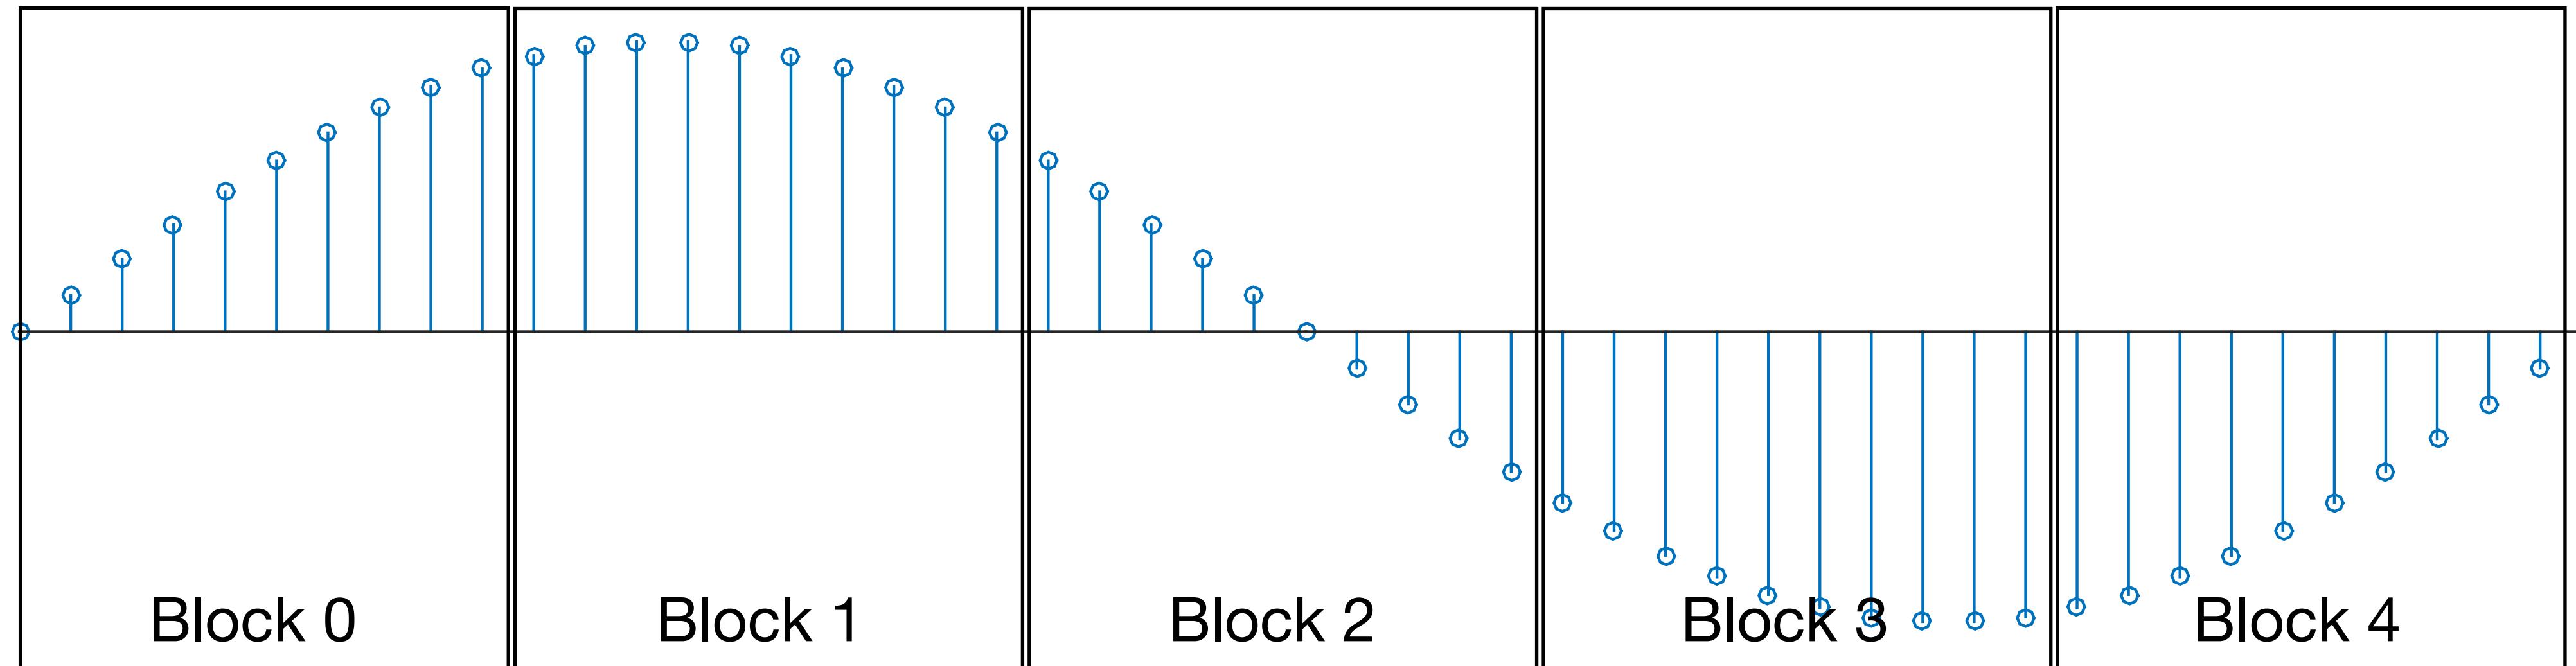
- The argument to the $\sin()$ function is known as the **phase**
 - Possible range of (unique) values: 0 to 2π
 - Analogous to the **y**-coordinate on the unit circle for a particular angle



- What does **frequency** correspond to in this analogy?
 - How fast we spin around the circle: the **derivative of the phase**

Real-time sine oscillator

- In Lecture 1, we made a real-time sine oscillator that generated one block at a time:

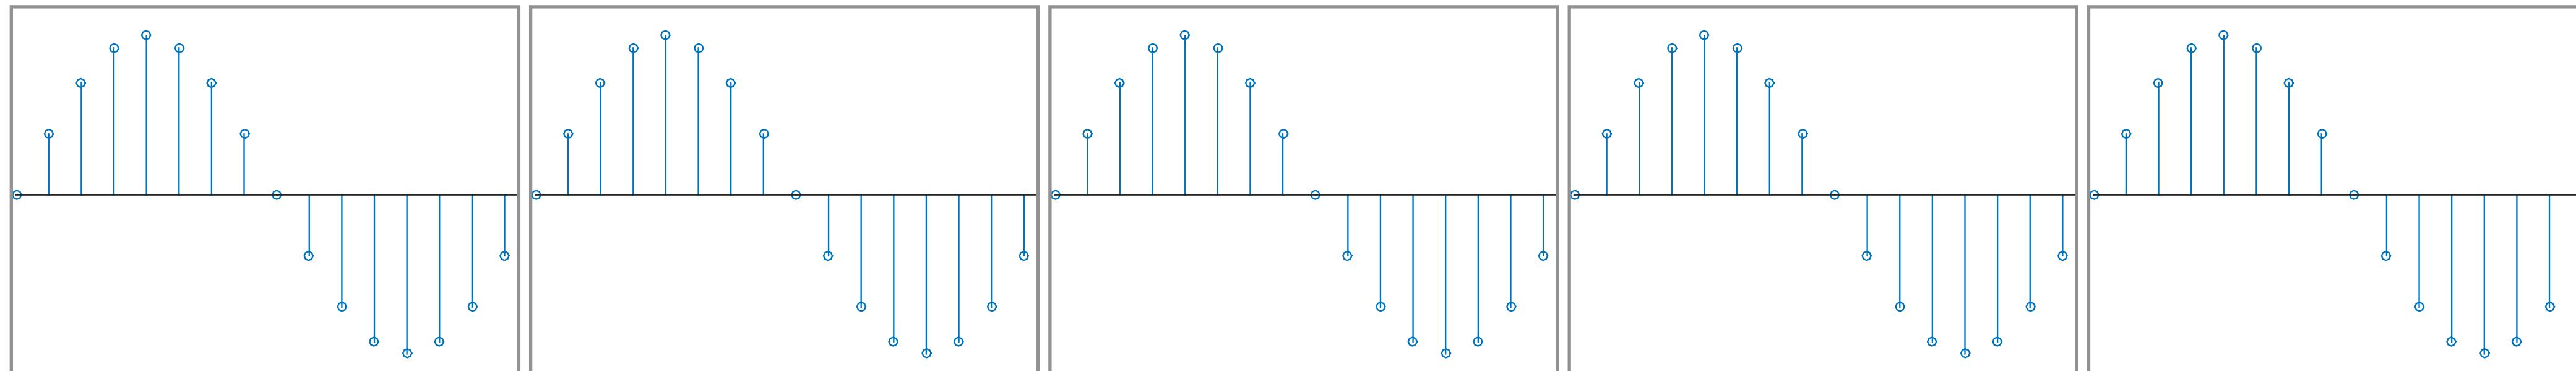


- We had to **remember the phase** between calls to `render()`
- The phase updated at each frame based on the frequency of the oscillator:

$$\phi = \phi + 2\pi f / f_s$$

Wavetable oscillators

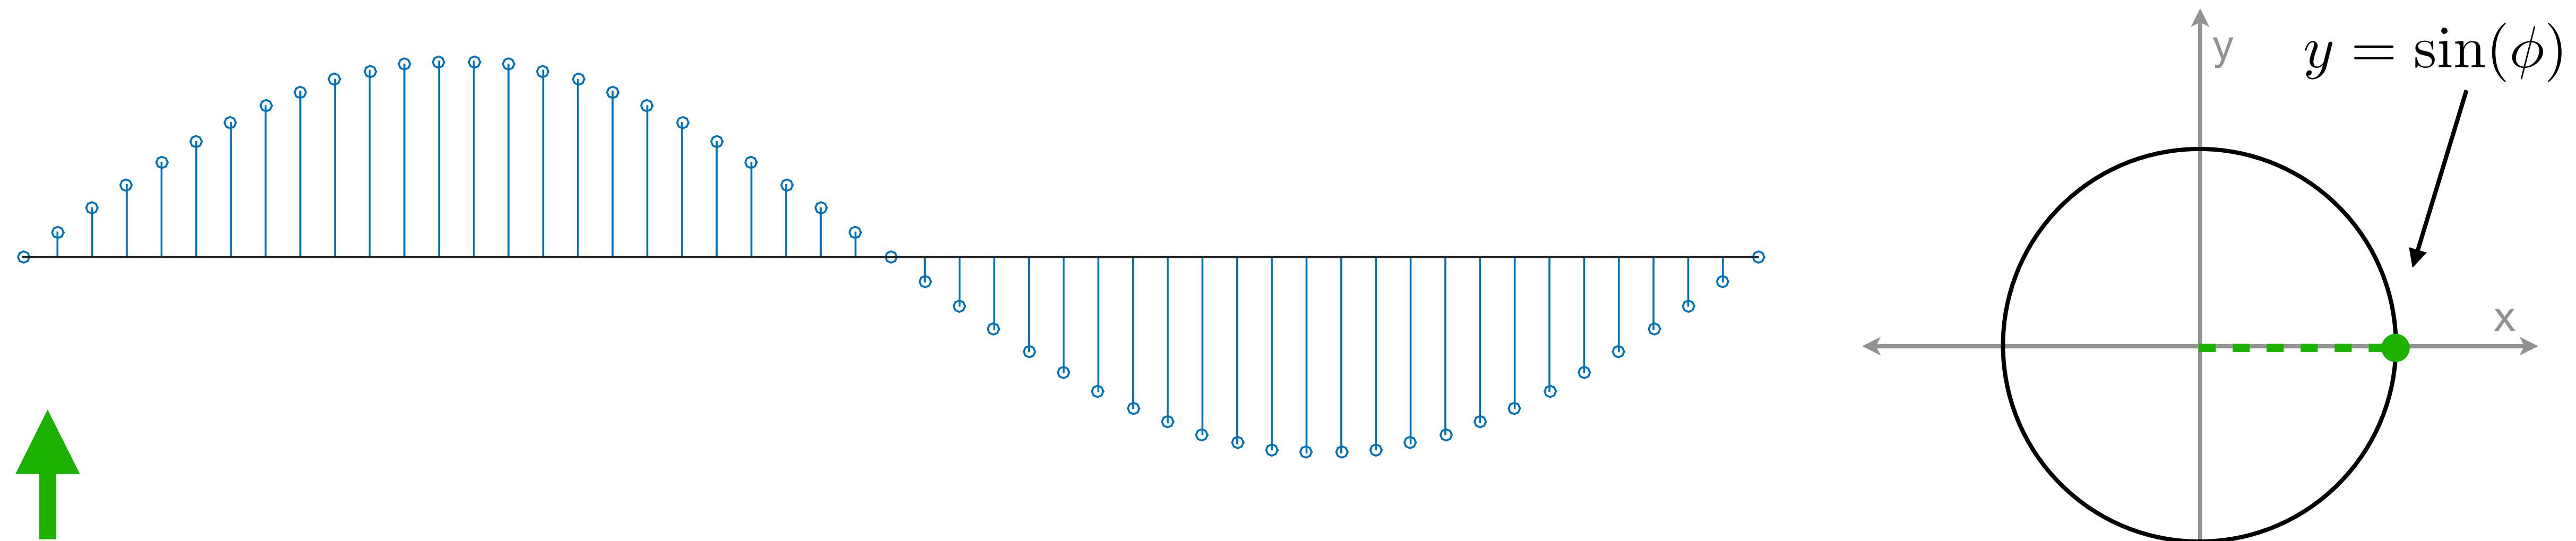
- A **wavetable oscillator** plays a repeating buffer of sound in a loop:



- ▶ Effectively, a **look-up table** of values
- ▶ As we saw in Lecture 3, we can use **interpolation** to get a higher-quality output
- Advantages of wavetables over using functions like `sin()`:
 - ▶ **Less computation time** per sample, because we can calculate the table at the beginning
 - ▶ More **flexibility** on the waveform (which could be nearly any shape)

Wavetable phase and frequency

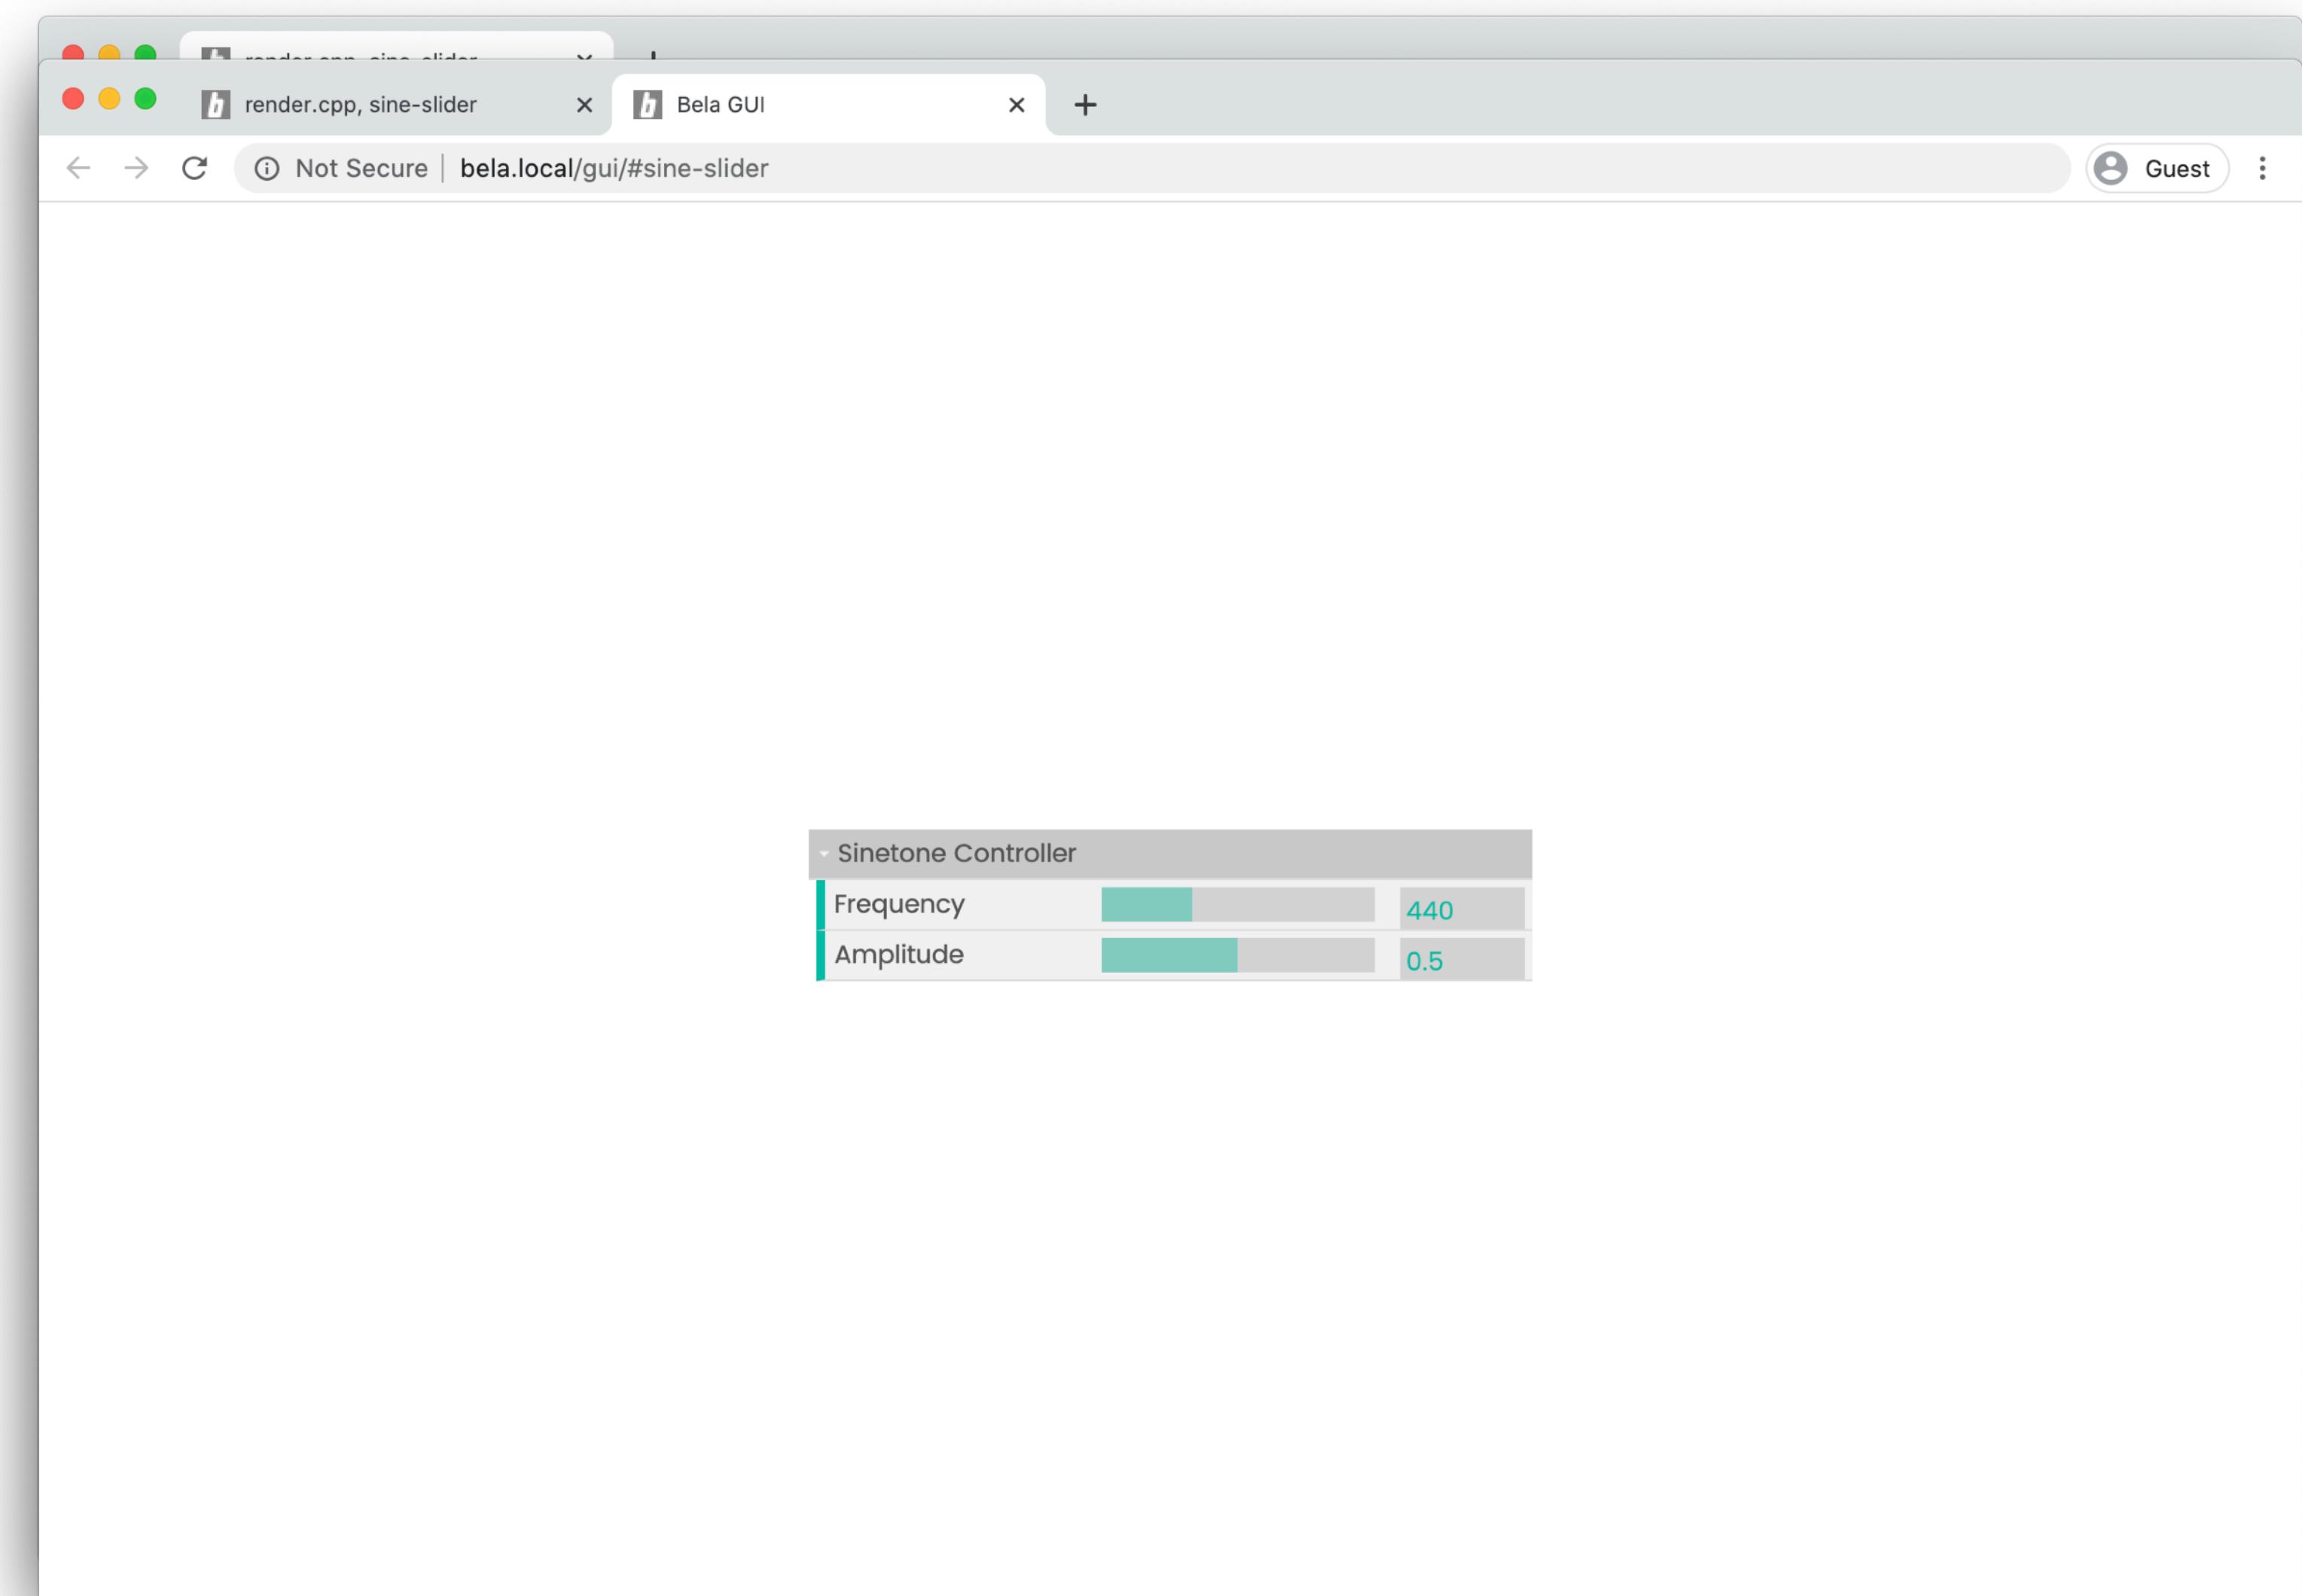
- We need to keep track of where we are in the wavetable
 - i.e. keep track of the **read pointer** in the source buffer
 - We'll need to remember this between calls to `render()`, like before
- The read pointer is analogous to the **phase** in our previous oscillator:



- If the read pointer corresponds to phase, **what determines frequency?**
 - Frequency is the **derivative** (rate of change) of phase
 - Therefore, frequency is given by **rate of change** of the read pointer

Adding controls

- Open the **sine-slider** example from the companion materials:
[github.com/BelaPlatform/
bela-online-course](https://github.com/BelaPlatform/bela-online-course)
- Run the program, then click the **GUI** button
- Adjusting the sliders should change the sound



GUI controls: initialisation

- Bela provides a simple browser-based GUI where you can create sliders to control parameters
- To use it:
 1. Include Gui.h and GuiController.h
 2. Make global Gui and GuiController variables
 3. Initialise the sliders you want in setup()
 4. Get the parameters in render()

```
#include <libraries/Gui/Gui.h>
#include <libraries/GuiController/GuiController.h>

// Browser-based GUI to adjust parameters
Gui gui;
GuiController controller;

// setup() only runs one time
bool setup(BelaContext *context, void *userData)
{
    // Set up the GUI
    gui.setup(context-> projectName);
    controller.setup(&gui, "Sinetone Controller");

    // Arguments: name, default, minimum, maximum, increment
    controller.addSlider("Frequency", 440, 220, 880, 0);
    controller.addSlider("Amplitude", 0.5, 0, 1, 0);

    return true;
}
```

GUI controls: parameter query

frequency and amplitude
are no longer global variables
(why does this still work?)

```
// render() is called every time there is a new block to calculate
void render(BelaContext *context, void *userData)
{
    float frequency = controller.getSliderValue(0); // Frequency is first slider
    float amplitude = controller.getSliderValue(1); // Amplitude is second slider

    // This for() loop goes through all the samples in the block
    for (unsigned int n = 0; n < context->audioFrames; n++) {
        // Increment the phase by one sample's worth at this frequency
        gPhase += 2.0 * M_PI * frequency / context->audioSampleRate;
        if(gPhase >= 2.0 * M_PI)
            gPhase -= 2.0 * M_PI;

        // Calculate a sample of the sine wave
        float out = amplitude * sin(gPhase);
    }
}
```

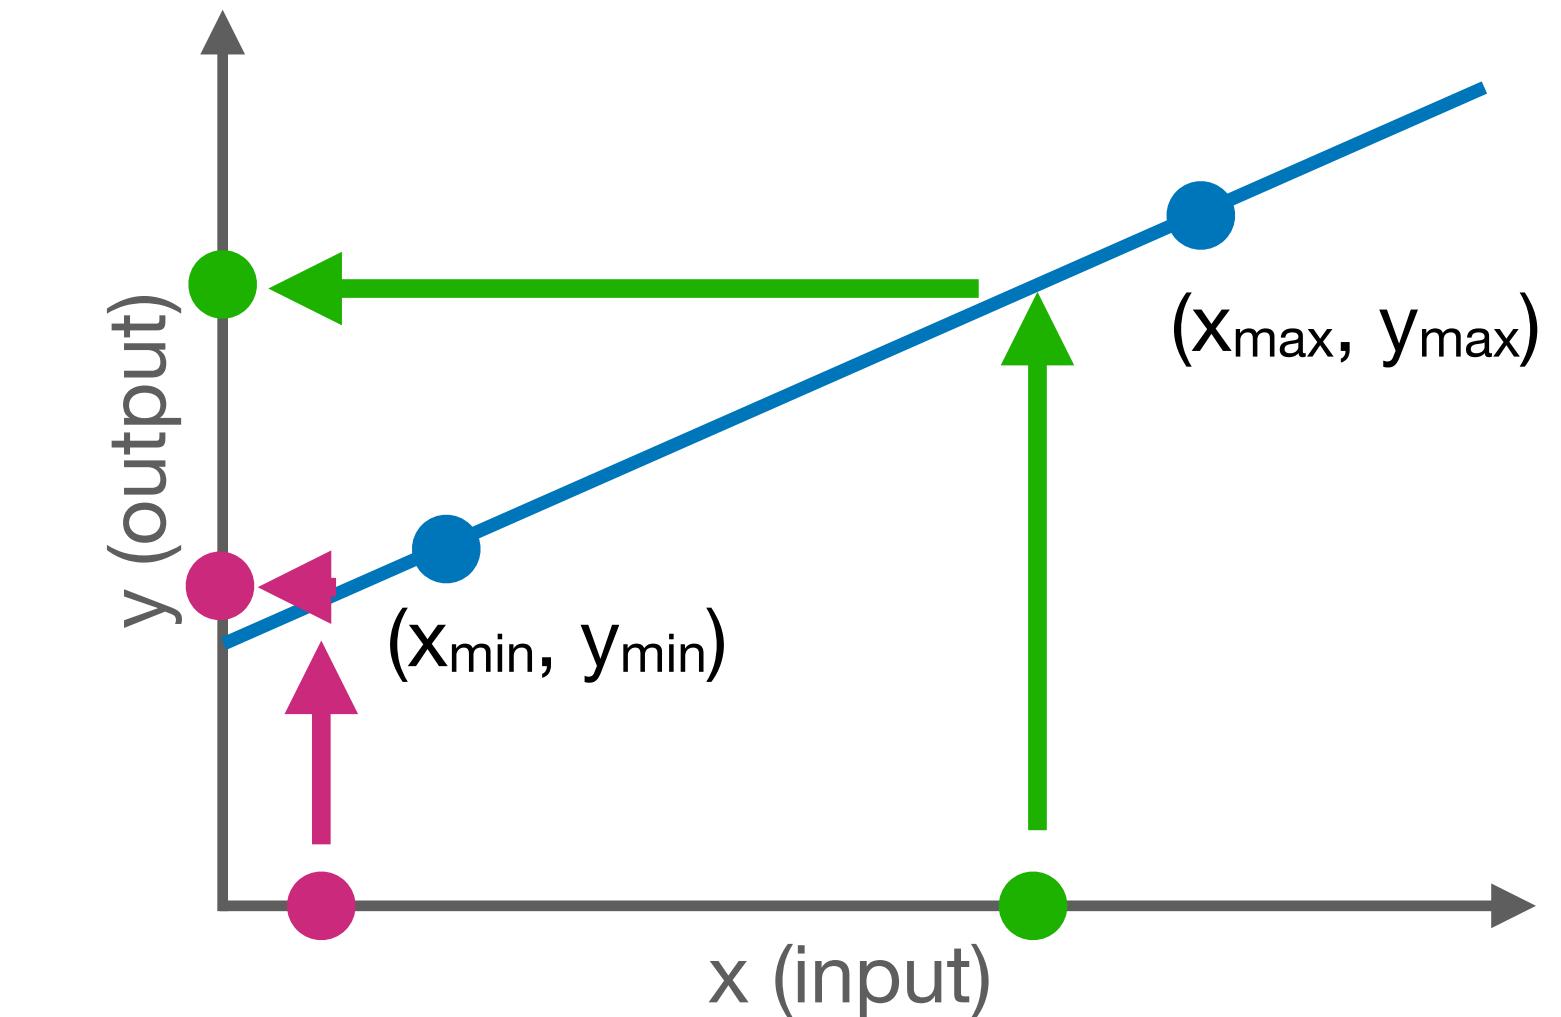
we query the sliders
once per block (not every sample)

Rescaling ranges

- It's common to want to rescale the range of a signal
 - For example, use a control with a 0-1 range to set the frequency of an oscillator
- Bela provides some useful functions for adjusting ranges
- `map()` rescales the input x from one range to another:

```
float map(float x, float in_min, float in_max, float out_min, float out_max)
{
    return (x - in_min) * (out_max - out_min) / (in_max - in_min) + out_min;
}
```

- Equivalently: $y = y_{min} + (y_{max} - y_{min}) \frac{x - x_{min}}{x_{max} - x_{min}}$
- Values outside the input range will be **extrapolated**
- **Task:** rescale the frequency range from 220-880 to a different range (say 110-440)



Constraining ranges

- Sometimes we want to limit the range of a signal
 - For example, by convention, audio signals should always be between -1 and 1
- `constrain()` clips the input `x` to stay within the specified range:

```
float constrain(float x, float min_val, float max_val)
{
    if(x < min_val) return min_val;
    if(x > max_val) return max_val;
    return x;
}
```

- Values within the range will be unchanged
- Values above or below will be constrained to the maximum or minimum

Decibel amplitude scale

- The `map()` function performs linear rescaling
 - However, many real-world phenomena are not linear!
- The **loudness** of a sound is logarithmic with respect to its raw amplitude
 - Human hearing is complex, and perceived loudness also depends on frequency and other factors, but this is the basic relationship
- The **decibel scale** is a logarithmic scale used across science and engineering to express ratios between values
 - For a given ratio A , we could express the same thing in linear or decibel terms:

$$A_{dB} = 20 \log_{10}(A_{linear}) \longleftrightarrow A_{linear} = 10^{\frac{A_{dB}}{20}}$$

- For example, a ratio of 0.1 (1/10) can be written as -20dB
- What is 0.01 (1/100) in dB?
- What is +60dB as a linear ratio?

Reference points for decibel scales

- Remember: the decibel scale is for ratios between two quantities
 - In every case using dBs, we should ask: “relative to what?”
- In digital signal processing, it’s common to see the term dBFS
 - Decibels Full-Scale
 - The maximum value (e.g. 1.0) is defined as 0 dBFS
 - Smaller values (e.g. 0.1, 0.01) correspond to negative numbers (-20dB, -40dB)
 - This is the most common use of dB if no other suffix is specified
- In analog audio systems, we’ll also see dBV and dBu
 - 0 dBV = 1.0V
 - e.g. -20dBV = 0.1V, etc.
 - 0 dBu = 0.775V (= 1mW into a standard 600Ω load)
- Decibels are also used in acoustics for sound pressure level (SPL)
 - The formula is slightly different (based on $10^*\log_{10}$ rather than $20^*\log_{10}$) because it’s measuring power rather than amplitude; we won’t go into this now

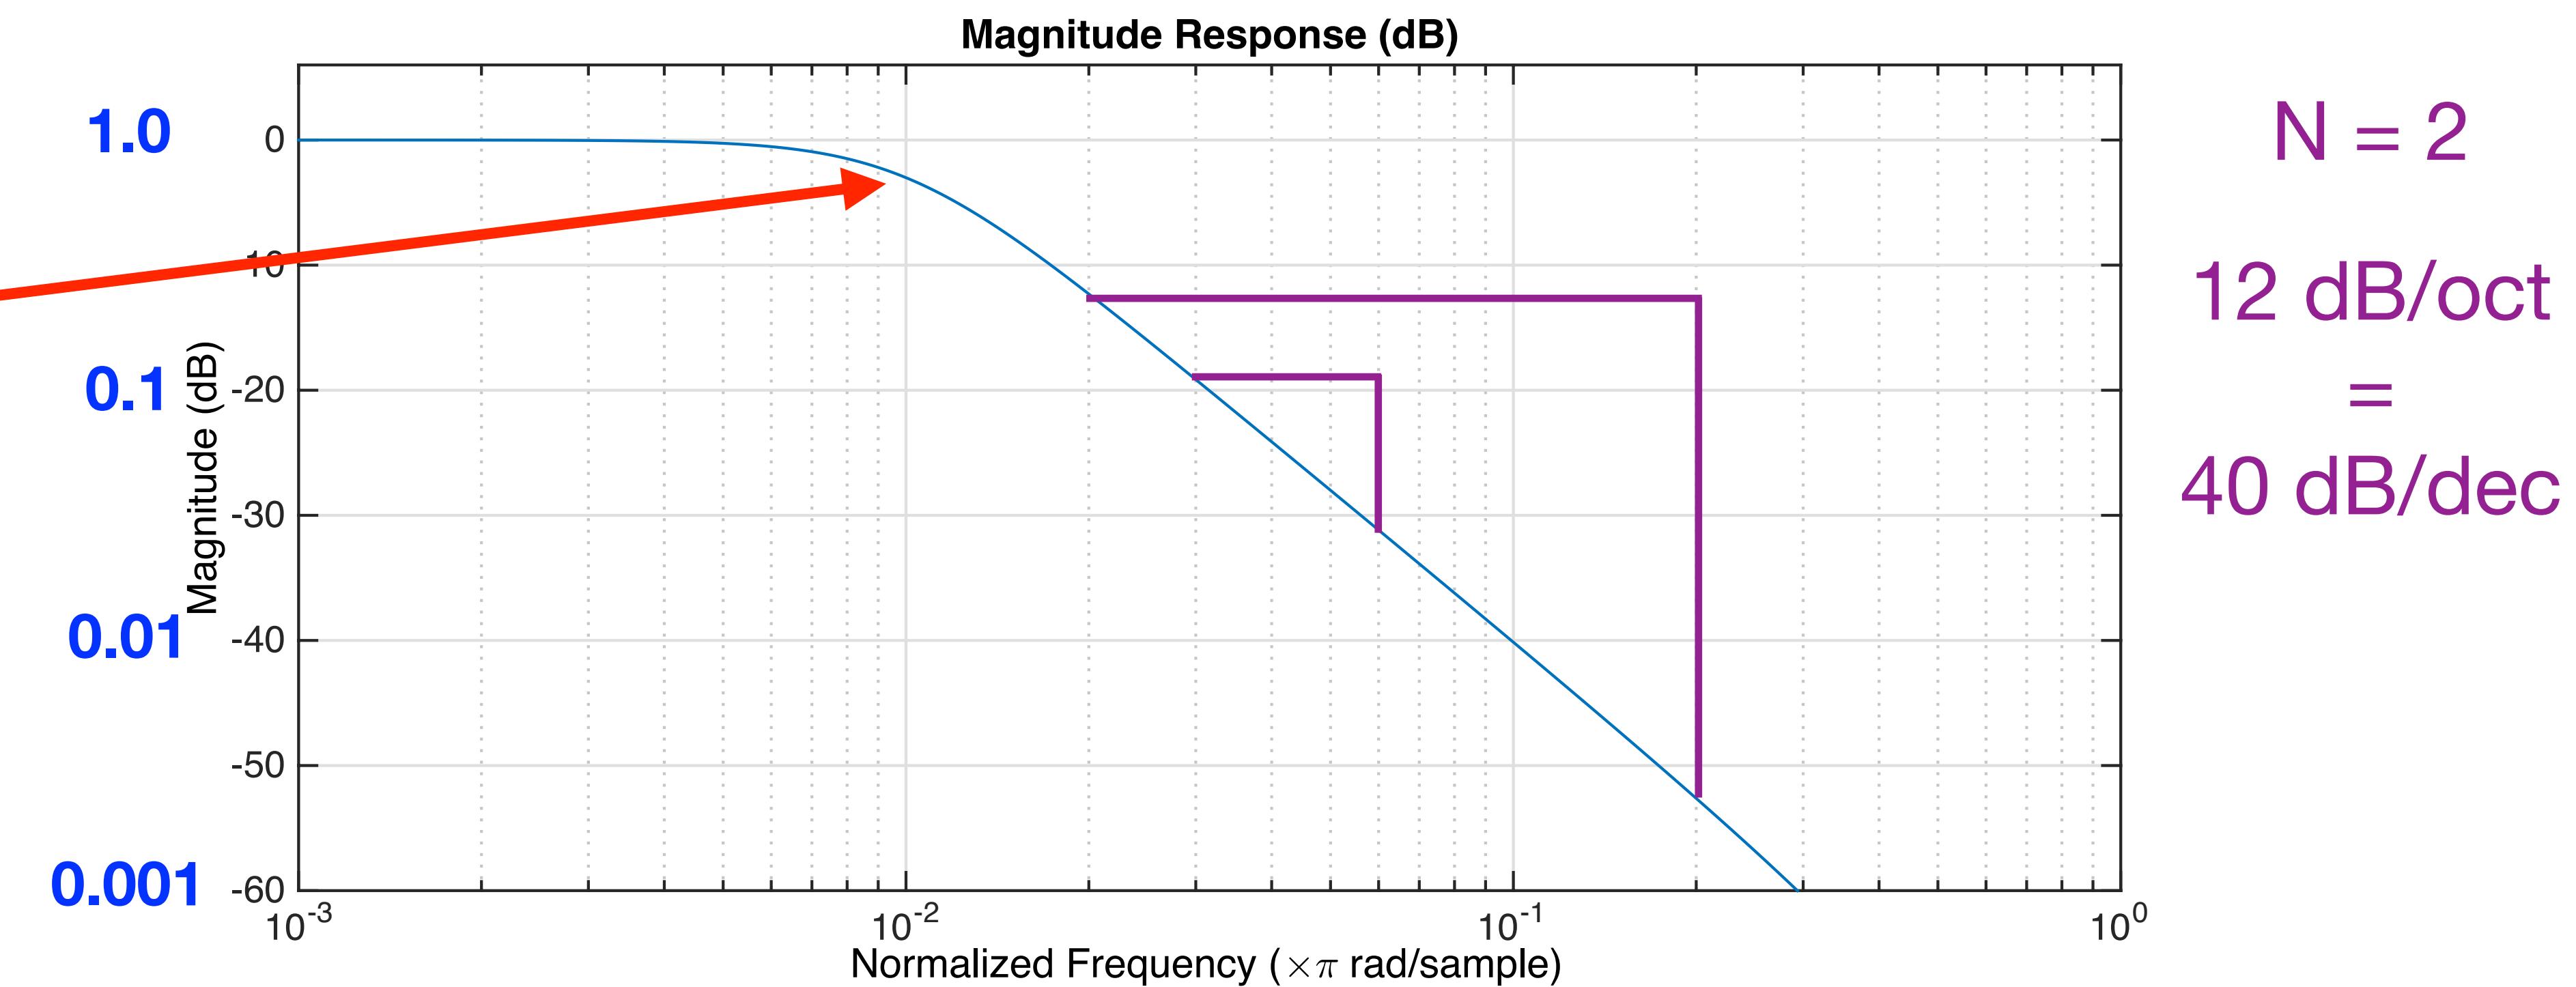
Uses of decibel amplitudes

- Decibels are used to label most **gain controls** and **level meters**
- Decibel scales are often used in specifying **filter magnitude response**
 - By convention, the critical frequency of many standard filter types is the -3dB point
 - The order of a filter determines the **rolloff**, measured in **dB/octave** or **dB/decade**

Butterworth
low-pass filter

$$f_c = 0.01\pi$$

-3dB at this point



Decibel scale task

- **Task:** in the [sine-slider](#) project, implement a decibel scale for amplitude
 - Make the range go from -40 dBFS to 0 dBFS
 - Option 1: change the range of the amplitude slider to -40 to 0
 - Option 2: use the `map()` function to rescale the range 0-1 to go from -40 to 0
 - In either case, implement the decibel formula to calculate the linear value of amplitude that we use to scale the signal: $A_{linear} = 10^{\frac{A_{dB}}{20}}$
 - Hint: `pow(x, y)` implements x^y

Multiple oscillators

- On analog synths, it's common to use **2 or more oscillators** for a single voice
 - The oscillators are tuned nearly, **but not exactly**, to the same frequency
 - Slight **detuning** results in **beat frequencies** according to the mathematical relationship:

$$\sin(x) + \sin(y) = 2 \sin\left(\frac{x+y}{2}\right) \cos\left(\frac{x-y}{2}\right)$$

- Practically speaking, we want 2 controls:
 - One for setting the **centre frequency** of both oscillators (same as before)
 - One for changing the **detuning** (the slight difference between the oscillators)
 - If the centre frequency is ***f*** and the detuning amount is ***r***, we should have:

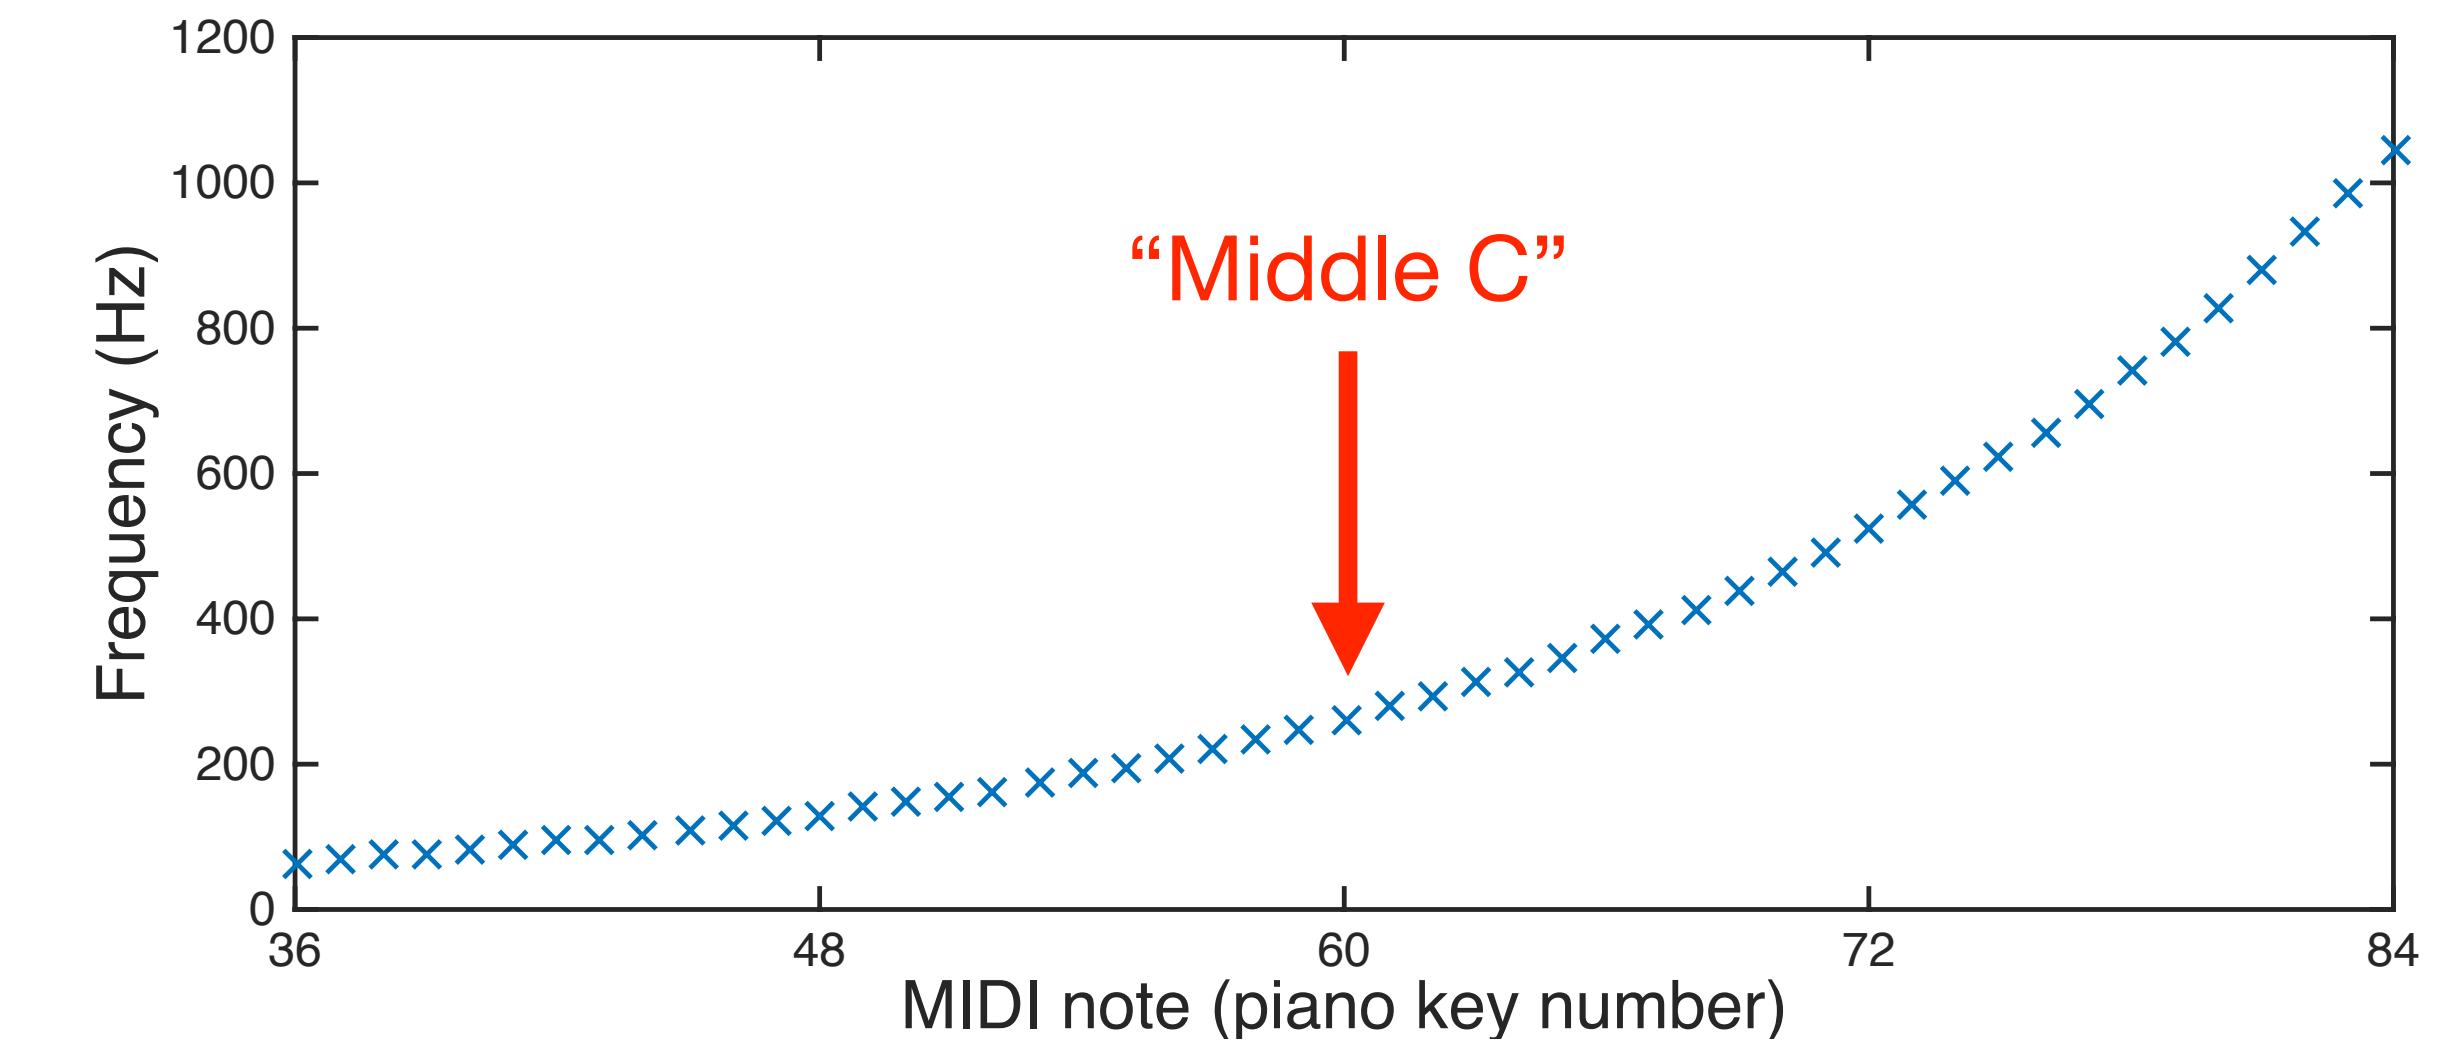
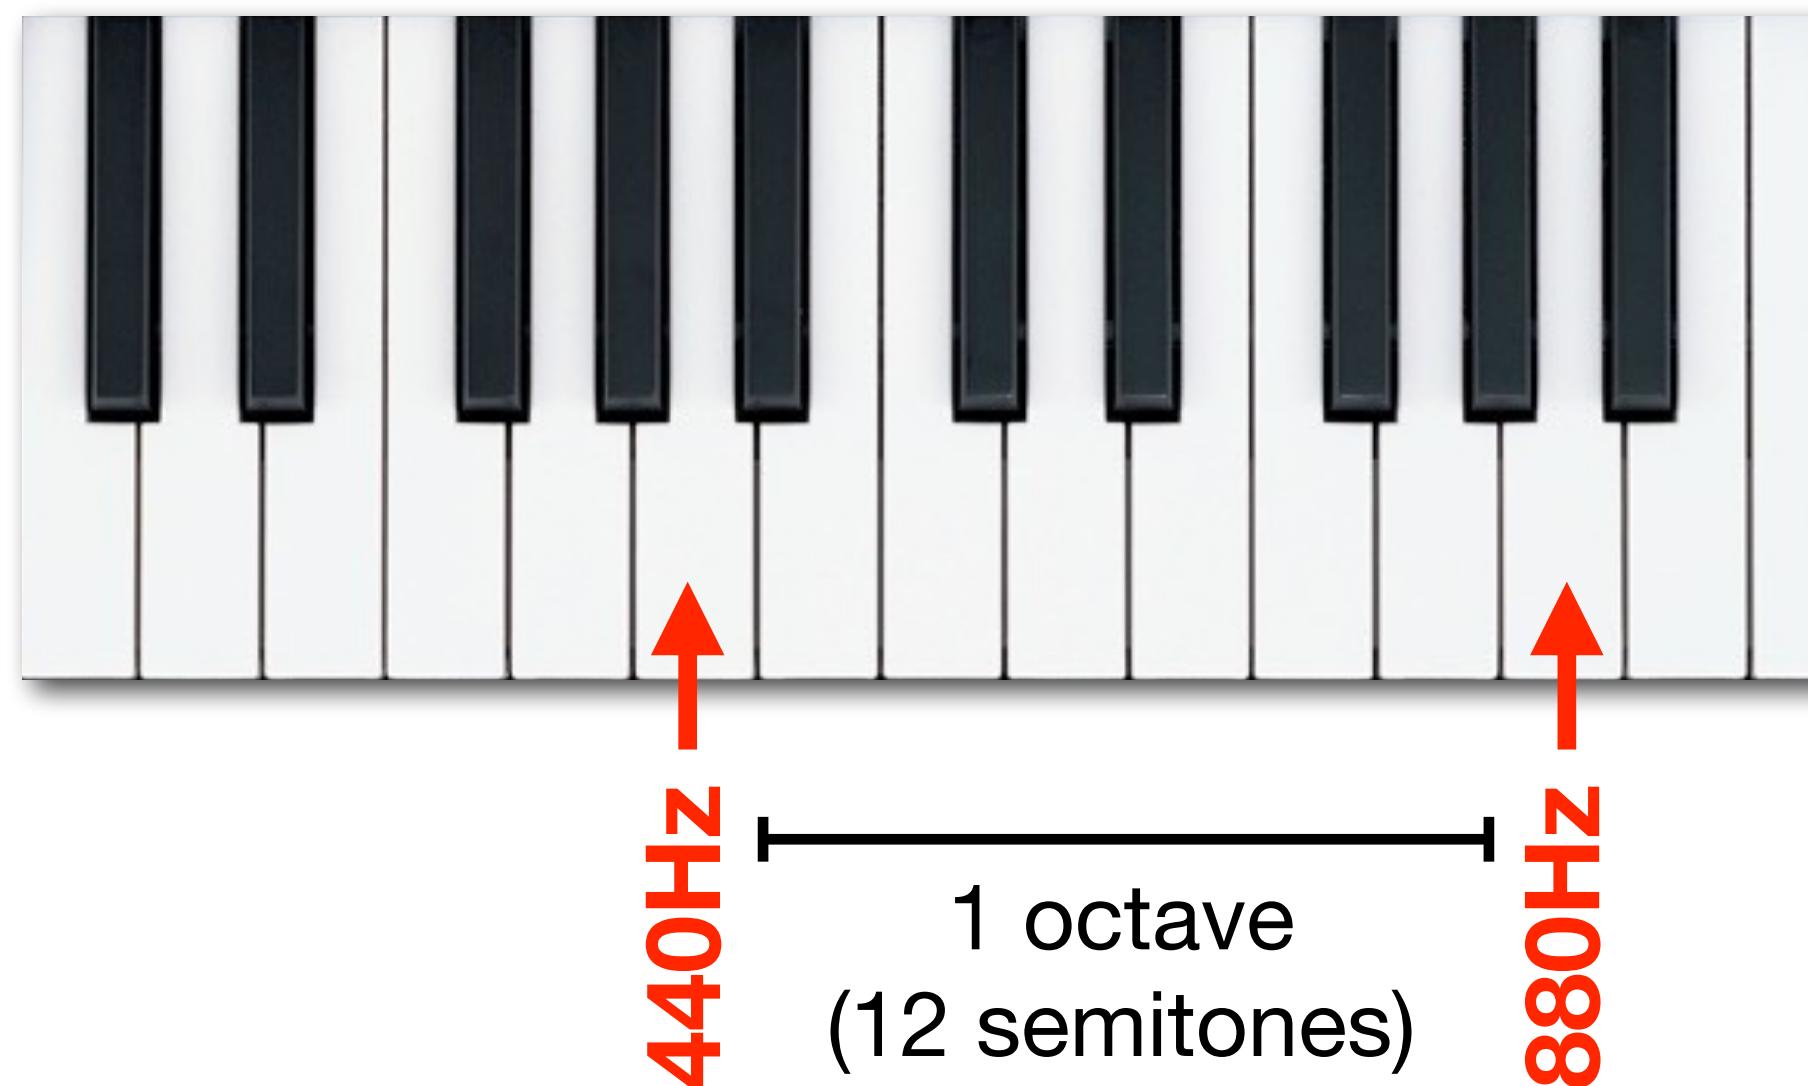
$$f_1 = (1 + r)f \quad f_2 = (1 - r)f$$

- In this case, $r = 0$ corresponds to no detuning (both oscillators the same frequency)

- **Task:** using the **wavetable-slider** example, implement two detuned oscillators
 - Add another GUI control for the detuning ratio

Frequency and pitch

- The perceived **pitch** of a sound is related to its **frequency**
- Pitch perception also operates on a **logarithmic scale**
 - Musical intervals are defined as **frequency ratios**
 - An octave is a 2:1 ratio; a perfect fifth is 3:2, a fourth is 4:3, etc.
 - To get from frequency to notes on a keyboard, we need a logarithmic calculation
 - However, by convention, this calculation is different from decibels

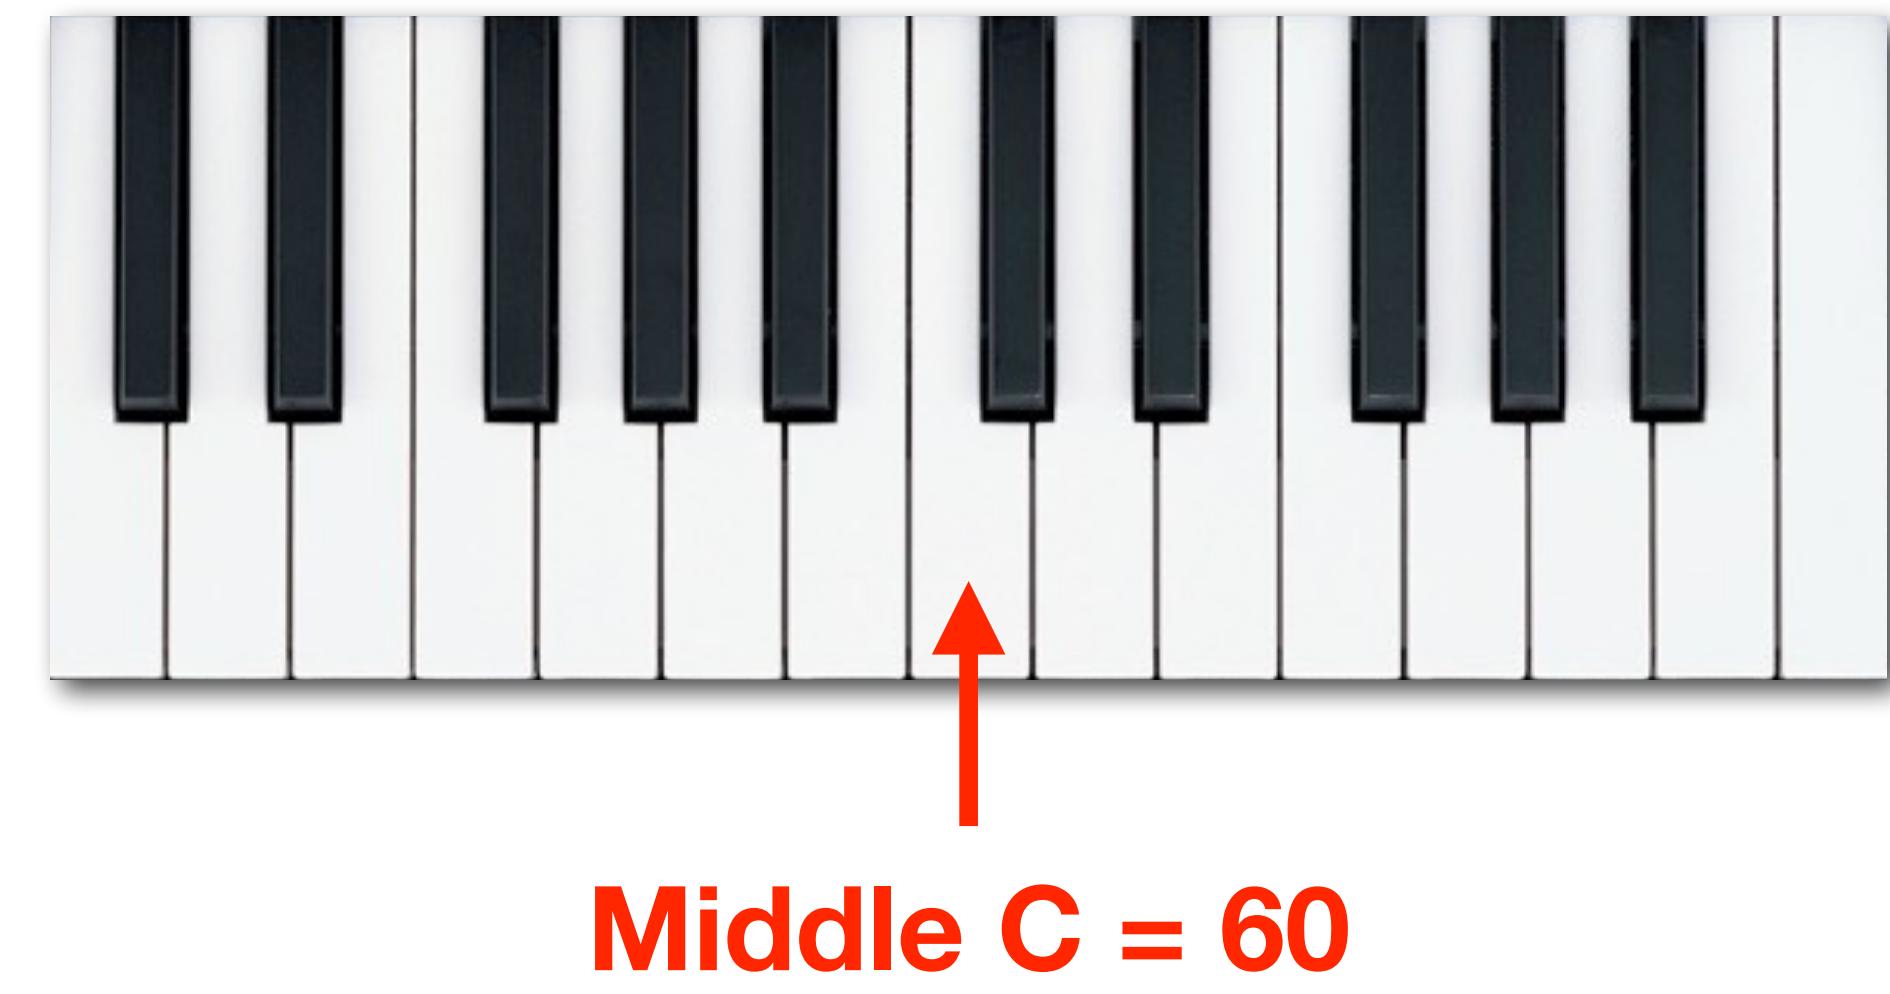


From musical pitch to frequency

- In audio synthesis, we often want to convert from **pitch** (“**note**”) to frequency
 - We want a **linearised scale**: an increment of x is always the same musical interval
 - On analog synths, a common standard is **1V/octave**
 - i.e. the control signal going up 1V makes the pitch go up by an octave (what happens to the frequency?)
- Given a reference frequency f_{ref} and an input V measured in octaves:
 - $f_{out} = f_{ref} 2^V$ (*octaves to frequency*)
 - What if V was measured in semitones instead? (12 semitones to the octave)
 - How do I convert from semitones to octaves?
 - $f_{out} = f_{ref} 2^{V/12}$ (*semitones to frequency*)
- **Task:** in the **wavetable-slider** project:
 - Change the **frequency** slider to be called “**pitch**”, with a range of **0-24 semitones**
 - Implement the frequency calculation above with $f_{ref} = 110\text{Hz}$
 - This should result in a more even pitch change as you move the slider

MIDI note number

- MIDI (Musical Instrument Digital Interface) is a standard protocol for controllers and synths
 - ▶ It defines a **note number** for each key of the piano keyboard, numbered sequentially from low to high
 - ▶ By definition, **middle C** is note number 60
 - It has frequency 261.63Hz
 - ▶ The **A above middle C** is note number 69
 - It has frequency 440Hz (a commonly used reference, sometimes called “**A440**”)
- How should we write the formula to go from MIDI note number to frequency?
 - ▶ For a note number N , how do we calculate number of semitones above A440? $N - 69$
 - Number of octaves? $(N - 69) / 12$
 - ▶ If note 69 corresponds to $f_{ref} = 440\text{Hz}$ how do we calculate the frequency for note N?



$$f_{out} = 440 \cdot 2^{\frac{N-69}{12}} \quad (\text{MIDI note number to frequency})$$

Quantising pitch

$$f_{out} = 440 \cdot 2^{\frac{N-69}{12}} \text{ (MIDI note number to frequency)}$$

- When using MIDI note numbers, each integer corresponds to a piano key
 - We can create a chromatic scale using only integer values
 - We can also feed in fractional note numbers to the formula, even if the MIDI standard doesn't allow this
- In the GUI, we can control the increment of a parameter:

```
controller.addSlider("Note Number", 60, 48, 72, 0);
```



Increment

- Default of 0 means a continuous range (no quantisation)
- Set increment to 1 to get a chromatic scale
- Set increment to 0.5 to hear quarter tones!

Bonus task: diatonic scales

- What if we wanted a different musical scale?
 - Diatonic scales have 7 notes to the octave in a characteristic pattern:



Interval: 2 2 1 2 2 2 1

- We can't implement this with a single value of increment
- Alternatively: use an array to implement the diatonic scale
 - Store the MIDI note numbers in an array, then use the GUI control as an index to the array

Keep in touch!

Social media:

@BelaPlatform

forum.bela.io

blog.bela.io

More resources and contact info at:

learn.bela.io/resources